

Open Science, Reproducibility, and R.

Andrew Stewart

Division of Neuroscience and Experimental Psychology,
University of Manchester

andrew.stewart@manchester.ac.uk

 @ajstewart_lang



Timings...

1000 - Start of morning session

1120 - 1130 Coffee

1130 - Morning session continues

1230 - Lunch

1345 - Start of afternoon session

1445 - 1500 Coffee

1500 - Afternoon session continues

1545 - Building a Binder for fully reproducible research

We have a replication
problem...

Replication and Reproducibility in Science

- Ioannidis (2005), *PLOS Medicine*, most published research findings are false.
- Prinz et al. (2011), *Nature Reviews Drug Discovery*, around 65% of cancer biology studies do not replicate.
- Button et al. (2013), *Nature Reviews Neuroscience*, small sample size undermines the reliability of neuroscience.
- MacLeod et al. (2014), *Lancet*, 85% of biomedical research resources are wasted.
- Baker (2015), *Nature*, 90% of scientists recognise a ‘reproducibility crisis’.
- Nosek & Errington (2017), *eLife*, out of first 5 replication attempts of preclinical cancer biology work, only 2 have replicated.
- Eisner (2018), *Journal of Molecular and Cellular Cardiology*. Reproducibility of science: Fraud, impact factors and carelessness.

Power Posing: Brief Nonverbal Displays Affect Neuroendocrine Levels and Risk Tolerance

Psychological Science
XX(X) 1–6
©The Author(s) 2010
Reprints and permission:
sagepub.com/journalsPermissions.nav
DOI: 10.1177/0956797610383437
<http://pss.sagepub.com>


Dana R. Carney¹, Amy J.C. Cuddy², and Andy J. Yap¹

¹Columbia University and ²Harvard University

Abstract

Humans and other animals express power through open, expansive postures, and they express powerlessness through closed, contractive postures. But can these postures actually cause power? The results of this study confirmed our prediction that posing in high-power nonverbal displays (as opposed to low-power nonverbal displays) would cause neuroendocrine and behavioral changes for both male and female participants: High-power posers experienced elevations in testosterone, decreases in cortisol, and increased feelings of power and tolerance for risk; low-power posers exhibited the opposite pattern. In short, posing in displays of power caused advantaged and adaptive psychological, physiological, and behavioral changes, and these findings suggest that embodiment extends beyond mere thinking and feeling, to physiology and subsequent behavioral choices. That a person can, by assuming two simple 1-min poses, embody power and instantly become more powerful has real-world, actionable implications.

Failed replications or effect sizes much smaller than in the original...

- Power posing
- Ego depletion
- Social priming
- Marshmallow test performance predicts future achievement
- Stanford prison experiment
- Growth mindset
- Any others you know of?

Why are so many studies not replicating?

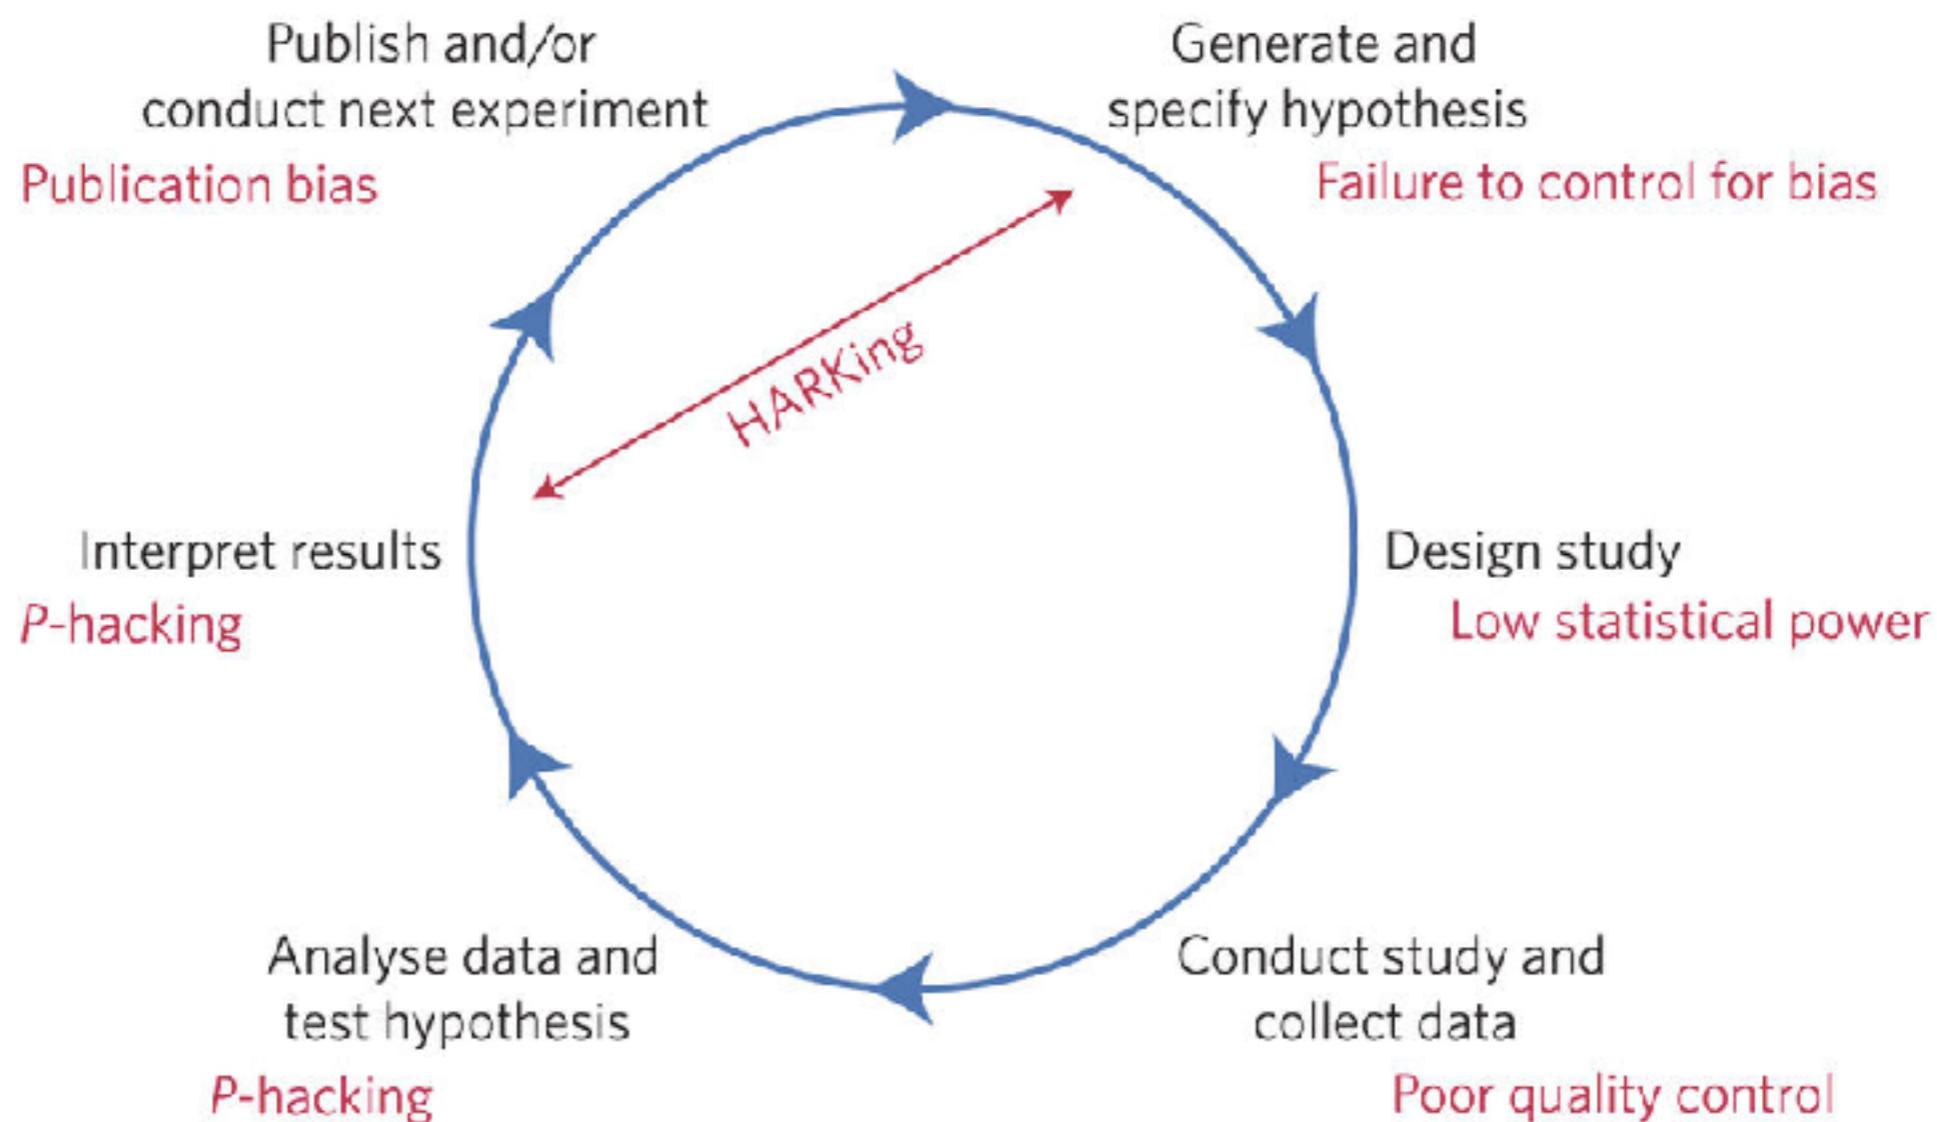
- There are too many studies with experimental power too low to detect the effect size of interest.
- One of the consequences of a low powered study is that when real effects are detected their magnitude is likely to be over-estimated.
- Studies which find the effect are published and studies that don't are not published - due to a bias to publish positive results.
- Future work may use the published effect size during *a priori* power analysis (and then fail to find the effect as the new study is effectively under-powered for what it's looking for).

Is there not just “good science” and “bad science”?

Without realising it, good scientists have been engaging in questionable research practices (QRPs) partly driven by an incentive structure that doesn't incentivise good scientific practice...

Problems include *p*-hacking, lack of power, HARKing, failing (refusal) to share data and code, too many researcher degrees of freedom...

From: A manifesto for reproducible science



Munafo et al. (2017), *Nature Human Behaviour*

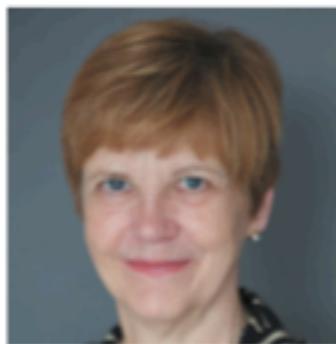
HARKing: Hypothesizing After the Results are Known

Norbert L. Kerr

*Department of Psychology
Michigan State University*

This article considers a practice in scientific communication termed HARKing (Hypothesizing After the Results are Known). HARKing is defined as presenting a post hoc hypothesis (i.e., one based on or informed by one's results) in one's research report as if it were, in fact, an a priori hypotheses. Several forms of HARKing are identified and survey data are presented that suggests that at least some forms of HARKing are widely practiced and widely seen as inappropriate. I identify several reasons why scientists might HARK. Then I discuss several reasons why scientists ought not to HARK. It is conceded that the question of whether HARKing's costs exceed its benefits is a complex one that ought to be addressed through research, open discussion, and debate. To help stimulate such discussion (and for those such as myself who suspect that HARKing's costs do exceed its benefits), I conclude the article with some suggestions for deterring HARKing.

ROBERT TAYLOR



Rein in the four horsemen of irreproducibility

Dorothy Bishop describes how threats to reproducibility, recognized but unaddressed for decades, might finally be brought under control.

More than four decades into my scientific career, I find myself an outlier among academics of similar age and seniority: I strongly identify with the movement to make the practice of science more robust. It's not that my contemporaries are unconcerned about doing science well; it's just that many of them don't seem to recognize that there are serious problems with current practices. By contrast, I think that, in two decades, we will look back on the past 60 years — particularly in biomedical science — and marvel at how much time and money has been wasted on flawed research.

How can that be? We know how to formulate and test hypotheses in controlled experiments. We can account for unwanted variation with statistical techniques. We appreciate the need to replicate observations.

Yet many researchers persist in working in a way almost guaranteed not to deliver meaningful results. They ride with what I refer to as the four horsemen of the reproducibility apocalypse: publication bias, low statistical power, *P*-value hacking and HARKing (hypothesizing after results are known). My generation and the one before us have done little to rein these in.

In 1975, psychologist Anthony Greenwald noted that science is prejudiced against null hypotheses; we even refer to sound work supporting such conclusions as 'failed experiments'. This prejudice leads to publication bias: researchers are less likely to write up studies that show no effect, and journal editors are less likely to accept them. Consequently, no one can learn from them, and researchers waste time and resources

be adequately powered. Other disciplines have yet to catch up.

I stumbled on the issue of *P*-hacking before the term existed. In the 1980s, I reviewed the literature on brain lateralization (how sides of the brain take on different functions) and developmental disorders, and I noticed that, although many studies described links between handedness and dyslexia, the definition of 'atypical handedness' changed from study to study — even within the same research group. I published a sarcastic note, including a simulation to show how easy it was to find an effect if you explored the data after collecting results (D. V. M. Bishop *J. Clin. Exp. Neuropsychol.* **12**, 812–816; 1990). I subsequently noticed similar phenomena in other fields: researchers try out many analyses but report only the ones that are 'statistically significant'.

This practice, now known as *P*-hacking, was once endemic to most branches of science that rely on *P* values to test significance of results, yet few people realized how seriously it could distort findings. That started to change in 2011, with an elegant, comic paper in which the authors crafted analyses to prove that listening to the Beatles could make undergraduates younger (J. P. Simmons *et al. Psychol. Sci.* **22**, 1359–1366; 2011). "Undisclosed flexibility," they wrote, "allows presenting anything as significant."

The term HARKing was coined in 1998 (N. L. Kerr *Pers. Soc. Psychol. Rev.* **2**, 196–217; 1998). Like *P*-hacking, it is so widespread that researchers assume it is good practice. They look at the data, pluck out a finding that looks exciting and write a paper to tell a story around this result. Of course, researchers should be free to explore their

MANY RESEARCHERS
PERSIST IN WORKING
IN A WAY ALMOST
GUARANTEED
NOT
TO DELIVER
MEANINGFUL
RESULTS.

Distinguishing between replicability and reproducibility (note, both are important!)

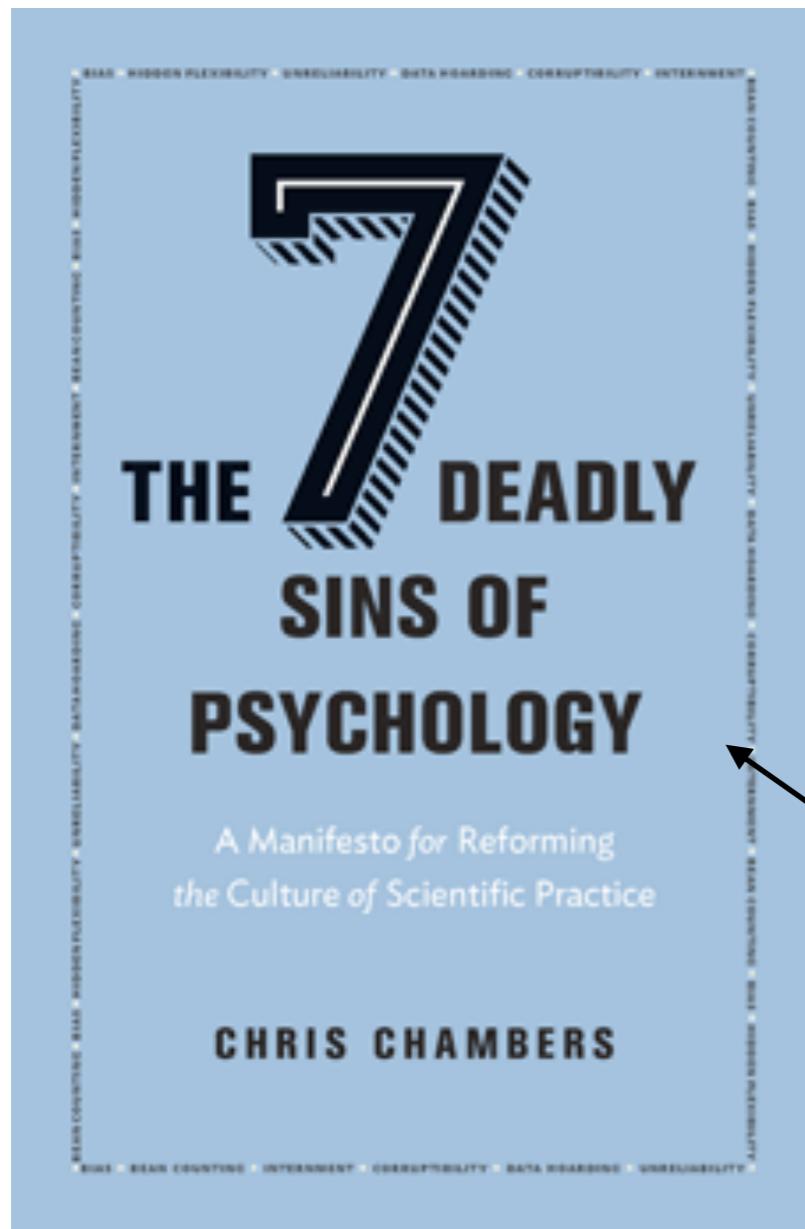
Replicable Science is when someone else can run a study the same as or conceptually equivalent to your one, and find a similar pattern of effects.

Reproducible Science is when someone else can take your data and your analysis code, run it and then find the same effects that you have reported.

How do we make our science more replicable?

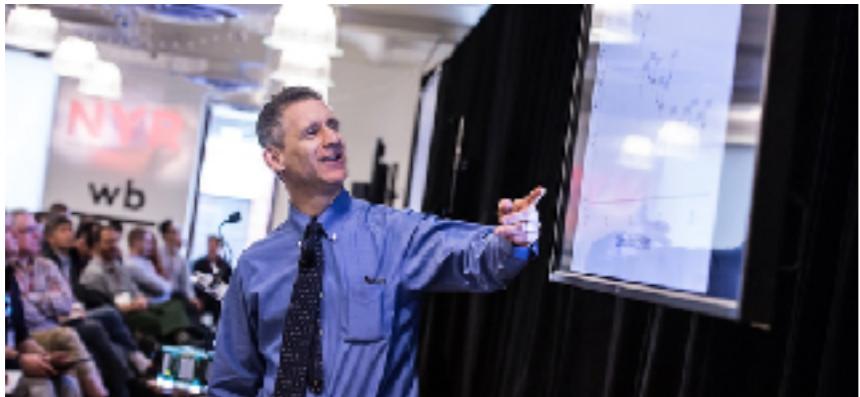
How do we make our science more reproducible?

A move towards open science...



Sins include p-hacking, lack of power, HARKing, failing (refusal) to share data and code, too many researcher degrees of freedom...

You really should read this book!



<http://www.stat.columbia.edu/~gelman/>

Andrew Gelman gives the following recommendations to researchers:

- Analyze all your data.
- Present all your comparisons.
- Make your data public.
- Put in the effort to take accurate measurements (low bias, low variance, and a large enough sample size).
- Do repeated-measures comparisons where possible.

Open Science recently recognised by G7 Science Ministers...

Focus: Incentives and the researcher ecosystem

Ambition: Foster a research environment in which career advancement takes into account Open Science activities, through incentives and rewards for researchers, and valuing the skills and capabilities in the Open Science workforce.

Recommendations:

At national levels: G7 nations should each engage with research stakeholders to identify and implement enhancements to research evaluation and reward systems that take into consideration the Open Science activities carried out by researchers and research institutions. Topics that could be discussed include:

- Recognizing Open Science practices during evaluation of research funding proposals, and research outcomes;
- Recognizing and rewarding research productivity and impact that reflect open science activities by researchers during career advancement reviews;
- Including credit for service activities such as reviewing, evaluating, and curation and management of research data; and,
- Developing metrics of Open Science practices.

Panel criteria and working methods

200. The sub-panels welcome research practice that supports reproducible science and the application of best practice. Examples include registered reports, pre-registration, publication of data sets, experimental materials, analytic code, and use of reporting checklists for publication purposes and those relating to the use of animals in research. These contribute to the evaluation of rigour for submitted outputs. Replication studies may be submitted as outputs and will be evaluated on the extent to which they contribute significant new knowledge, improved methods, or advance theory or practice¹.

346...

Within the context of the institution's strategy, how the submitting unit is progressing towards an open research environment, including where this goes above and beyond the REF open access policy requirements, and wider activity to encourage the effective sharing and management of research data, as appropriate to the discipline. Consideration of reproducibility should also be included where relevant to the discipline.

and is forming part of Universities' teaching manifestos.

Teaching with Open Science commitment:

To teach the practices and skills of open research and science in our undergraduate and postgraduate degree programmes

- a. Promote open science in our teaching.
- b. Design a Research Methods curriculum that teaches skills for open science and uses open science to enhance teaching (for example: teach R and use open data to practice analysis skills).
- c. Learn about and adopt open educational practices in our teaching.
- d. Produce and promote tools for helping student researchers adopt open practices, including training and guidance suitable to their level of study.
- e. Author, share and use open educational resources to promote teaching with open science beyond our School and Institution.
- f. Support our colleagues to learn the skills of teaching Open Science.

**How do you do Open
Science?**

Before Data Collection

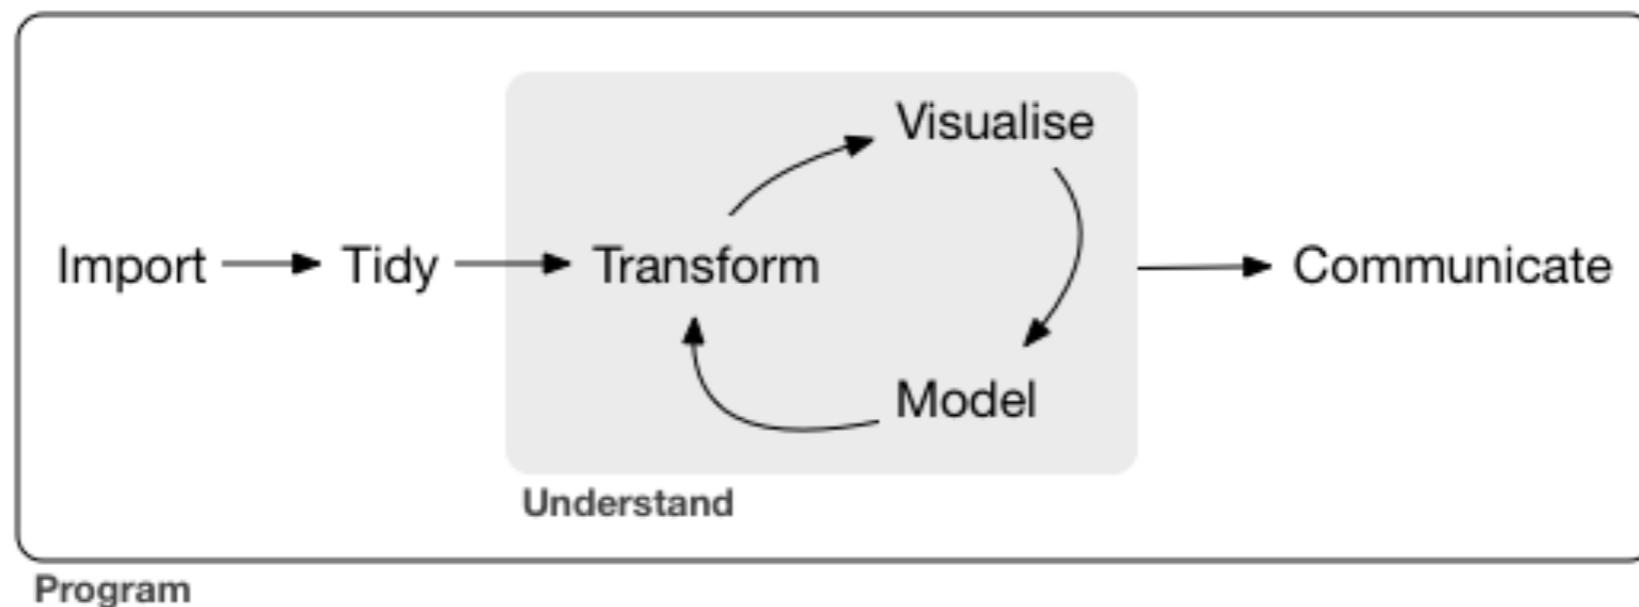
- Specify your hypotheses and analysis plan.
- **Pre-register** your hypotheses and analysis plan at osf.io
- Consider data simulation so that you can write your analysis script before you have your real data.
- Consider submitting as a **registered report** - currently **186** journals now support this route. This involves acceptance in principle before you have even started collecting your data.

Registered Reports



After Data Collection

- You need to use analysis software that allows for open sharing and reproducibility of the entire data wrangling/analysis/write-up workflow.



Hadley Wickham and Garrett Grolemund

- You can share your data at osf.io or on GitHub:

[ajstewartlang / Comprehension-of-indirect-requests-is-influenced-by-their-degree-of-imposition](#)

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

Branch: master [Comprehension-of-indirect-requests-is-influenced-by-their-degree-of-imposition / RP.csv](#) Find file Copy path

 ajstewartlang Made consistent the labelling of factors in data files and in paper 7b3b3b1 on 29 Mar 2017

0 contributors

1681 Lines (1681 sloc) 69.7 KB										
	Raw	Blame	History							
1	P.s	Item	Condition	Probmanip	Speaker	statement	response	final	Meaning	Imposition
2	1	1	1	1708	302	1389	1867	1206	Indirect	High
3	1	2	2	1166	296	1377	1671	828	Indirect	Low
4	1	3	3	1393		1484	1950	1812	Direct	High
5	1	4	4	2463	630	1691	1866	966	Direct	Low
6	1	5	1	1662	267	1332	1477	1346	Indirect	High
7	1	6	2	1446	444	1004	1067	797	Indirect	Low
8	1	7	3	2159	501	739	1231	2240	Direct	High
9	1	8	4	1459		1086	946	978	Direct	Low
10	1	9	1	3302		1503	900	1736	Indirect	High

- alongside your analysis code

```
--  
26 FPs$Meaning <- as.factor(FPs$Meaning)  
27 FPs$Imposition <- as.factor(FPs$Imposition)  
28  
29 #this sets up the contrasts so that the intercept in the mixed LMM is the grand mean (i.e., the mean of all conditions)  
30 my.coding <- matrix (c(.5, -.5))  
31  
32 contrasts (FPs$Meaning) <- my.coding  
33 contrasts (FPs$Imposition) <- my.coding  
34  
35 #construct the models with crossed random effects for subjects and items for the pre-critical, critical and post-critical region  
36 fpmodelprec <- lmer (Prohmanip ~ Meaning*Imposition + (1+Meaning*Imposition | P.s) + (1+Meaning+Imposition | Item), data=FPs, REML=F  
37 summary (fpmodelprec)  
38 lsmeans (fpmodelprec, pairwise~Meaning*Imposition, adjust="none")  
39  
40 fpmodelc <- lmer (statement ~ Meaning*Imposition + (1+Meaning*Imposition | P.s) + (1+Meaning*Imposition | Item), data=FPs, REML=T  
41 summary (fpmodelc)  
42 lsmeans (fpmodelc, pairwise~Meaning*Imposition, adjust="none")  
43  
44 fpmodelpostc <- lmer (response ~ Meaning*Imposition + (1+Meaning*Imposition | P.s) + (1+Meaning+Imposition | Item), data=FPs, REML=F  
45 summary (fpmodelpostc)  
46 lsmeans (fpmodelpostc, pairwise~Meaning*Imposition, adjust="none")  
47  
48 #Regression Path Analysis  
49 #Read in Regression Path data  
50 RPs <- read.csv("~/RPs.csv")  
51  
52 RPs$Meaning <- as.factor(RPs$Meaning)  
53 RPs$Imposition <- as.factor(RPs$Imposition)  
54  
55 contrasts (RPs$Meaning) <- my.coding  
56 contrasts (RPs$Imposition) <- my.coding  
57  
58 #construct the models with crossed random effects for subjects and items for the pre-critical, critical and post-critical region  
59 rpnodelprec <- lmer (Prohmanip ~ Meaning*Imposition + (1+Meaning*Imposition | P.s) + (1+Meaning*Imposition | Item), data=RPs, REML=F
```

And preserve it with a DOI via Zenodo

The screenshot shows a web browser window with the URL zenodo.org/account/settings/github. The page is titled "GitHub Repositories".

The left sidebar has a "Settings" section with the following items:

- Profile
- Change password
- Security
- Linked accounts
- Applications
- Shared links
- GitHub** (selected)

The main content area is titled "GitHub Repositories" and shows the status "(Updated 21 seconds ago) Sync now...". It features three steps to get started:

- 1 Flip the switch**: Select the repository you want to preserve, and toggle the switch below to turn on automatic preservation of your software. The switch is currently set to **ON**.
- 2 Create a release**: Go to GitHub and [create a release](#). Zenodo will automatically download a .zip ball of each new release and register a DOI.
- 3 Get the badge**: After your first release, a DOI badge that you can include in GitHub README will appear next to your repository below.

Below this, there's a "Repositories" section with a note: "If your organization's repositories do not show up in the list, please ensure you have enabled [third-party access](#) to the Zenodo application. Private repositories are not supported." A specific repository entry is shown:

[ajstewartlang/Affective-Theory-of-Mind-Inferences](#)

Qualitative

- You can pre-register your analysis protocol in advance - much like with quantitative research.
- Although there are some caveats around making qualitative data fully open (espec. in light of confidentiality issues), there is a good discussion of some possible solutions here:

[https://www.memedpublish.org/manuscripts/1986](https://www.mededpublish.org/manuscripts/1986)

Using R for Reproducible Data Science

“Hadley Wickham, the Man Who Revolutionized R”



Chief Scientist at RStudio, author of key R packages incl. `ggplot2`, `tidyverse`, `dplyr` - all components of the tidyverse.

Data Science is a growing employment destination for Psychologists.



AMERICAN PSYCHOLOGICAL ASSOCIATION

ABOUT APA

TOPICS

PUBLICATIONS & DATABASES

PSYCHOLOGY HELP CENTER

NEWS & EVENTS

SCIENCE

[Home](#) // [gradPSYCH Magazine](#) // [January 2013 gradPSYCH](#) // [Hot jobs: Big-data psychologists](#)

CAREER CENTER

Hot jobs: Big-data psychologists

Wanted: Scientists who can help companies make sense of consumer and employee data.



By Rebecca Voelker

Print version: page 18

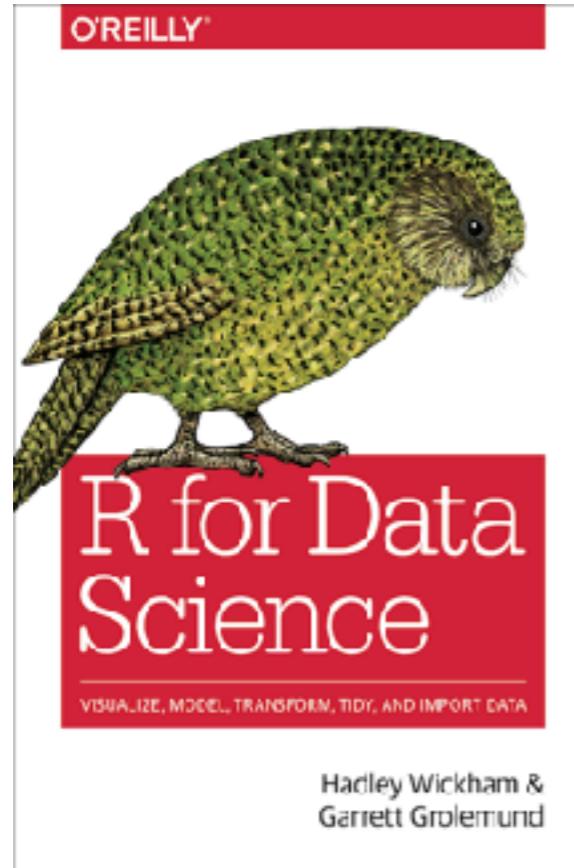
Every time you use an Internet search engine, sign up for a company "rewards" program or swipe your credit card, that information is saved and stored. The industries that amass these billions of bytes of data are increasingly hiring psychologists to help make sense of it, says Douglas Reynolds, PhD, president of APAs Div. 14 (Society for Industrial and Organizational Psychology).

"Big data, and what it means for business, is a hot topic right now," says Reynolds, who also serves as vice president of assessment technology at Development Dimensions International Inc., a human resources consulting firm. "We're in high demand."

Psychologists' ability to interpret numbers and human behavior makes them key members of many industry analytics teams, adds Suzie Weaver, PhD, a psychologist and senior analytic consultant with Epsilon, a global marketing and analytics company. "Analytics is at the core of everything we do, whether it's in research, the academic sphere or the business world."

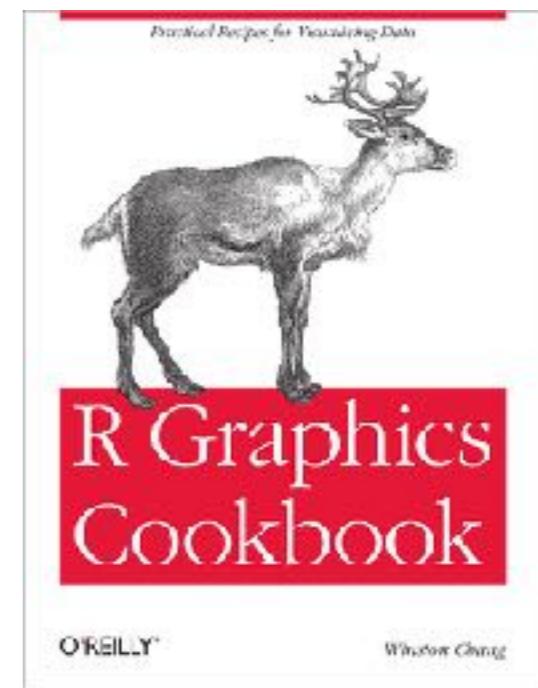
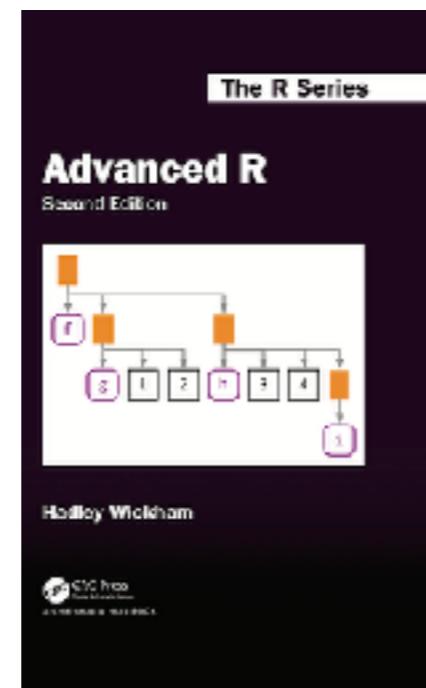
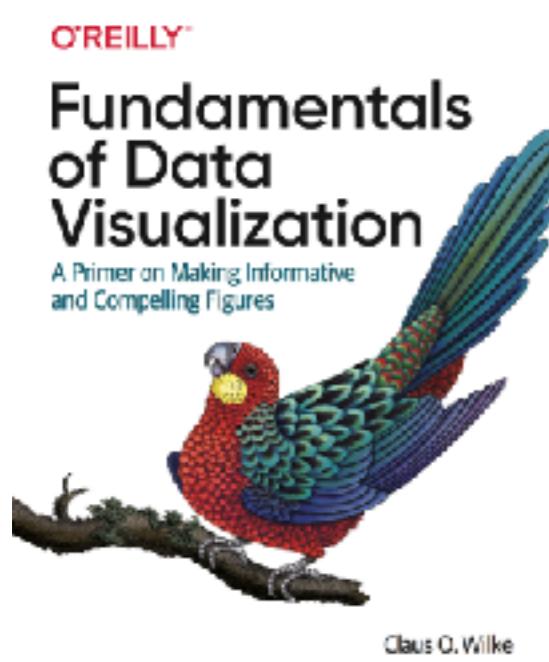
What role can R play in Open Science?

- R scripts are easy to share allowing for reproducibility and easy public sharing of data and code.
- R is free, open source software that is much more flexible and powerful than SPSS.
- There is an active R community continuously updating statistical tests and packages that run in R.
- As R is a programming language, it forces you to know your data.



Available electronically for free at:

<http://r4ds.had.co.nz>



How to create BBC style graphics

Load all the libraries you need

Install the `bbplot` package

How does the `bbplot` package work?

Save out your finished chart

Make a line chart.

Make a multiple line chart

Make a bar chart

Make a stacked bar chart

Make a grouped bar chart

Make a dumbbell chart

Make a histogram

Make changes to the legend

Make changes to the axes

Add annotations

Work with small multiples

Do something else entirely

BBC Visual and Data Journalism cookbook for R graphics

Last updated: 2015-01-24

How to create BBC style graphics

At the BBC data team, we have developed an R package and an R cookbook to make the process of creating publication-ready graphics in our in-house style using R's ggplot2 library a more reproducible process, as well as making it easier for people new to R to create graphics.

The cookbook below should hopefully help anyone who wants to make graphics like these.



Handy list of Psychology groups that teach R, plus links to course materials - list compiled by Andy Wills at Plymouth.

[View on GitHub](#) 

rminr

Research Methods in R

Teaching Research Methods in R

This is a crowd-sourced list of uses of R to teach research methods in Psychology, and a link to Creative Commons teaching materials, where these are available. The year teaching in R was adopted at undergraduate and postgraduate level is also recorded, where known. Where there are no materials, but the organization's name has a link, this is a link to evidence that R is used.

If you'd like to add to this list, please submit a [pull request](#). Or, if you're not sure how to do that, just email me: andy@willslab.co.uk

Universities

University	Country	UG	PG	Link
Harrisburg University of Science and Technology	U.S.A.		2018	PG
Missouri State	U.S.A.		2017	PG
Nottingham Trent University	U.K.	2012	2010	
University of Edinburgh	U.K.	2018	2018	
University of Glasgow	U.K.	2015	2010	UG, PG
University of Lancaster	U.K.		2014	
University of Lincoln	U.K.		2018	PG
University of Manchester	U.K.		2018	PG
University of Plymouth	U.K.	2018 (Year 1) - 2020 (Year 3)	2017	UG, PG
University of Sussex	U.K.	2019		

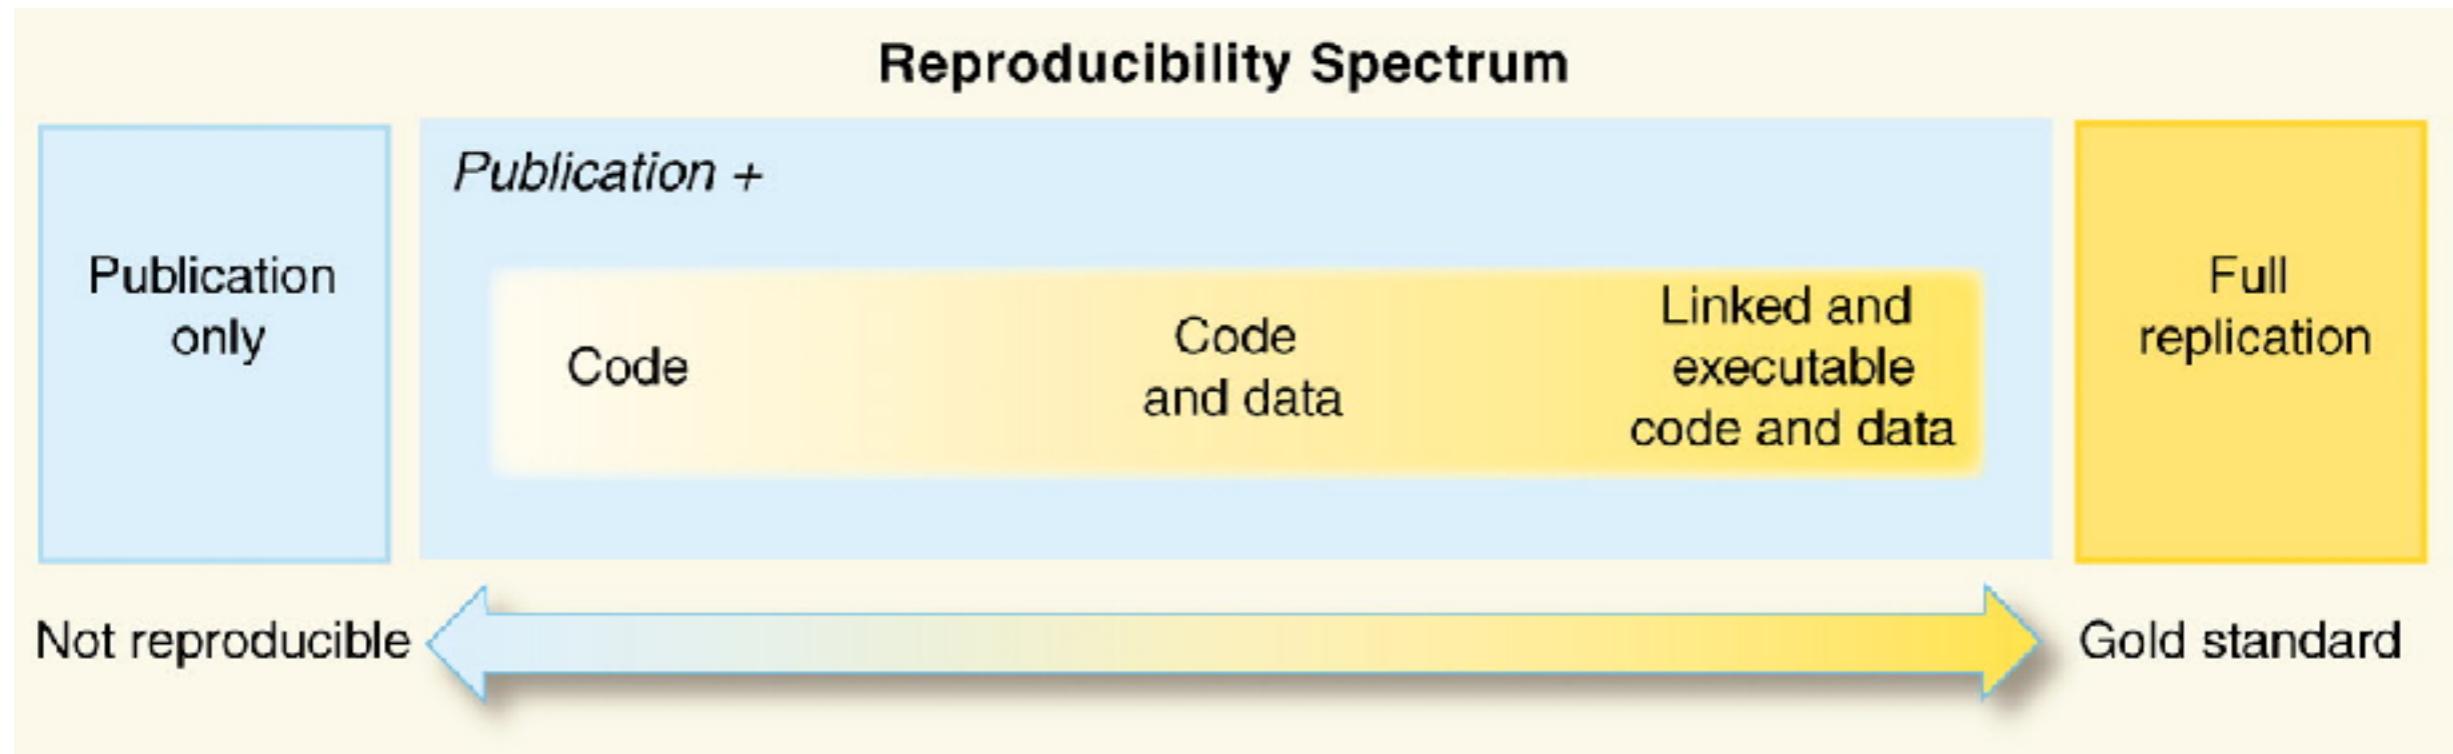
<https://ajwills72.github.io/rminr/rminrinpsy.html>

Reproducible Research in Computational Science

Roger D. Peng

[+ See all authors and affiliations](#)

Science 02 Dec 2011;
Vol. 334, Issue 6060, pp. 1226-1227
DOI: 10.1126/science.1213847



"*You can't do reproducible research in a GUI*", Hadley Wickham (probably!)

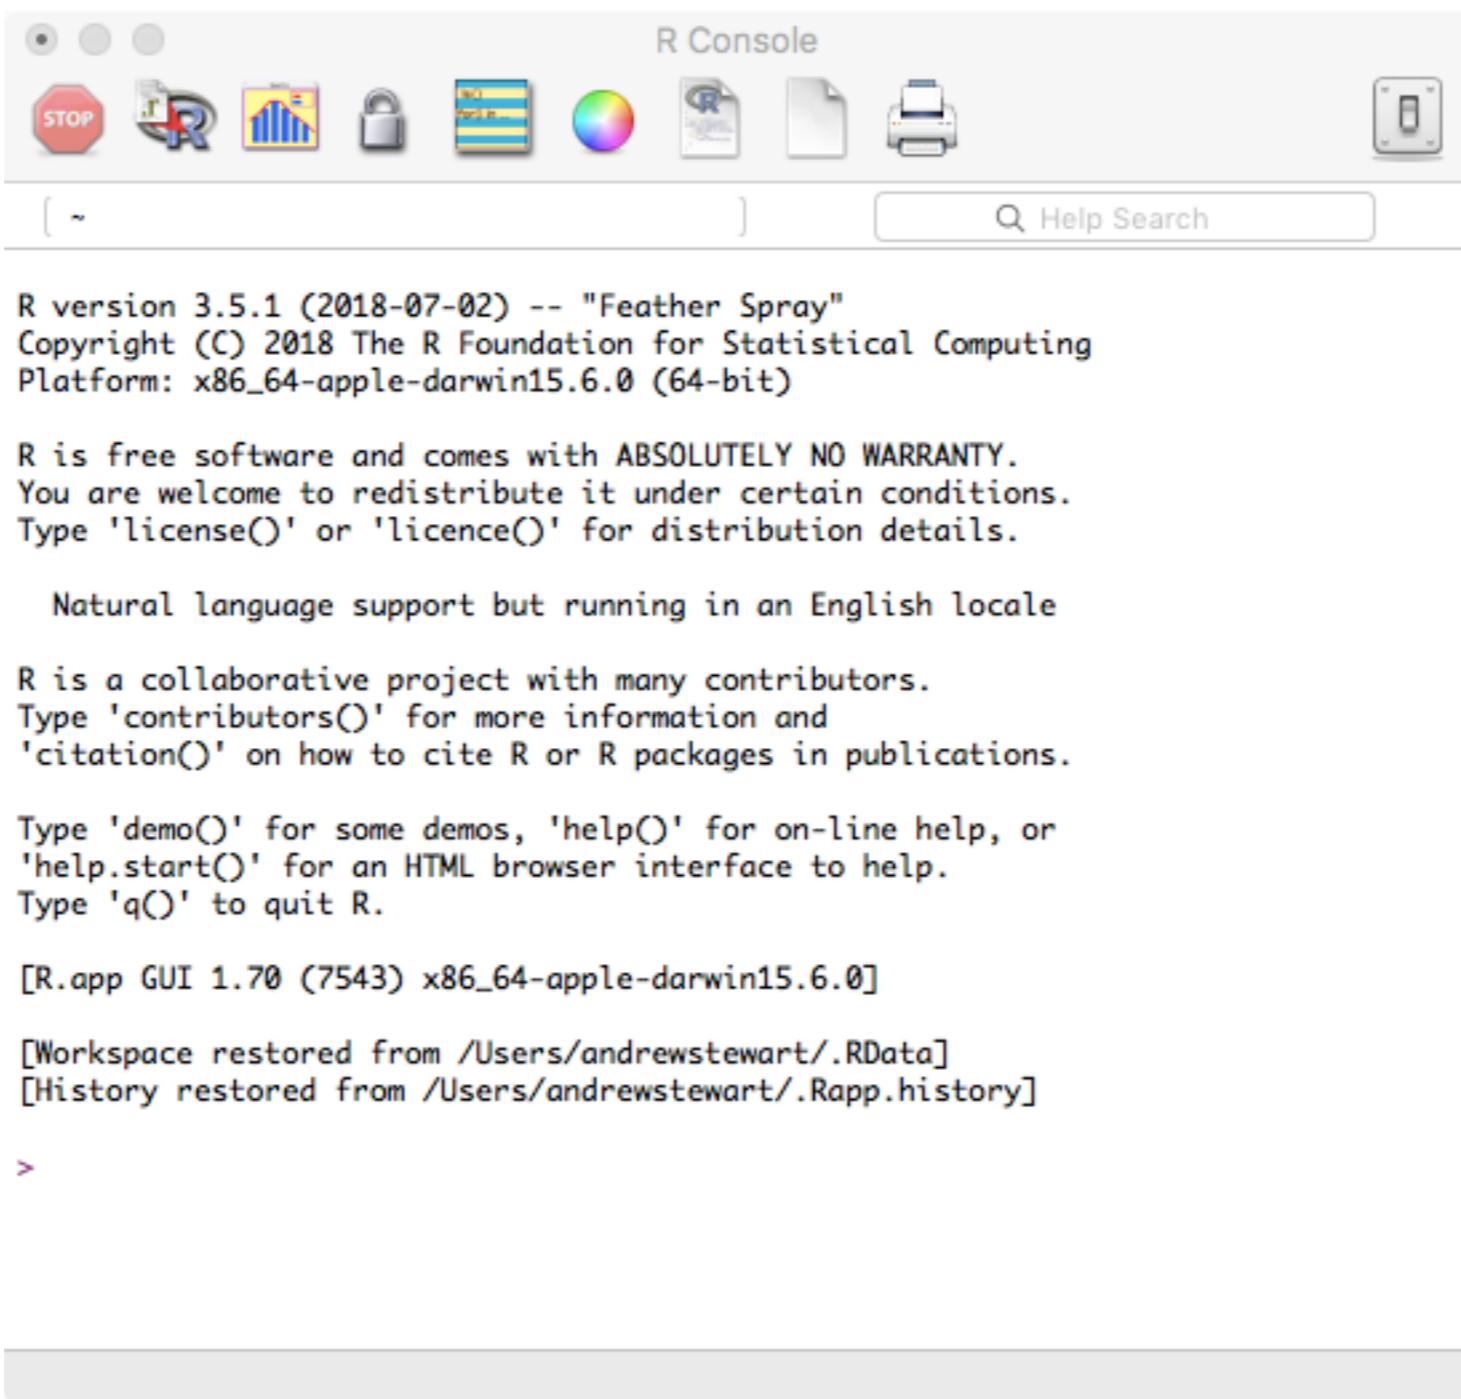
Discussion / Q & As about reproducibility

Starting R

- R is open source (i.e., free to download and add to). Main R site is:
- *www.r-project.org*
- From here you can download R (from one of the CRAN¹ mirrors for Windows, Mac and UNIX).
- R updates regularly (you need to update manually).

¹CRAN = *The Comprehensive R Archive Network*

- If you load R directly it looks something like this:



R version 3.5.1 (2018-07-02) -- "Feather Spray"
Copyright (C) 2018 The R Foundation for Statistical Computing
Platform: x86_64-apple-darwin15.6.0 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[R.app GUI 1.70 (7543) x86_64-apple-darwin15.6.0]

[Workspace restored from /Users/andrewstewart/.RData]
[History restored from /Users/andrewstewart/.Rapp.history]

>



- Rather than using the R interface, you should use the RStudio graphical interface.
- Download it from www.rstudio.com
- When you use RStudio, it looks like this:

RStudio File Edit Code View Plots Session Build Debug Tools Window Help

Thu 15:21

~/Desktop/Air Work/MRes 2016:17/R workshop MRes/R workshop directory/Workshop - RStudio

Addins

Workshop script.R

```

1 library(lme4)
2 library(lsmeans)
3 library(pbkrtest)
4
5 #Read in First Pass Data
6 FPs <- read.csv("~/Desktop/Air Work/R analyses/Indirect Request Expt/Experiment 1 - probability of success - Libby's data/FP")
7
8 #This sets up the contrasts so that the intercept in the mixed LMM is the grand mean (i.e., the mean of all conditions)
9 my.coding <- matrix(c(.5, -.5))
10
11 contrasts(DVSContext)<-matrix(c(.5, -.5))
12 contrasts(DVSSentence)<-matrix(c(.5, -.5))
13
14 #construct the models with crossed random effects for subjects and items for the pre-critical, critical and post-critical regt
15
16 model.full <- lmer(RT ~ Context*Sentence + (1|Context*Sentence |Subject) + (1|Context*Sentence |Item), data=DV, REML=TRUE)
17 model.null <- lmer(RT ~ (1+Context*Sentence |Subject) + (1+Context*Sentence |Item), data=DV, REML=TRUE)
18
19 summary(model.full)
20 lsmeans(model.full, pairwise=Context*Sentence, adjust="none")
21
22 contrasts <- lmer(CatStatement ~ MixedEffectsList + (1|MixedEffectsList |Subject) + (1|MixedEffectsList |Item), data=R)
23
24
```

The following object is masked from 'package:stats':

```

step
> library("lsmeans", lib.loc="/Library/Frameworks/R.framework/Versions/3.3/Resources/library")
Loading required package: estimability
Attaching package: 'lsmeans'

The following object is masked from 'package:lmerTest':
lsmeans
> model.full <- lmer(RT ~ Context*Sentence + (1|Context*Sentence |Subject) + (1|Context*Sentence |Item), data=DV, REML=TRUE)
```

Error in strsplit(keys, " ", "") : non-character argument

Environment History

DV 1680 obs. of 5 variables

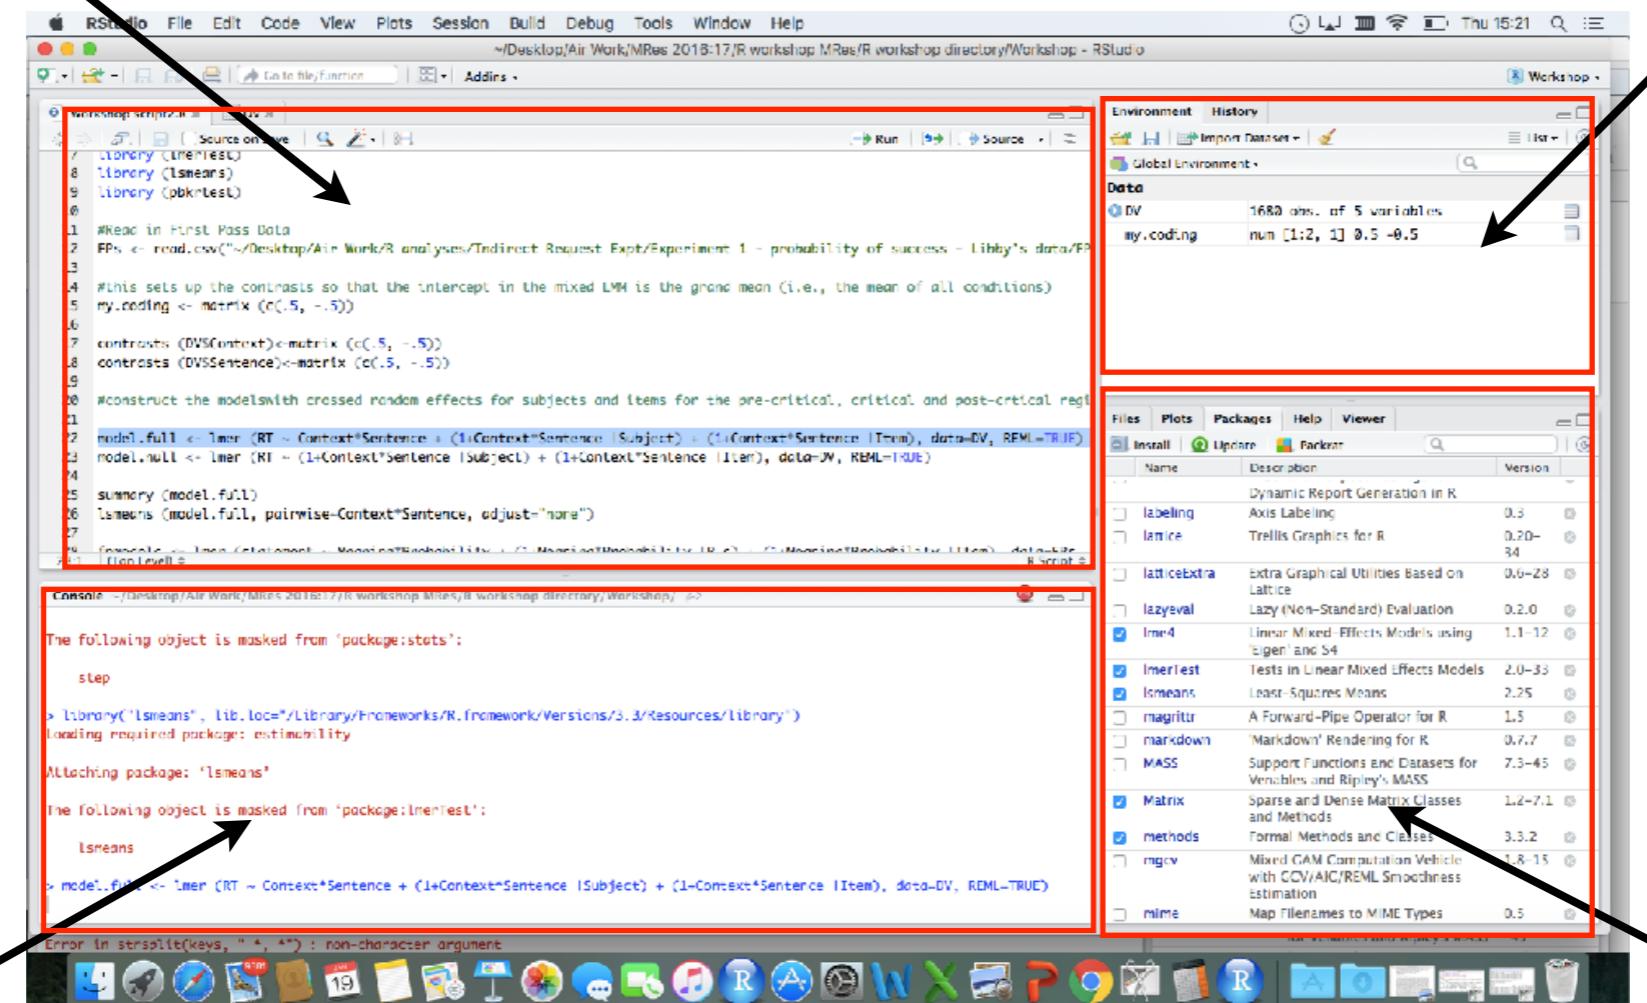
Global Environment

Files Plots Packages Help Viewer

Install Update Backtrack

Name	Description	Version
labeling	Dynamic Report Generation in R	0.3
lattice	Axis Labeling	0.20-
latticeExtra	Trellis Graphics for R	0.24
lazyeval	Extra Graphical Utilities Based on Lattice	0.6-28
lme4	Lazy (Non-Standard) Evaluation	0.2.0
lmerTest	Linear Mixed-Effects Models using 'Eigen' and 'S4'	1.1-12
lsmeans	Tests in Linear Mixed Effects Models	2.0-33
magrittr	Least-Squares Means	2.25
markdown	A Forward-Pipe Operator for R	1.5
MASS	'MarkDown' Rendering for R	0.7.7
Matrix	Support Functions and Datasets for Venables and Ripley's MASS	7.3-45
methods	Sparse and Dense Matrix Classes and Methods	1.2-7.1
mgcv	Formal Methods and Classes	3.3.2
mime	Mixed GAM Computation Vehicle with GCV/AIC/REML Smoothness Estimation	1.8-15
readr	Map Filenames to MIME Types	0.5

This is where you build your script and where data can be seen.

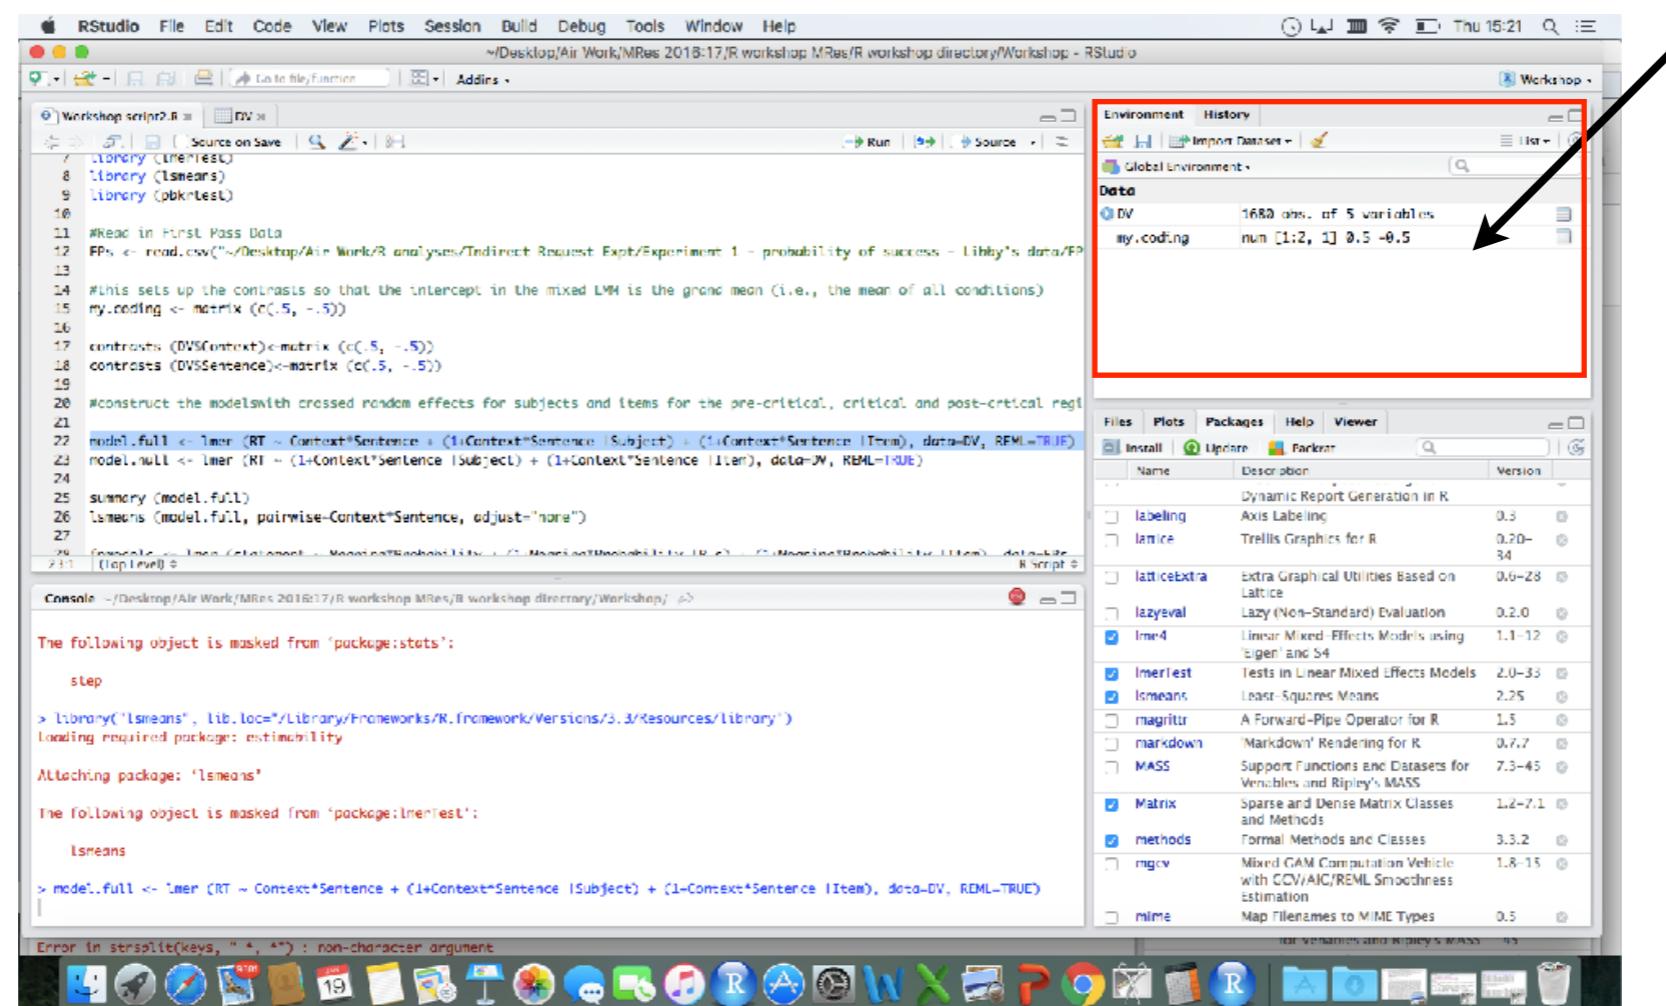


This is where you type commands.

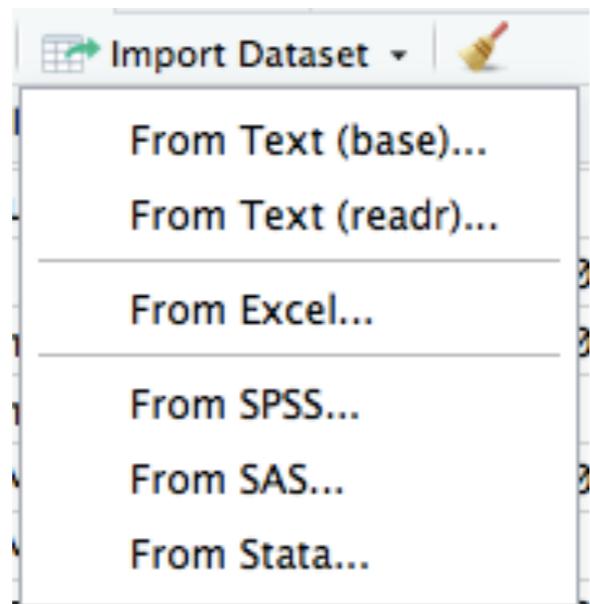
This contains information about your variables and open data sets.

This lists the packages you have loaded, and has tabs for help, graphs etc.

Here is where you import your data.



- Importing data (CSV (Text), Excel, SPSS, SAS and Stata format). CSV can be delimited by commas, tabs etc. Most of the time you'll use the `readr` import function...



Import Text Data

File/Url: ~/Desktop/Air Work/R analyses/R courses/R course/R Work Folder/priming data.txt

Data Preview:

Subject (integer)	Priming Condition (character)	RT (integer)
1	LONG	500
2	LONG	551
3	LONG	479
4	LONG	563
5	LONG	522
6	SHORT	399
7	SHORT	423
8	SHORT	444
9	SHORT	410
10	SHORT	398

You can change the type of each variable by clicking on the down arrow.

You should change this to type Factor (LONG vs. SHORT).

Previewing first 50 entries.

Import Options:

Name: <input type="text" value="priming_data"/>	<input checked="" type="checkbox"/> First Row as Names	Delimiter: <input type="button" value="Tab"/>	Escape: <input type="button" value="None"/>
Skip: <input type="text" value="0"/>	<input checked="" type="checkbox"/> Trim Spaces	Quotes: <input type="button" value="Default"/>	Comment: <input type="button" value="Default"/>
	<input checked="" type="checkbox"/> Open Data Viewer	Locale: <input type="button" value="Configure..."/>	NA: <input type="button" value="Default"/>

Code Preview:

```
library(readr)
priming_data <- read_delim("~/Desktop/Air Work/R analyses/R courses/R course/R Work Folder/priming data.txt",
  "\t", escape_double = FALSE, trim_ws = TRUE)
View(priming_data)
```

Can change here how the columns are delimited.

This is the code that corresponds to what you're getting RStudio to do.

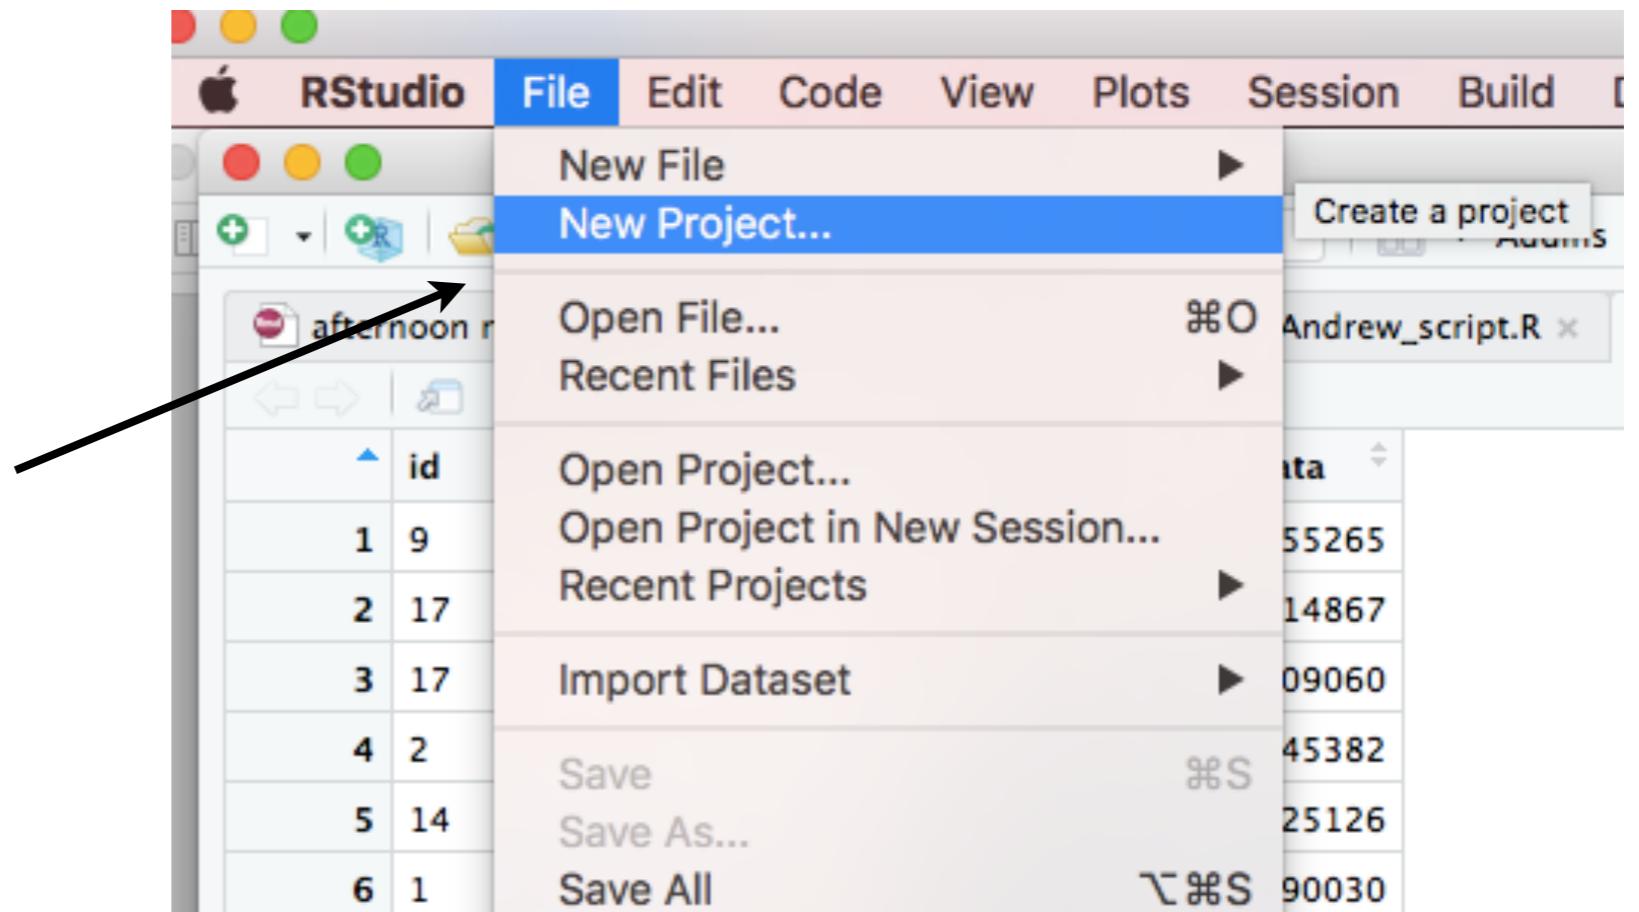
- RStudio now looks like this:

The screenshot shows the RStudio interface with the following components:

- Data View:** Displays the "priming_data" data frame with 10 rows and 3 columns: Subject, Priming Condition, and RT.
- Environment View:** Shows the Global Environment with "priming_data" listed.
- Console View:** Displays R code and its output. The code reads a CSV file into "priming_data" and then views it. It also attempts to subset by "Priming Condition" and "RT" but fails because "priming_data" is not found.
- Packages View:** Shows a list of installed packages, including labeling, lattice, latticeExtra, lazyeval, lme4, lmerTest, lsmeans, magrittr, markdown, MASS, Matrix, methods, mgcv, mime, minqa, multcomp, and munsell.

When Starting R

Always create a new project - this will keep all your files nice and organised.

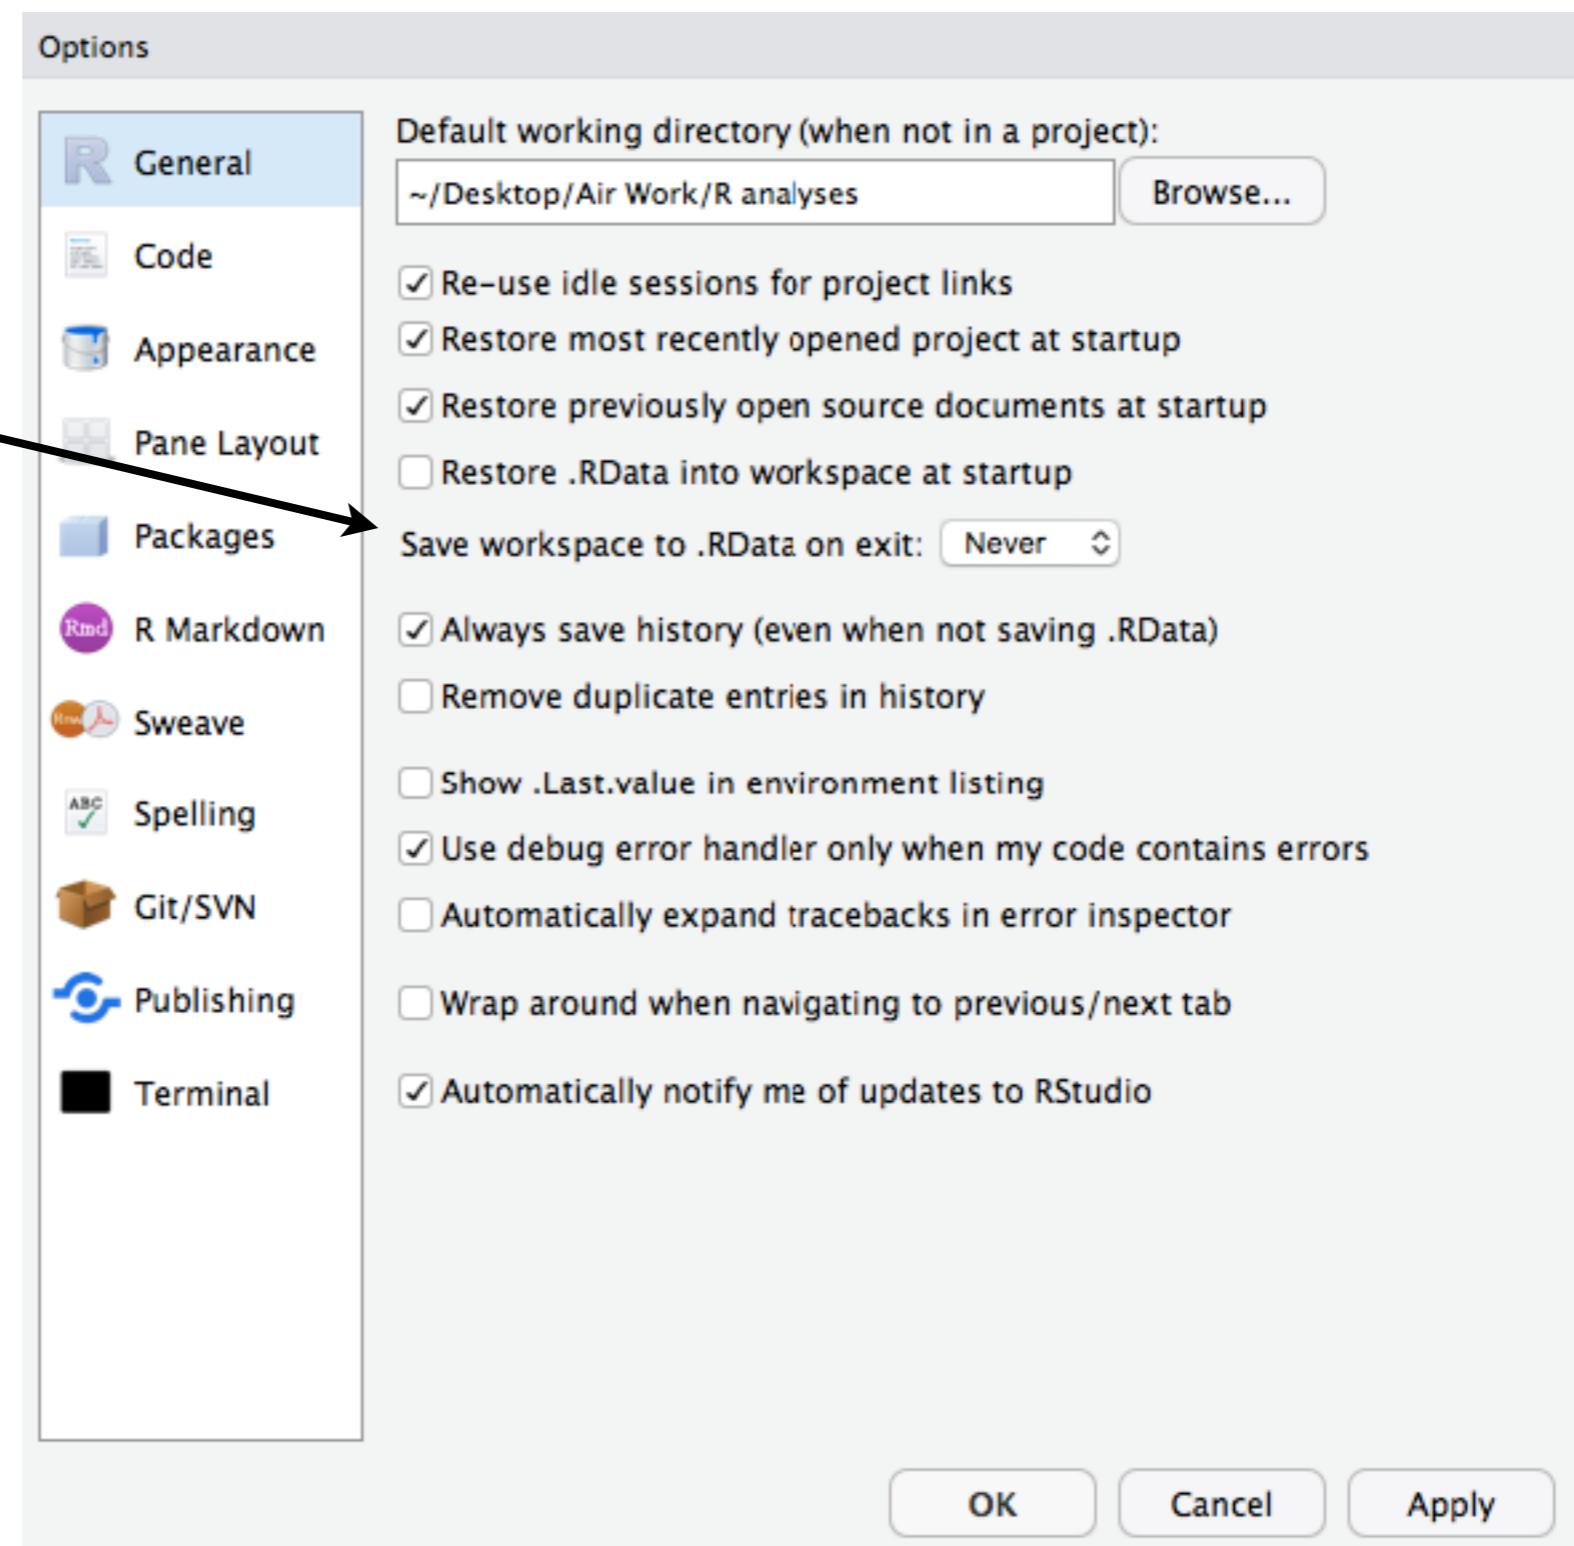


When you create an R Project in a folder, it's a good idea to keep that folder itself nicely organised:

```
name_of_project
| --name_of_project.Rproj
| --raw_data
|   |--data.csv
| --tidied_data
|   |--tidy_data.csv
| --R_scripts
|   |--data_tidy.R
|   |--data_vis.R
|   |--analysis.R
| --markdown_scripts
|   |--analysis.Rmd
| --output_reports
|   |--analysis.html
|   |--analysis.pdf
```

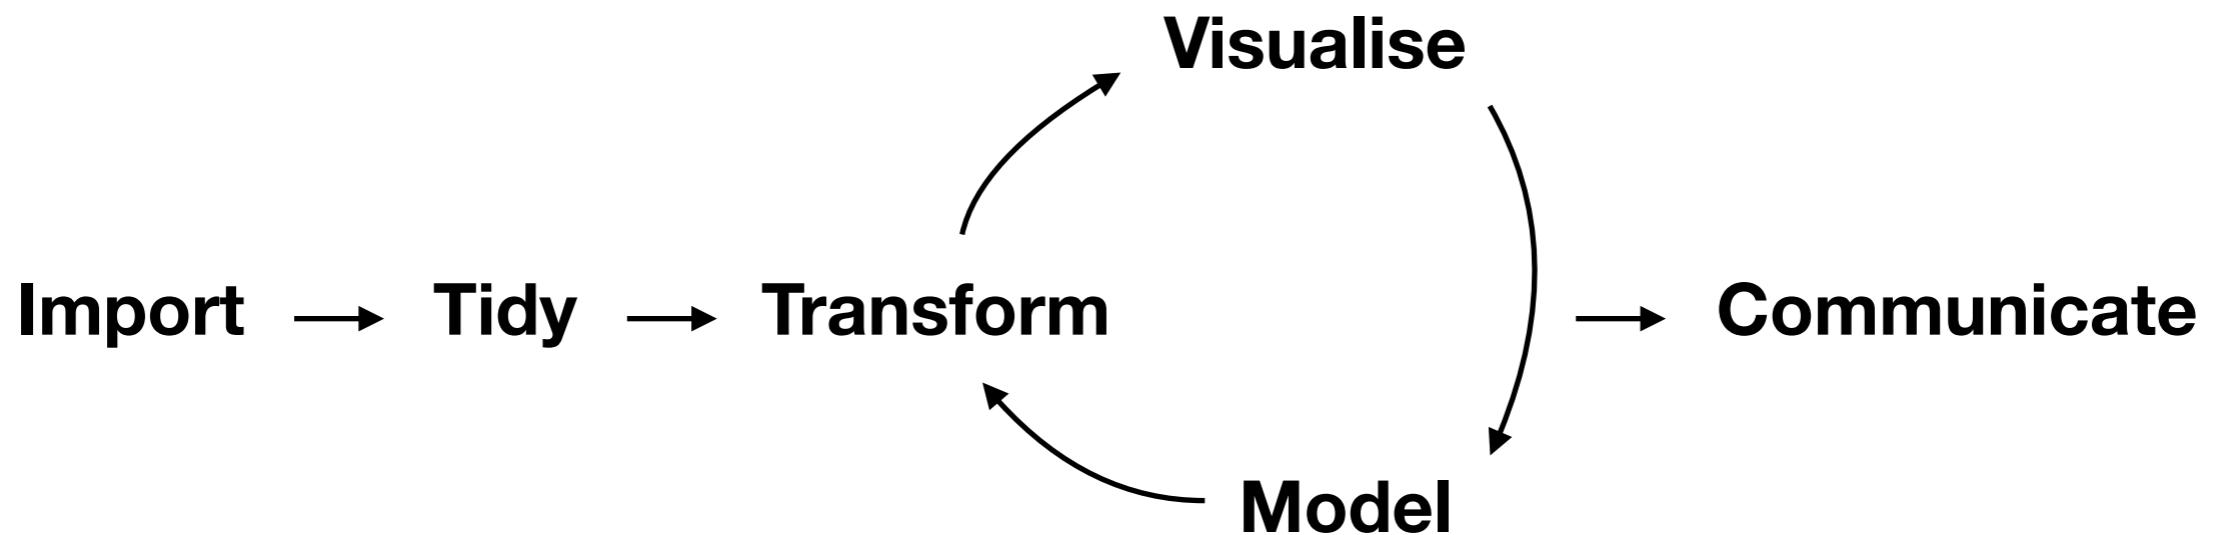
When Starting R

Under preferences,
make sure you
uncheck the saving
workspace
option...

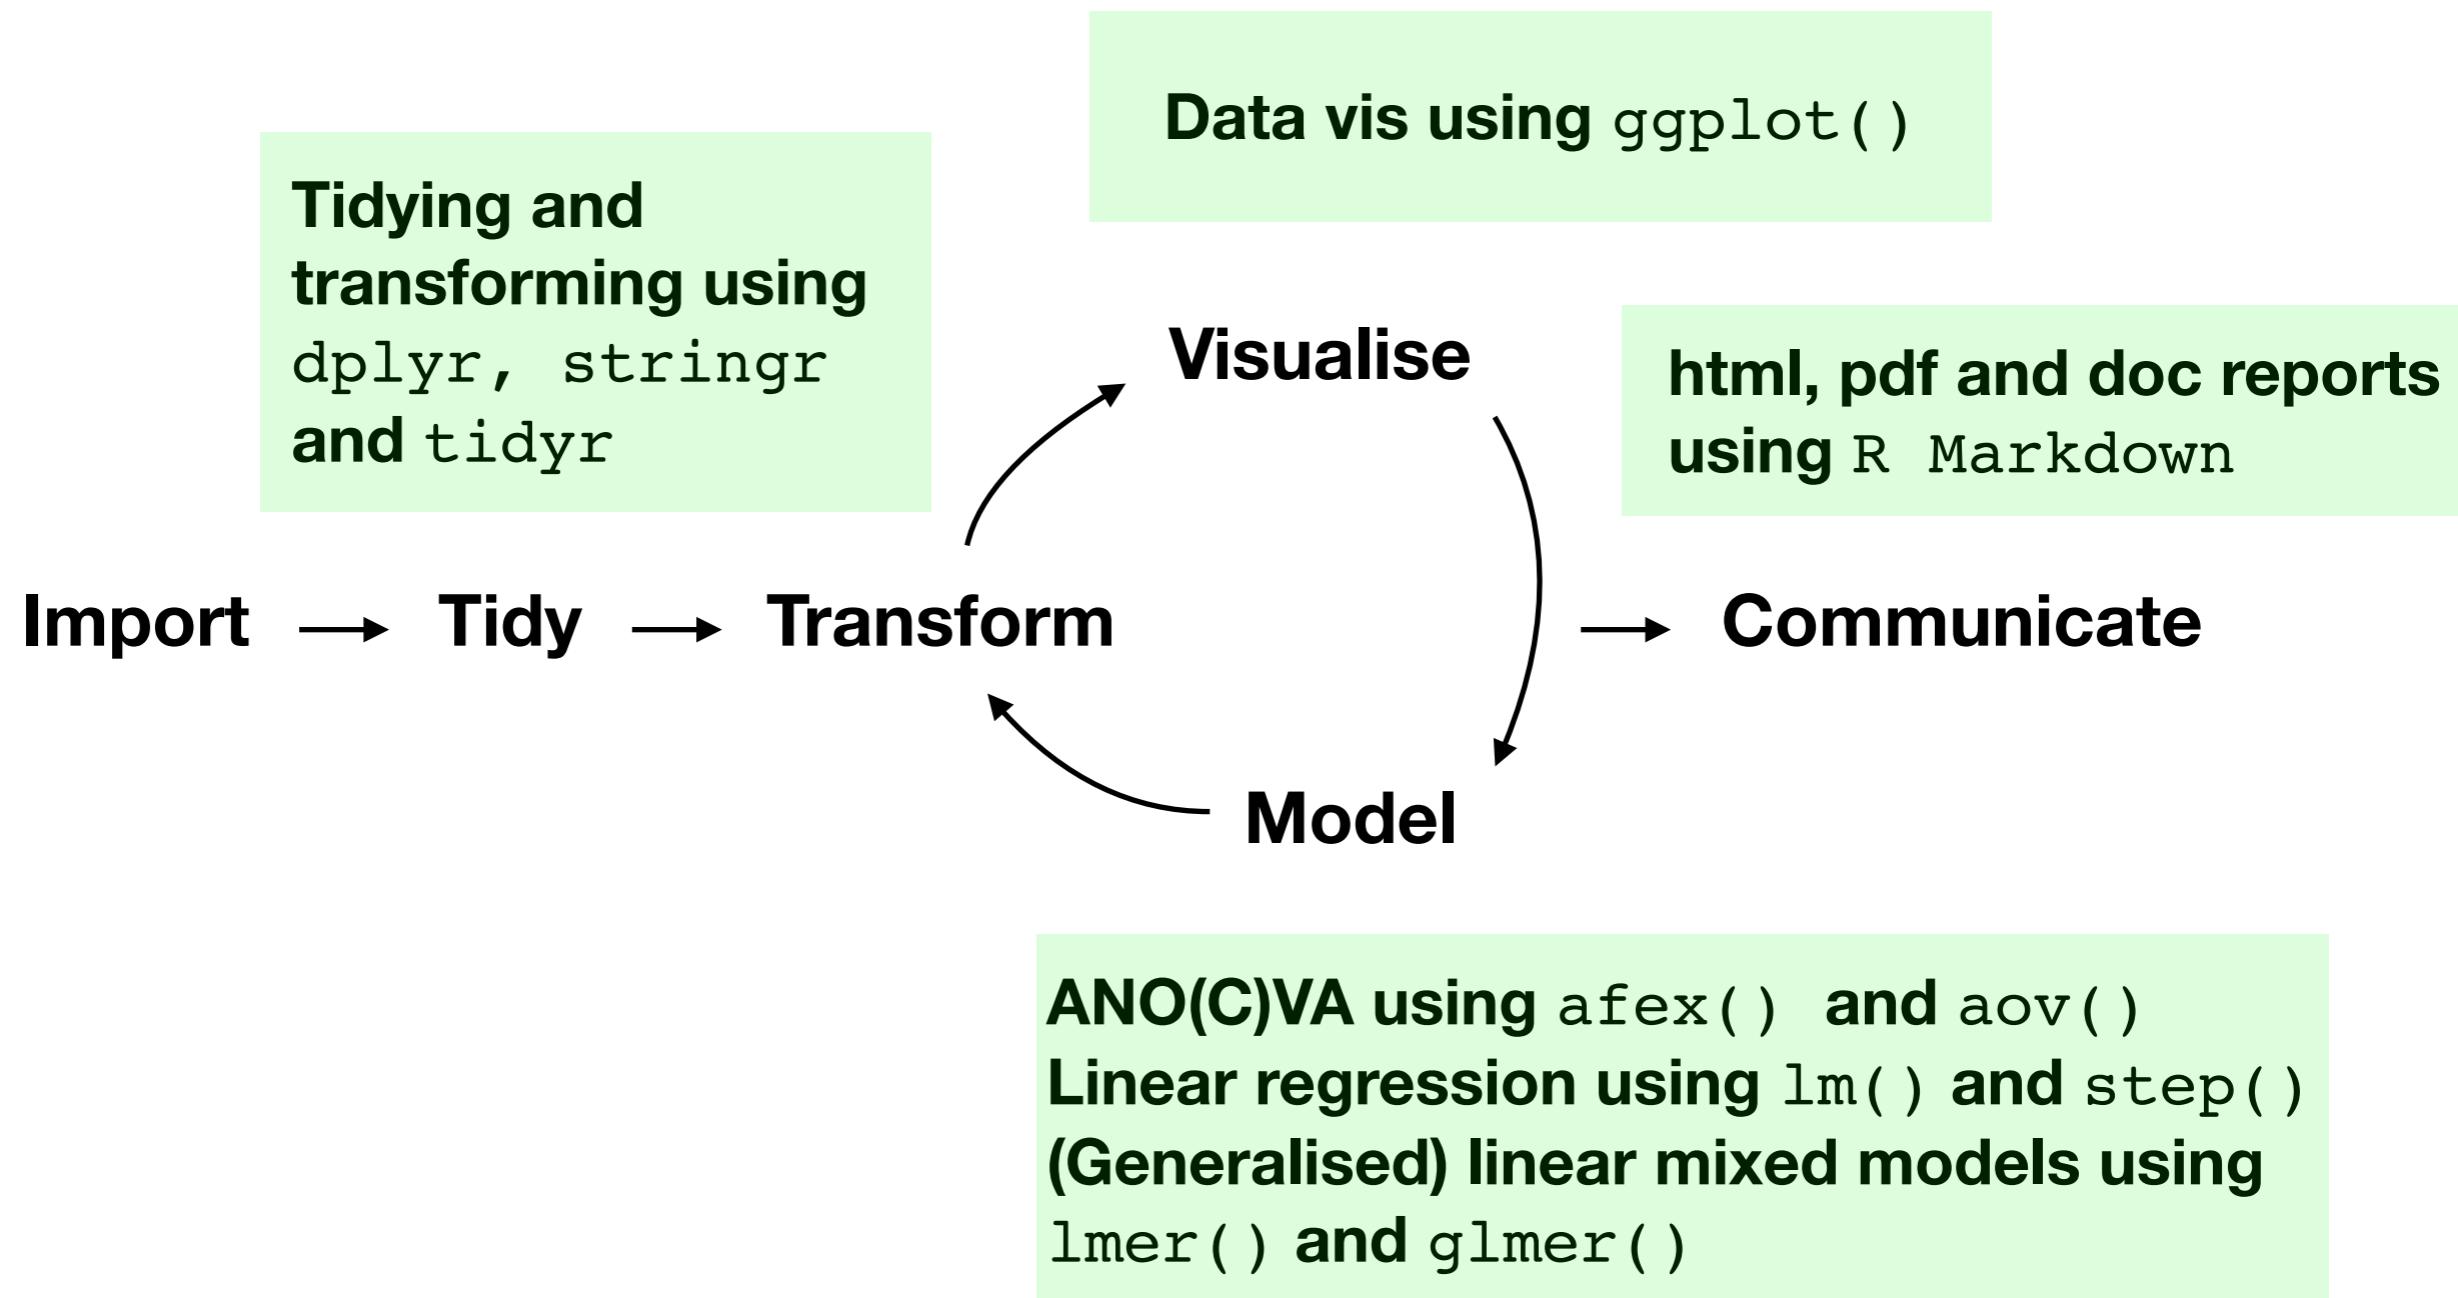


Analysis Workflow in the Tidyverse

(Garrett Grolemund and Hadley Wickham) - from Data to Write-up



A Reproducible Workflow



Packages in R

R has a number of functions already built in. These are part of “base R”. For much of what we do we need to load particular packages to let us use additional functions. We do this by:

```
> install.packages ("packagename")  
  
> library(packagename)
```

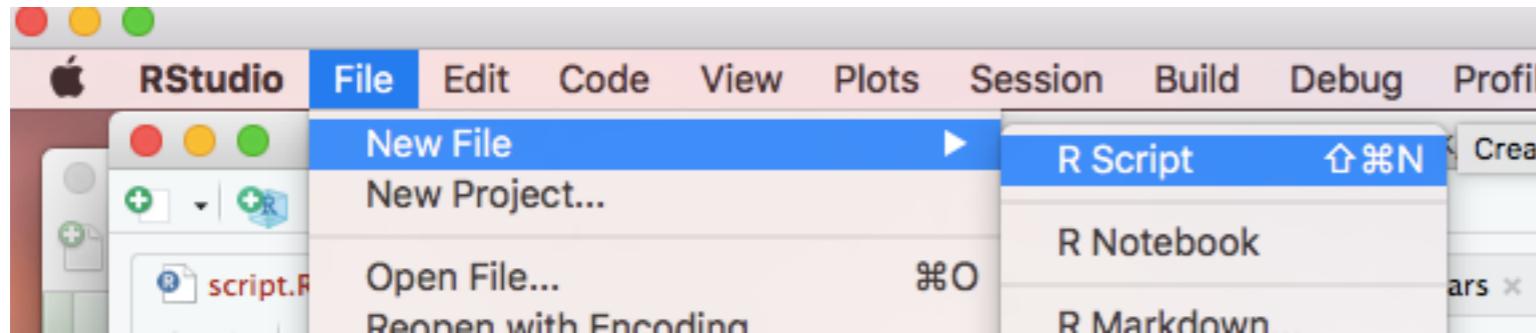
Every time you start R afresh, you need to load the packages you will use by using the `library` function. If you see notation like the following, it means we’re using a function (e.g., `summarise`) from within a package (e.g., `dplyr`)

```
dplyr::summarise
```

You don't need to re-invent the wheel...

- For any problem you want to solve, chances are others have had the same problem, solved it and have created an R package to do exactly what you want...
- One of the many great things about R is that you can add freely available packages to your library.
- ~15,000 R packages currently available
- The sites r-bloggers.com and rweekly.org are good places to find out about new packages.

Writing Scripts in R



Ultimately you will be writing scripts that will allow yourself and others to recreate your analysis just by running the script - the code at the start of the script will load the packages your analysis needs, then load your data, tidy, transform and visualise your data, and then code for your model...

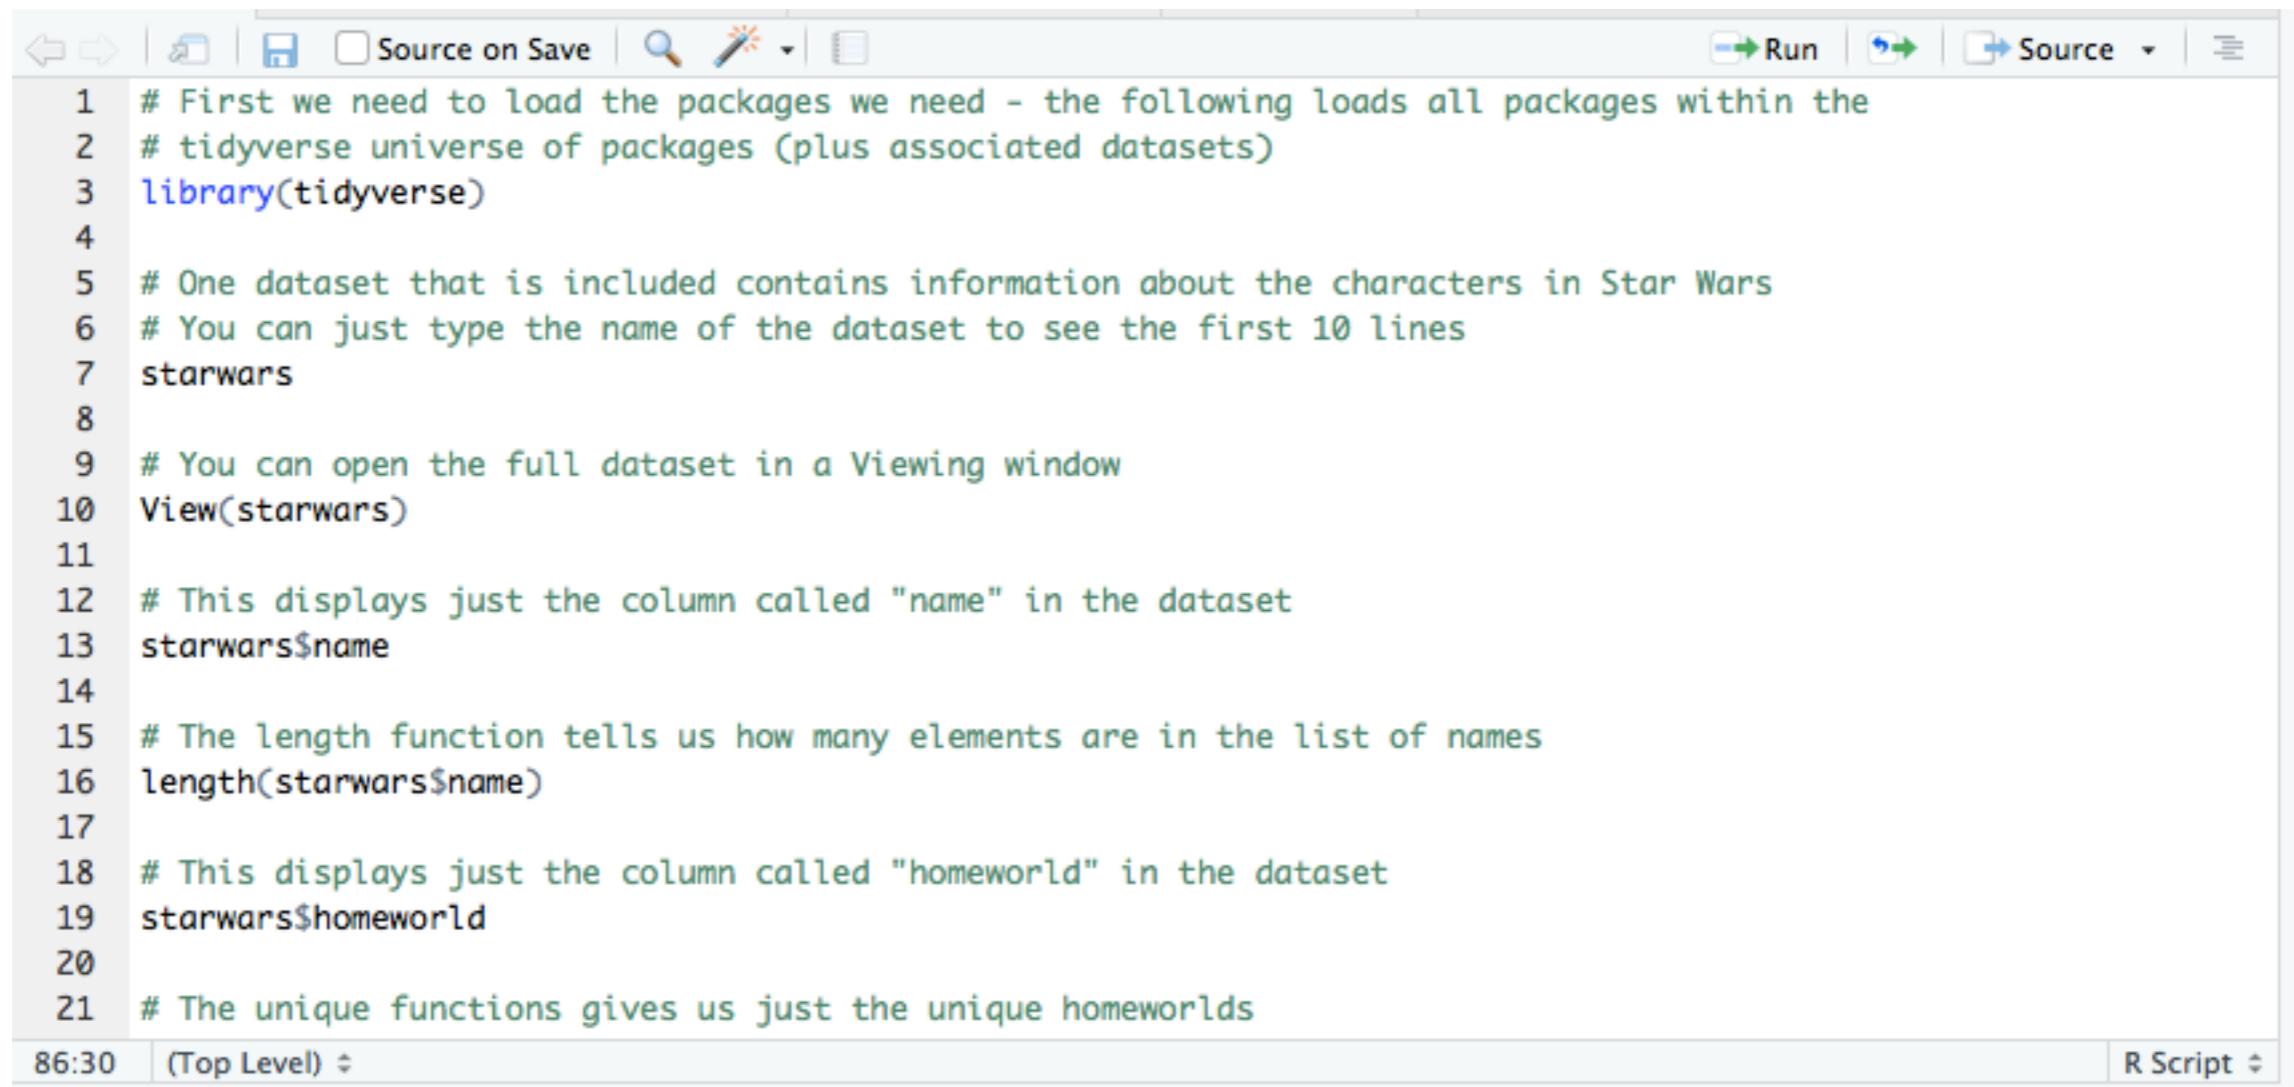
You should write your code as a script, and use the Console window for single command instructions relatively rarely...

Writing Scripts in R

Scripts should have lots of comments, indicated by a # symbol - this helps others read your code, and also yourself when you look back at it later.

Scripts can be saved and uploaded to (e.g.) OSF, GitHub, or submitted as an electronic supplement alongside your journal submission. Even when journals don't require you to adhere to Open Science practices, it's a good idea to make your code and data available during the reviewing process - and publicly available once your paper is published.

If you aren't sharing your data and code, you aren't engaged in generating reproducible or open research...

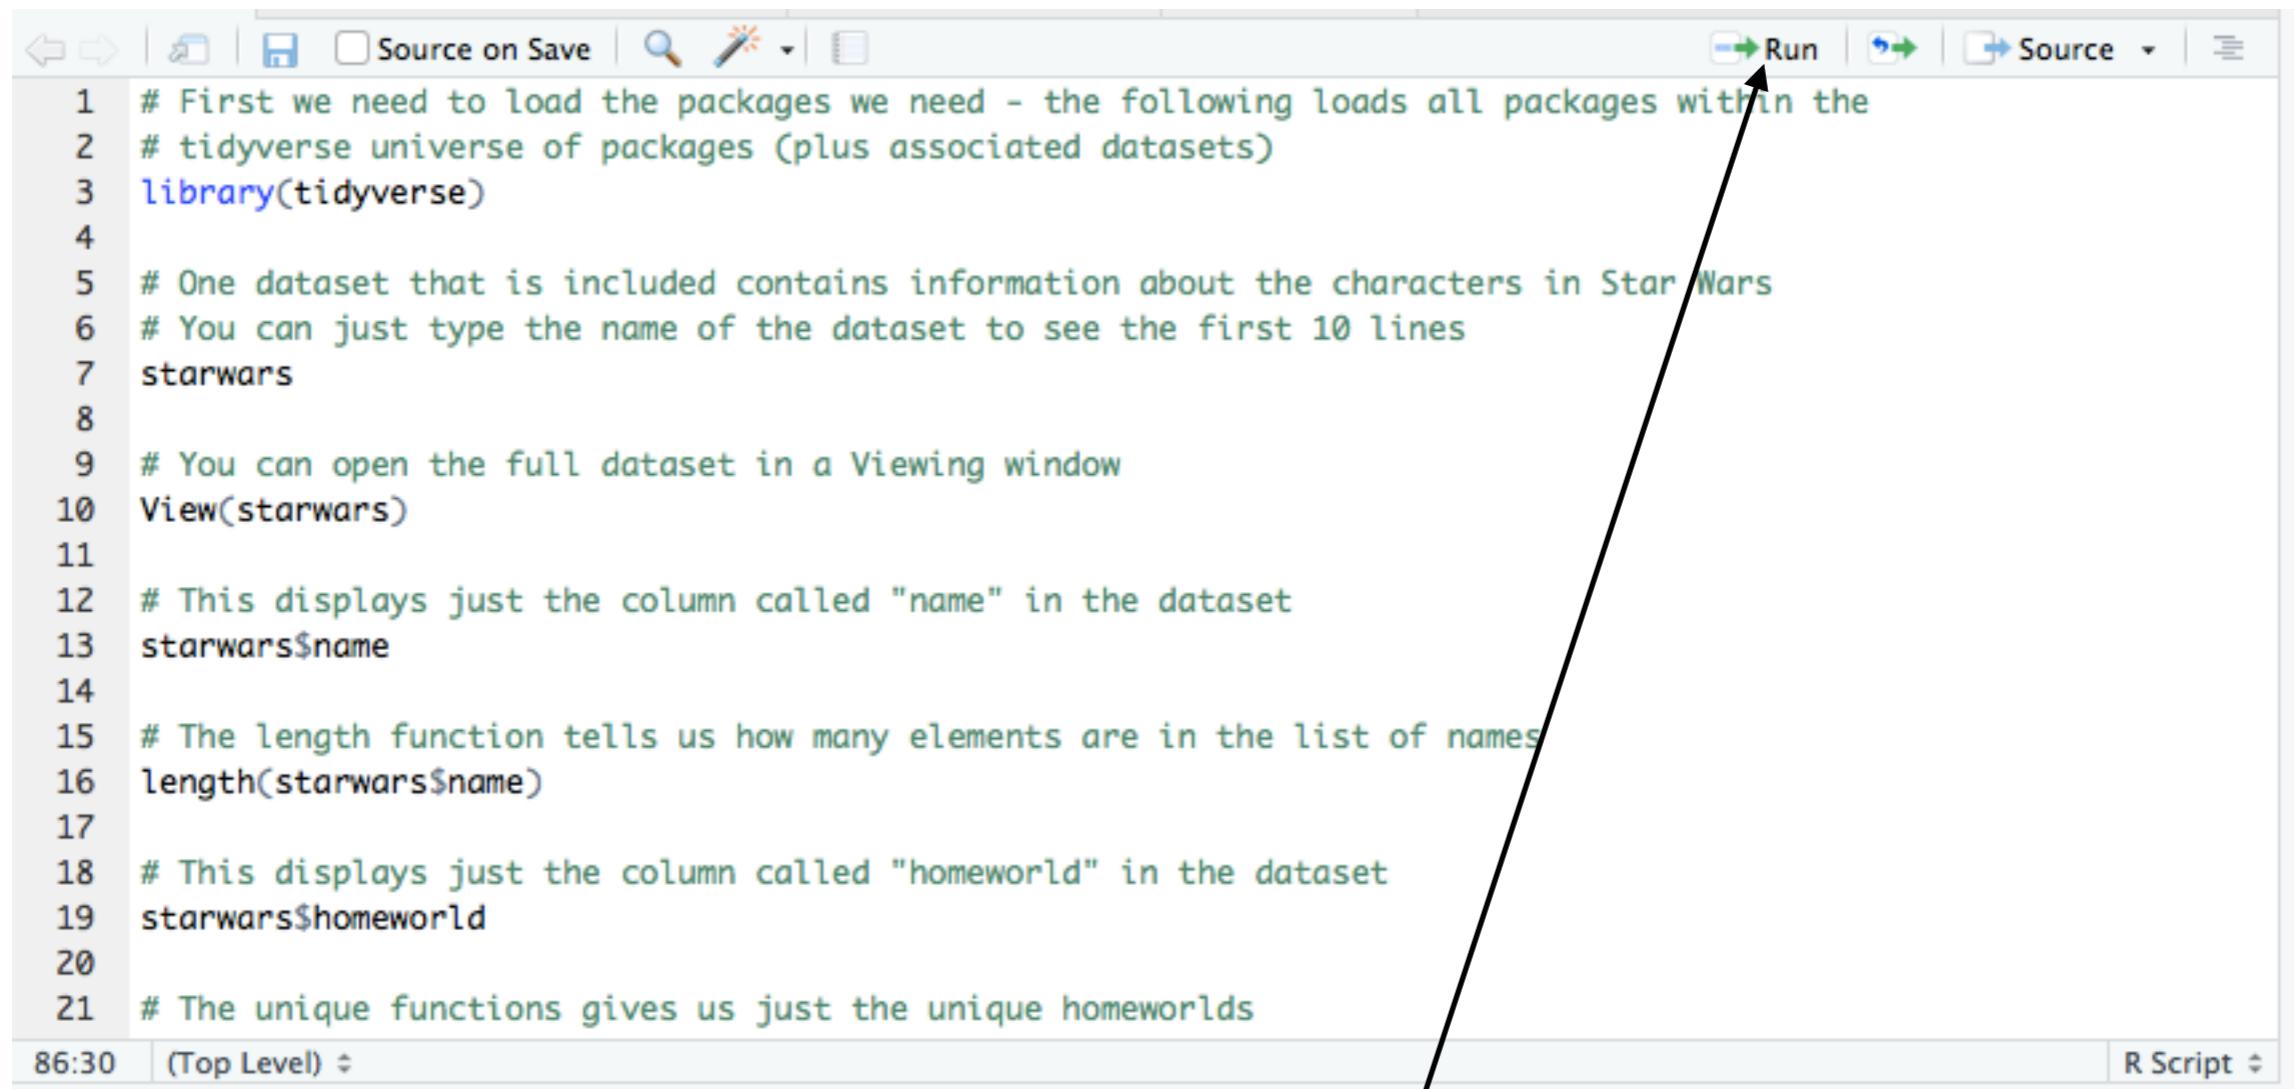


The screenshot shows the RStudio interface with an R script editor. The script contains numbered comments explaining various R commands:

```
1 # First we need to load the packages we need - the following loads all packages within the
2 # tidyverse universe of packages (plus associated datasets)
3 library(tidyverse)
4
5 # One dataset that is included contains information about the characters in Star Wars
6 # You can just type the name of the dataset to see the first 10 lines
7 starwars
8
9 # You can open the full dataset in a Viewing window
10 View(starwars)
11
12 # This displays just the column called "name" in the dataset
13 starwars$name
14
15 # The length function tells us how many elements are in the list of names
16 length(starwars$name)
17
18 # This displays just the column called "homeworld" in the dataset
19 starwars$homeworld
20
21 # The unique functions gives us just the unique homeworlds
```

The status bar at the bottom shows the time as 86:30, the project level as (Top Level), and the file type as R Script.

Lots of comments explaining what each section of code does, and lots of white space.



```
1 # First we need to load the packages we need - the following loads all packages within the
2 # tidyverse universe of packages (plus associated datasets)
3 library(tidyverse)
4
5 # One dataset that is included contains information about the characters in Star Wars
6 # You can just type the name of the dataset to see the first 10 lines
7 starwars
8
9 # You can open the full dataset in a Viewing window
10 View(starwars)
11
12 # This displays just the column called "name" in the dataset
13 starwars$name
14
15 # The length function tells us how many elements are in the list of names
16 length(starwars$name)
17
18 # This displays just the column called "homeworld" in the dataset
19 starwars$homeworld
20
21 # The unique functions gives us just the unique homeworlds
```

Once a script is written, you can run the entire script by highlighting it all (CMD-A in OSX) and clicking on ‘Run’, or you can highlight a section and run only that. CMD-Return will also Run the highlighted code.

Like this...

Style Guide

File names (both scripts and data files) should be meaningful:

`regression_model.R` - Good

`somemodel.R` - Bad

Variable names should be meaningful - consider using `snake_case` or `CamelCase`, but never use whitespace in variable names:

`age` - Good

`A1` - Bad

`expt1_data` - Good

`Expiriment1datawithoutliersremoved` - Bad

Style Guide

Place spaces around operators like + , - , = , <-

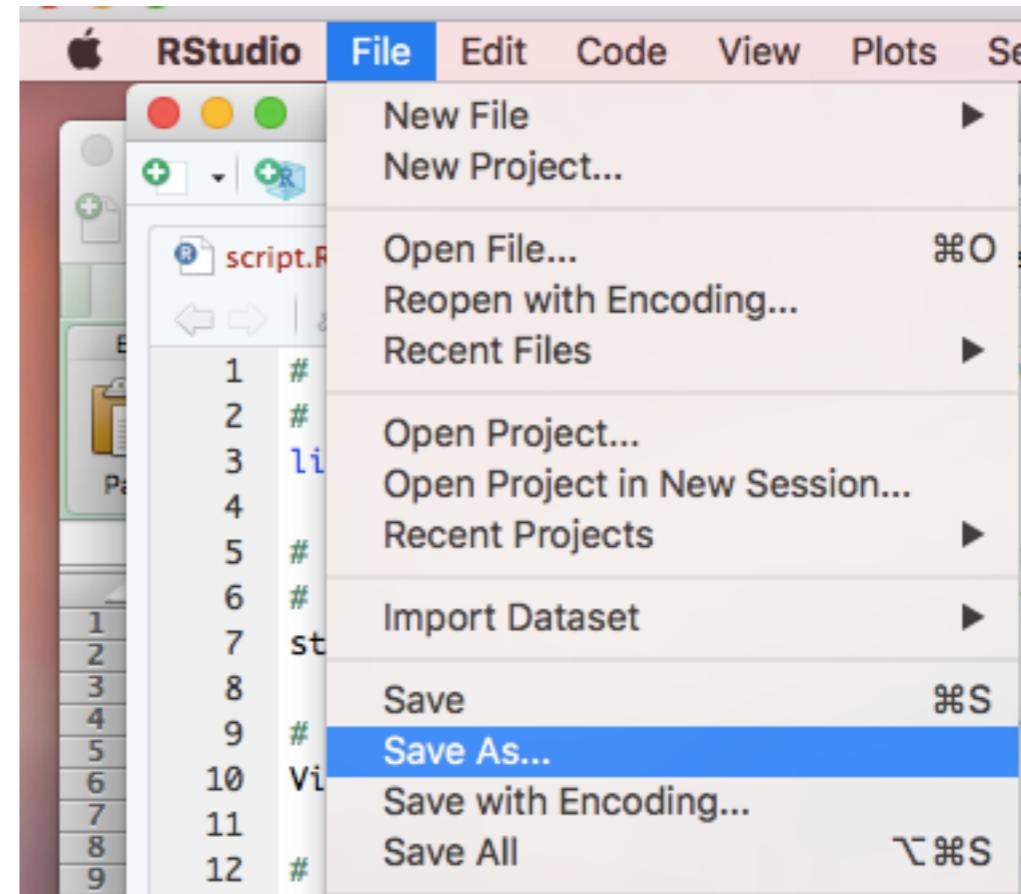
Whitespace used sensibly in your script makes your code easier to read. Separate discrete sections of code in your script with a blank line.

When writing scripts, continue on a subsequent line rather than writing one very long line.

Comment, comment, comment...

```
# First we load our datafile.
```

Once your
script is
finished, don't
forget to save
it...



Sharing your analysis

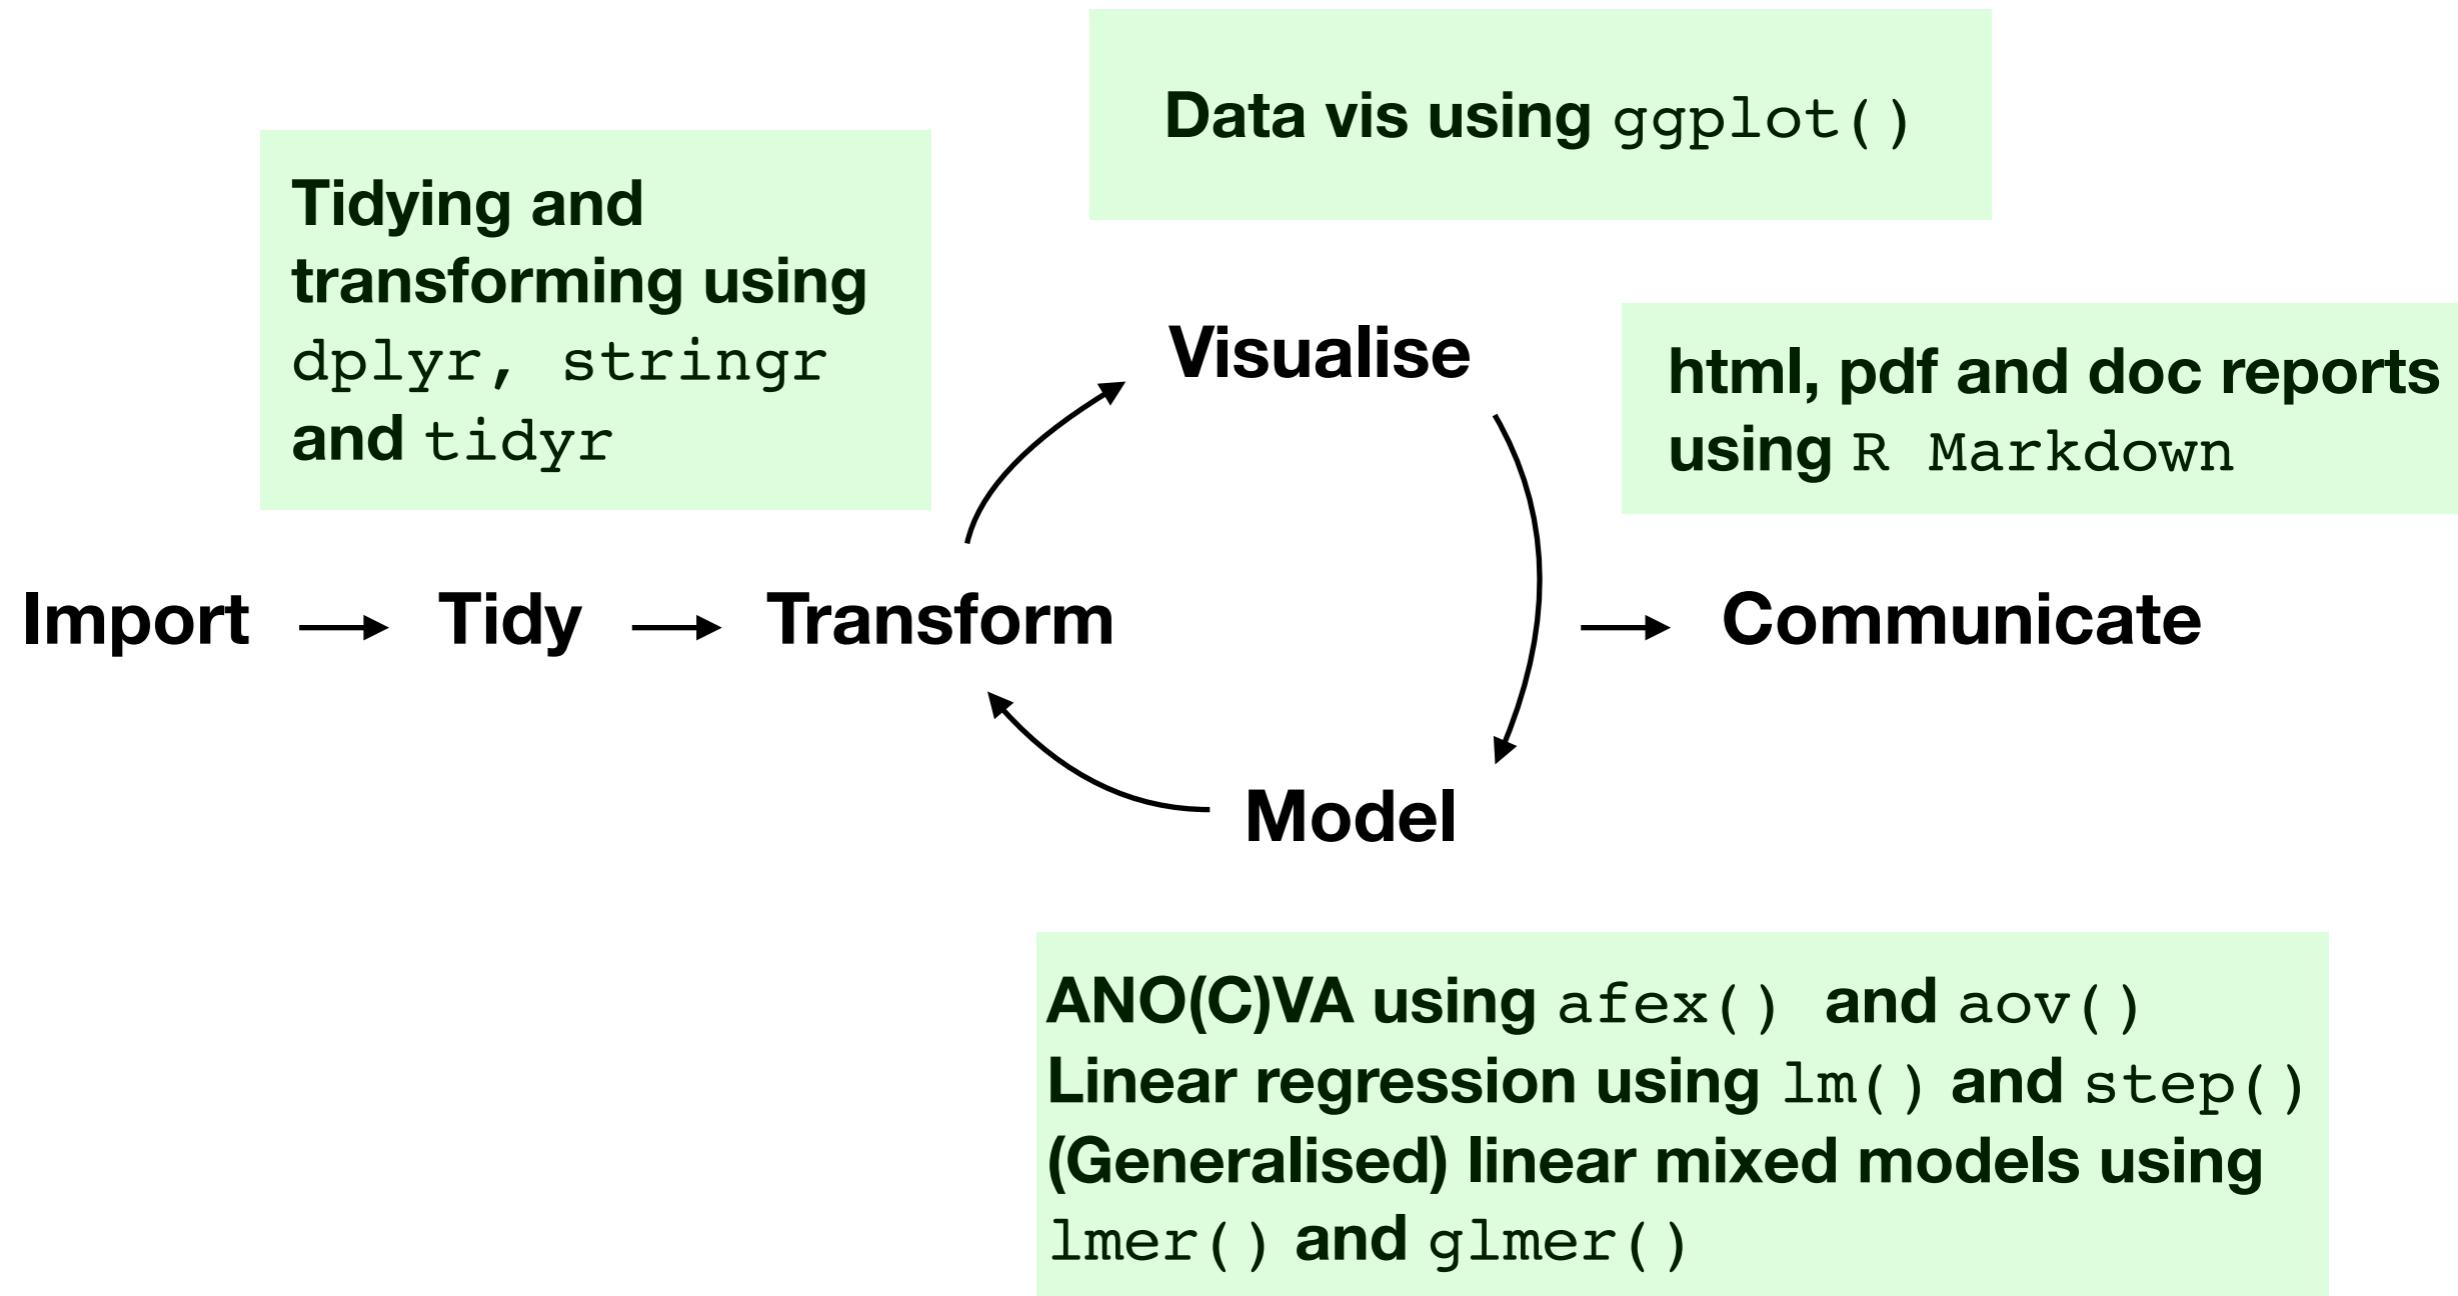
Journals are increasingly requesting analysis code and data to be published alongside the published paper (and sometimes at the review stage).

You can share your analysis and data even at the submission stage using something like GitHub, OSF or (even better) build a Binder.

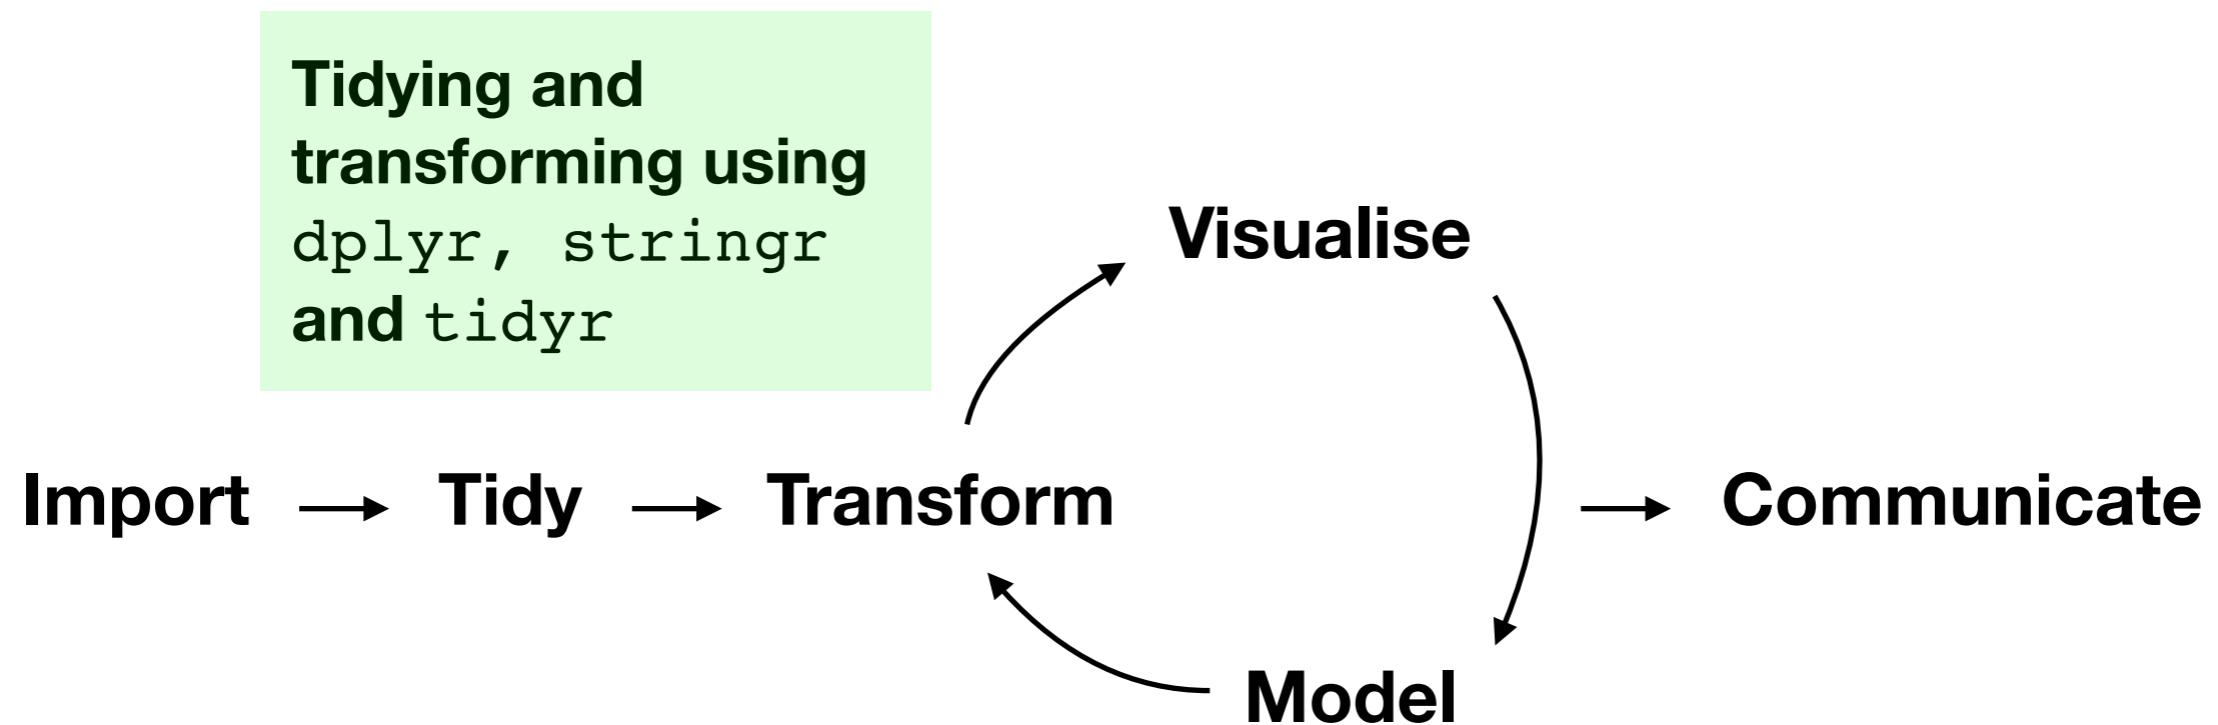
Git is a powerful tool that allows for the version control in collaborative projects, and the sharing of code and projects.

You can set up a GitHub account, and install GitHub Desktop on your own computer.

A Reproducible Workflow



A Reproducible Workflow



Let's get ready to code...

On your computer, fire up RStudio and install the tidyverse:

```
> install.packages("tidyverse")
```

Then create a new Project in a new folder, and fire up a new script...

On the first line of your script type:

```
library(tidyverse)
```

Then run the script...

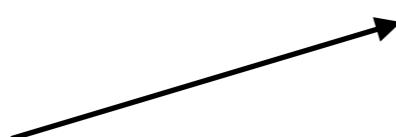
Let's now load two datasets that I've created:

```
data1 <- read_csv("https://bit.ly/2OTkuNz")  
dataRT <- read_csv("https://bit.ly/31xd6sG")
```

Tidying and Transforming Data

Imagine we have two rectangular datasets - one (called data1) contains a large number of records of individual participants with measures of Working Memory, IQ, and reading comprehension.

If we type into the Console:
> data1
the data frame is displayed like
this...



	id	wm	iq	comp
	<int>	<int>	<int>	<int>
1	1	43	72	16
2	2	51	109	18
3	3	55	107	18
4	4	38	102	20
5	5	52	121	17
6	6	52	92	16
7	7	47	68	21
8	8	47	97	23
9	9	47	93	22
10	10	45	101	17
# ... with 9,990 more rows				

To get more information about the structure of our data frame we can type:

```
> str(data1)
Classes 'tbl_df', 'tbl' and 'data.frame': 10000 obs. of 4 variables:
$ id   : int 1 2 3 4 5 6 7 8 9 10 ...
$ wm   : int 43 51 55 38 52 52 47 47 47 45 ...
$ iq   : int 72 109 107 102 121 92 68 97 93 101 ...
$ comp: int 16 18 18 20 17 16 21 23 22 17 ...
```

So we have 10,000 observations with 4 variables associated with each observation - all of them of type integer.

If you ever need help about a function (e.g. str), just type:

```
>?str
```

or

```
>help(str)
```

in the Console window.

Imagine that 48 of these 10,000 people also took part in a reading time experiment and we have their reading data (called `dataRT`) for Simple Sentence and Complex Sentence reading conditions:

```
> str(dataRT)
Classes 'tbl_df', 'tbl' and 'data.frame': 48 obs. of 3 variables:
 $ id           : int  6400 457 8291 4998 2579 9122 1138 5138 5244 3160 ...
 $ simple_sentence : int  1902 1797 2080 1856 1997 1868 2154 1933 1900 1929 ...
 $ complex_sentence: int  2341 2503 2731 2375 2177 2284 2441 2349 2371 2372 ...
```

We are maybe interested in analysing the data of these 48 people in the data frame called `dataRT` but covarying out the effect of IQ captured in our data frame called `data1`.

Problem - how can we combine these two data frames so that we end up with one data frame of 48 people, their reading times plus their individual difference measures?

Manually, in Excel we could open the two data frames as spreadsheets and cut and paste cases where the id number matches...

Probably ok for 48 participants, but what if you had 200 or 2,000?

In R, we can use the `inner_join` function from the `dplyr` package where we join the two data frames matched by ID.

```
> dataRT_all <- inner_join(data1, dataRT, by = (c("id")) )  
> dataRT_all  
    id   wm   iq  comp simple_sentence complex_sentence  
1   95   47   94    19                 2154                  2441  
2  400   45  118    18                 1824                  2456  
3  457   42  100    22                 1857                  2324  
4 1138   41   77    18                 1902                  2341  
5 1587   54   67    21                 1844                  2320  
6 1805   52  109    19                 2224                  2256  
7 1864   57  111    19                 1880                  2391  
8 2006   44  110    19                 2091                  2456  
9 2183   55  125    23                 1926                  2218  
10 2318  51   91    21                 1960                  2440
```

We can use the assignment symbol `<-` to assign the output of this `inner_join` function to a new variable I'm calling `dataRT_all`. We can ask for the structure of this new data frame using the `str()` function:

```
> dataRT_all <- inner_join(data1, dataRT, by = (c("id")))
> str(dataRT_all)
Classes 'tbl_df', 'tbl' and 'data.frame': 48 obs. of 6 variables:
 $ id           : int  95 400 457 1138 1587 1805 1864 2006 2183 2318 ...
 $ wm           : int  47 45 42 41 54 52 57 44 55 51 ...
 $ iq            : int  94 118 100 77 67 109 111 110 125 91 ...
 $ comp          : int  19 18 22 18 21 19 19 19 23 21 ...
 $ simple_sentence: int  2154 1824 1857 1902 1844 2224 1880 2091 1926 1960 ...
 $ complex_sentence: int  2441 2456 2324 2341 2320 2256 2391 2456 2218 2440 ...
```

So we have created a new data frame of 48 participants consisting of their reading times and their individual difference measures from two separate (and different sized) data frames...with one line of code...

Now imagine we find the distributions of reading times for our two conditions are positively skewed (and we discover the residuals are non-normal). We could log transform these two columns and have two new columns in our data frame - let's call them `log_simple` and `log_complex`. We can use the `mutate` function in the `dplyr` package to create two new columns.

```
> data_transformed <- mutate(dataRT_all, log_simple = log(simple_sentence) ,  
+ log_complex = log(complex_sentence))  
> data_transformed  
# A tibble: 48 x 8  
  id      wm      iq      comp simple_sentence complex_sentence log_simple log_complex  
  <int> <int> <int> <int>          <int>          <int>       <dbl>       <dbl>  
1 95     47     94     19          1960         2440    7.58     7.80  
2 400    45    118     18          2186         2200    7.69     7.70  
3 457    42    100     22          1797         2503    7.49     7.83  
4 1138   41     77     18          2154         2441    7.68     7.80  
5 1587   54     67     21          1936         2395    7.57     7.78  
6 1805   52    109     19          1864         2560    7.53     7.85  
7 1864   57    111     19          1930         2540    7.57     7.84  
8 2006   44    110     19          2230         2267    7.71     7.73  
9 2183   55    125     23          1857         2324    7.53     7.75  
10 2318   51     91     21          1918         2739    7.56     7.92  
# ... with 38 more rows
```

Perhaps we have a reason to exclude a particular participant - number 2006 for example. We can use the filter function in `dplyr` to keep those participants where the ID number does not equal 2006.

```
filtered_data <- filter(data_transformed, id != 2006)
```

`!=` stands for “not equal to”- here are other useful logical operators in R:

`<` less than

`<=` less than or equal to

`>` greater than

`>=` greater than or equal to

`==` exactly equal to

`!=` not equal to

We can now apply our logical vector to our dataRT_all data frame and create a new filtered data frame (which I am calling filtered_data):

```
> filtered_data <- filter(data_transformed, id != 2006)
> filtered_data
# A tibble: 47 x 8
  id    wm    iq   comp simple_sentence complex_sentence log_simple log_complex
  <int> <int> <int> <int>          <int>          <int>      <dbl>       <dbl>
1 95     47    94    19        1960        2440      7.58       7.80
2 400    45   118    18        2186        2200      7.69       7.70
3 457    42   100    22        1797        2503      7.49       7.83
4 1138   41    77    18        2154        2441      7.68       7.80
5 1587   54    67    21        1936        2395      7.57       7.78
6 1805   52   109    19        1864        2560      7.53       7.85
7 1864   57   111    19        1930        2540      7.57       7.84
8 2183   55   125    23        1857        2324      7.53       7.75
9 2318   51    91    21        1918        2739      7.56       7.92
10 2324  43   120    20        1891        2426      7.54       7.79
# ... with 37 more rows
```

We could then run an ANCOVA over the log transformed RTs while covarying out the individual participant effects...

Problem - imagine our data are in the wrong ‘shape’ - they are in **Wide format** (each row is one *participant*) but we need them in **Long or Tidy format** (each row is one *observation*).

In SPSS, most data will be in Wide format with each experimental condition its own column:

```
> dataRT  
# A tibble: 48 x 3  
      id simple_sentence complex_sentence  
   <int>           <int>           <int>  
1    6400            1902            2341  
2     457             1797            2503  
3    8291            2080            2731  
4    4998            1856            2375  
5    2579            1997            2177  
6    9122            1868            2284  
7   1138             2154            2441  
8   5138             1933            2349  
9   5244             1900            2371  
10   3160             1929            2372  
# ... with 38 more rows
```

For many analyses in R, data need to be in Long format with each row being one observation. So, we want to transform our dataRT data frame so it looks like this:

id	condition	rt
...

To do this we can use the `gather()` function in the `tidyverse` package.

```
> data_long <- gather(dataRT, "condition", "rt", c("simple_sentence",  
"complex_sentence"))
```

The first parameter is the name of the data frame we want to reshape, the second is the name of the new ‘Key’ column, the third is the name of the new value column and the fourth the names of the columns we want to collapse.

We can use this to create a new data frame called `data_long` which looks like this:

```
> data_long <- gather(dataRT, "condition", "rt", c("simple_sentence",
  "complex_sentence"))
> data_long
# A tibble: 96 x 3
  id condition       rt
  <int> <chr>     <int>
1 6400 simple_sentence 1902
2 457  simple_sentence 1797
3 8291 simple_sentence 2080
4 4998 simple_sentence 1856
5 2579 simple_sentence 1997
6 9122 simple_sentence 1868
7 1138 simple_sentence 2154
8 5138 simple_sentence 1933
9 5244 simple_sentence 1900
10 3160 simple_sentence 1929
# ... with 86 more rows
```

And in reverse we can use the `spread()` function to go from Long to Wide data format:

```
> data_wide <- spread(data_long, "condition", "rt")
> data_wide
# A tibble: 48 x 3
  id complex_sentence simple_sentence
  <int>             <int>             <int>
1    95              2440              1960
2   400              2200              2186
3   457              2503              1797
4  1138              2441              2154
5  1587              2395              1936
6  1805              2560              1864
7  1864              2540              1930
8  2006              2267              2230
9  2183              2324              1857
10 2318              2739              1918
# ... with 38 more rows
```

We're now back to where we started with data in Wide format:

This is just a small example of functions in the `dplyr` and `tidyverse` packages that allow you to tidy, transform, and reshape your data. All of your code for doing this should appear at the start of your analysis script so that others (and you in 5 years or 5 days time) can see exactly what you did.

This allows for fully reproducible data preparation in the first part of your analysis workflow (important for Open Science and reproducibility).

Generating Descriptives - using dplyr

- You can use the `group_by()` and `summarise()` functions in the `dplyr` package to generate descriptives.
- In the following example, we are also using the pipe operator `%>%` which passes a value into an expression or function call from left to right:

```
> data_long %>%
  group_by(condition) %>%
  summarise(Mean = mean(rt), Min = min(rt), Max = max(rt), SD = sd(rt))
# A tibble: 2 x 5
  condition      Mean   Min   Max     SD
  <chr>        <dbl> <int> <int>  <dbl>
1 complex_sentence 2405.  2177  2739  132.
2 simple_sentence 1957.  1694  2356  147.
```

Tidying Up Some Real World Messy Data

Let's load another dataset that I created:

```
my_data <- read_csv("https://bit.ly/2YYcP4B")
```

Tidying Up Some Real World Messy Data

- We ran a reaction time experiment with 24 participants and 4 conditions - they are numbered 1-4 in our datafile.

```
> my_data
# A tibble: 96 x 3
  participant condition     rt
        <int>      <int> <int>
  1             1          1    879
  2             1          2   1027
  3             1          3   1108
  4             1          4    765
  5             2          1   1042
  6             2          2   1050
  7             2          3    942
  8             2          4    945
  9             3          1    943
 10            3          2    910
# ... with 86 more rows
```

- But actually it was a repeated measures design where we had one factor (Prime Type) with two levels (A vs. B) and a second factor (Target Type) with two levels (A vs. B)
- We want to recode our data frame so it better matches our experimental design.
- First we need to recode our 4 conditions like this:

```
# Recode condition columns follows:  
# Condition 1 = prime A, target A  
# Condition 2 = prime A, target B  
# Condition 3 = prime B, target A  
# Condition 4 = prime B, target B  
  
my_data <- my_data %>%  
  mutate(condition = recode(condition,  
    "1" = "primeA_targetA",  
    "2" = "primeA_targetB",  
    "3" = "primeB_targetA",  
    "4" = "primeB_targetB"))
```

- Now our data frame looks like this:

```
> my_data
# A tibble: 96 x 3
  participant condition       rt
  <int> <chr>     <int>
1 1 primeA_targetA 879
2 1 primeA_targetB 1027
3 1 primeB_targetA 1108
4 1 primeB_targetB 765
5 2 primeA_targetA 1042
6 2 primeA_targetB 1050
7 2 primeB_targetA 942
8 2 primeB_targetB 945
9 3 primeA_targetA 943
10 3 primeA_targetB 910
# ... with 86 more rows
```

- We then need to separate out our Condition column into two - one for our first factor (Prime), and one for our second factor (Target).

```
> my_data <- separate(my_data, col = "condition", into = c("prime",
  "target"), sep = "_")
> my_data
# A tibble: 96 x 4
  participant prime  target       rt
  <int> <chr>  <chr>     <int>
1          1 primeA targetA    879
2          1 primeA targetB   1027
3          1 primeB targetA   1108
4          1 primeB targetB    765
5          2 primeA targetA  1042
6          2 primeA targetB  1050
7          2 primeB targetA   942
8          2 primeB targetB   945
9          3 primeA targetA   943
10         3 primeA targetB   910
# ... with 86 more rows
```

- This is looking good - we now have our two factors coded separately and our data are in tidy format (i.e., one observation per row).
- How long would this have taken you in Excel? Would it have been reproducible?

- Perhaps we want to go from the data in long format, to wide format.

```
> my_data <- unite(my_data, col = "condition", c("prime", "target"), sep = "_")
> wide_data <- spread(my_data, key = "condition", value = "rt")
> wide_data
# A tibble: 24 x 5
  participant primeA_targetA primeA_targetB primeB_targetA primeB_targetB
      <int>        <int>        <int>        <int>        <int>
1         1          879        1027        1108        765
2         2         1042        1050        942         945
3         3          943        910         952        900
4         4          922        1006        1095        988
5         5          948        908         916       1241
6         6         1013        950         955       1045
7         7          930        855        1057        897
8         8          998        906        1110        952
9         9          929        949         837        883
10        10          781        865         970        953
# ... with 14 more rows
```

- No matter what format your data are in originally, you can use functions from the `dplyr` and `tidyverse` packages to quickly get it into whatever format you need for analysis.

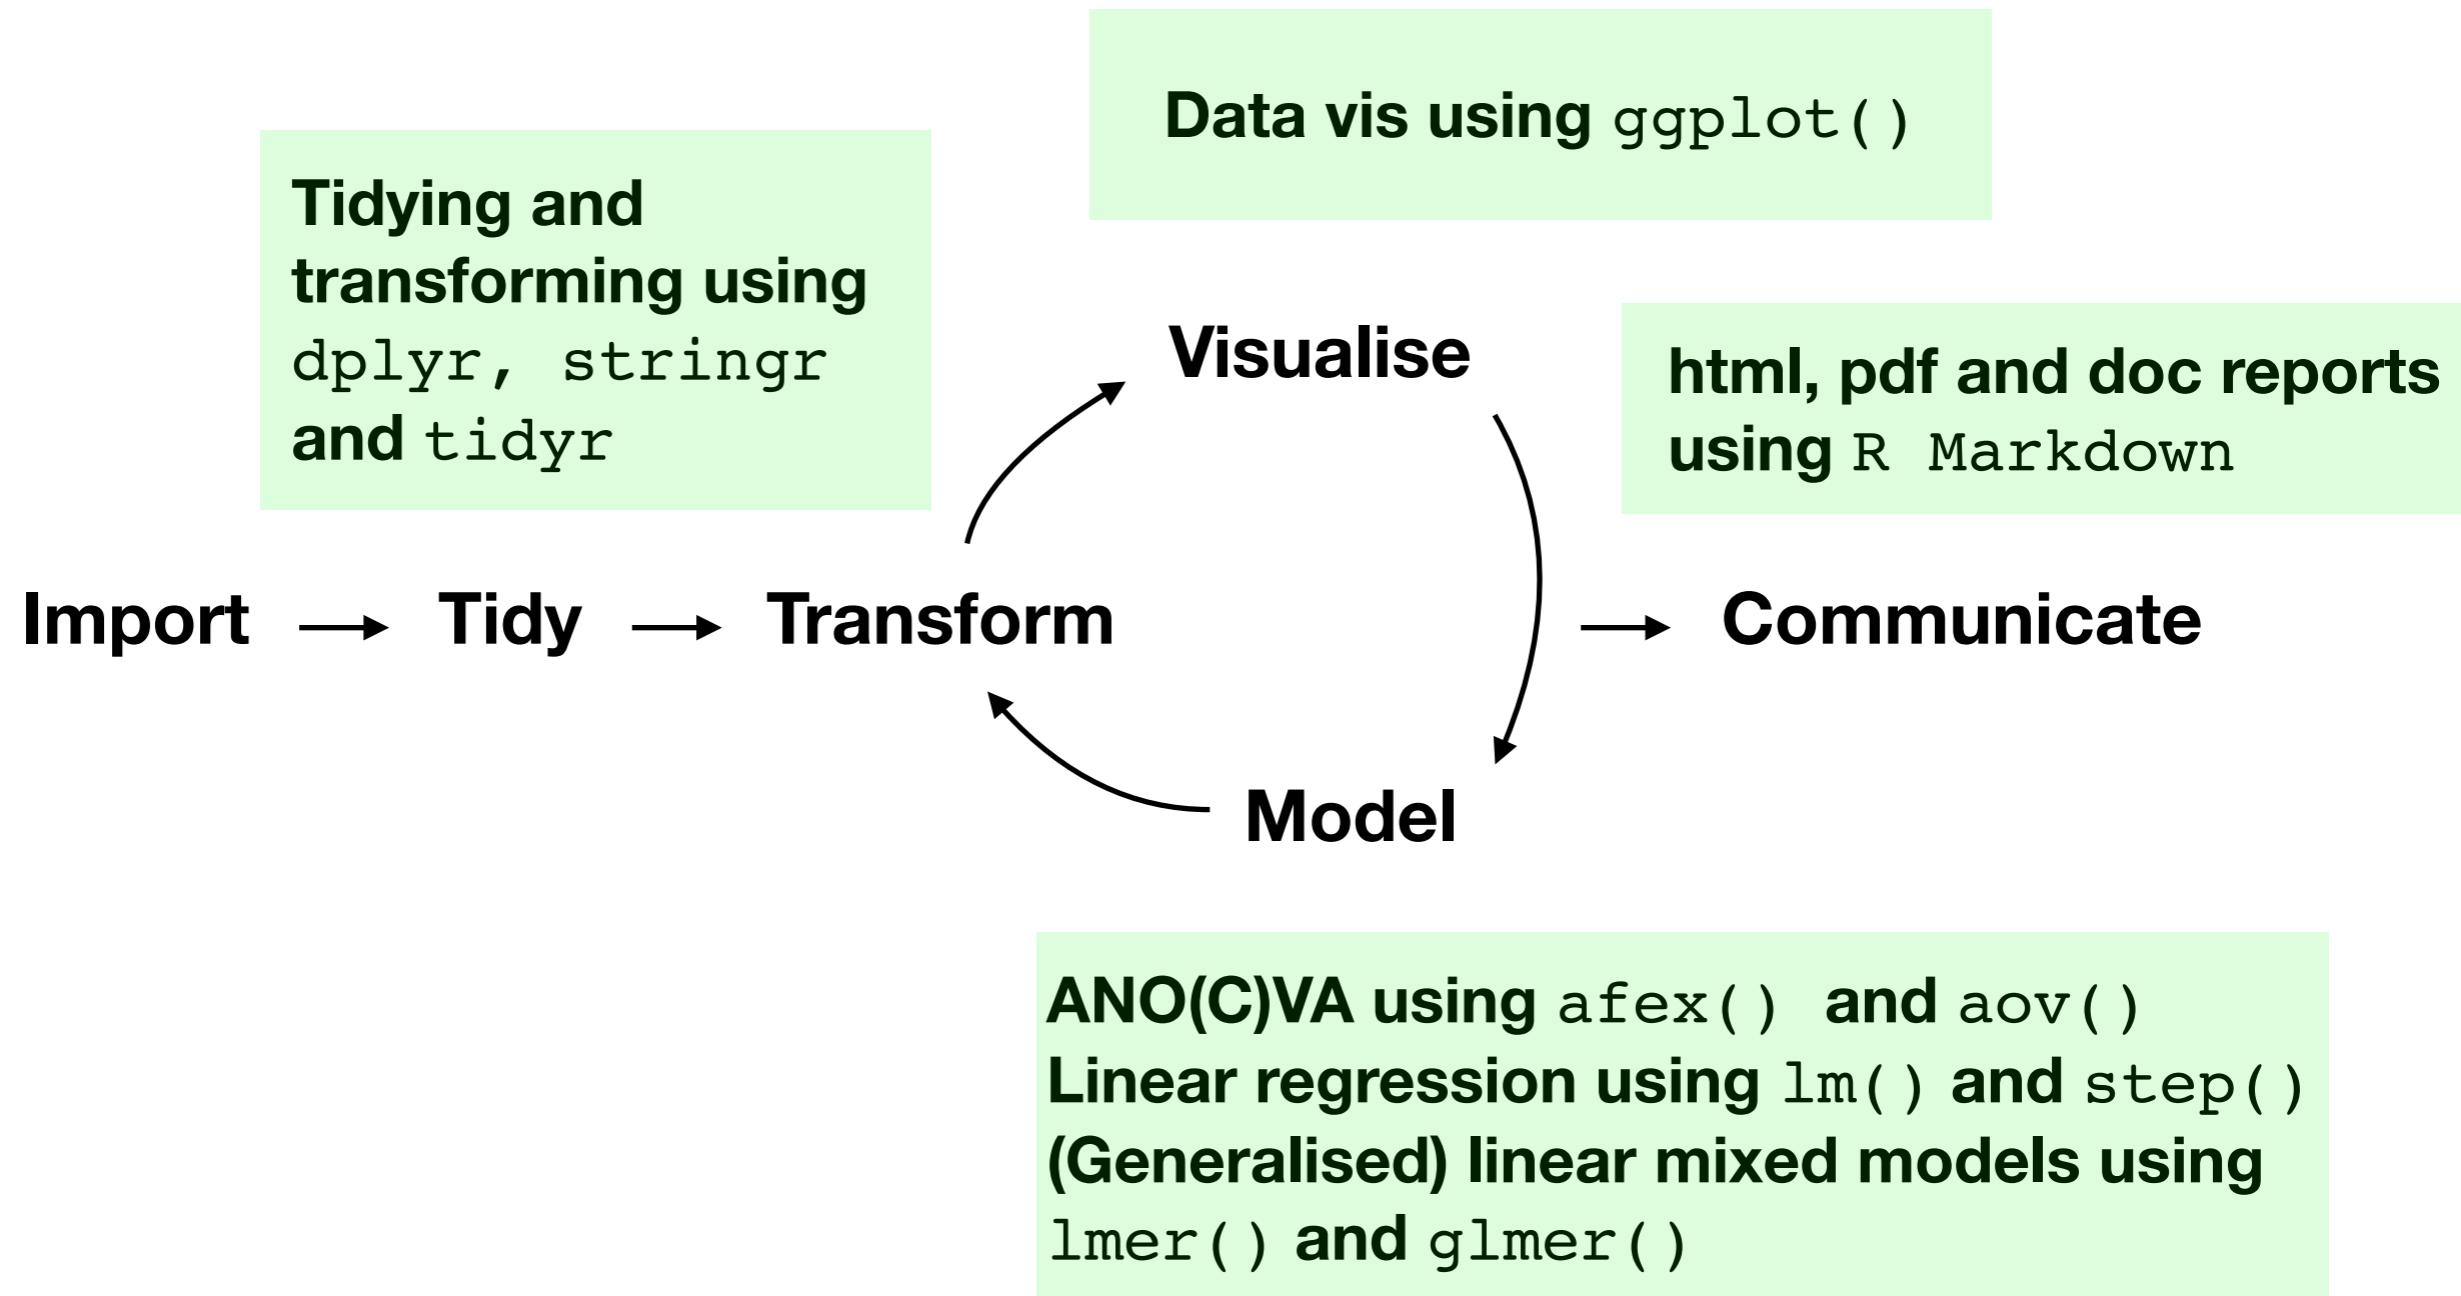
Or using the pipe...

- We could alternatively have used the `%>%` operator to combine all the last few operations which would have avoided the need to create temporary variables.

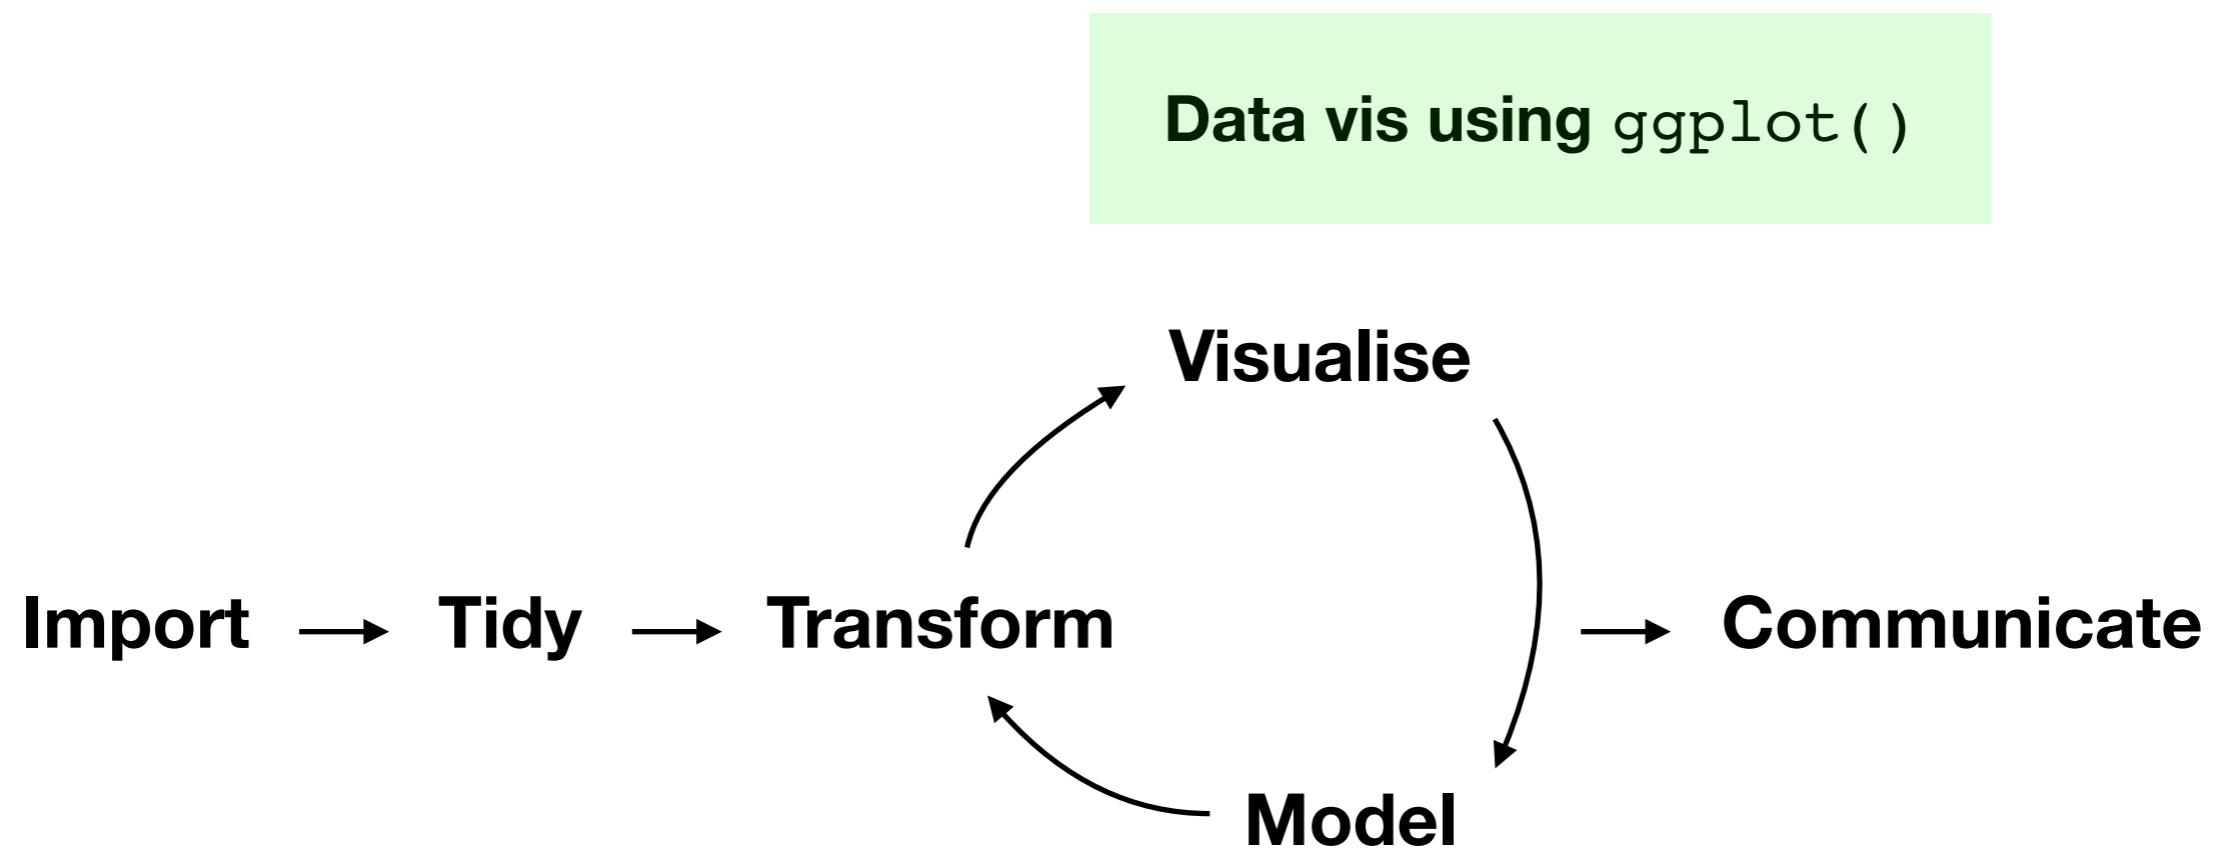
```
my_data %>%
  unite(col = "condition", c("prime", "target"), sep = "_") %>%
  spread(key = "condition", value = "rt")
```

- Take `my_data` and then `unite` and then `spread`...

A Reproducible Workflow



A Reproducible Workflow

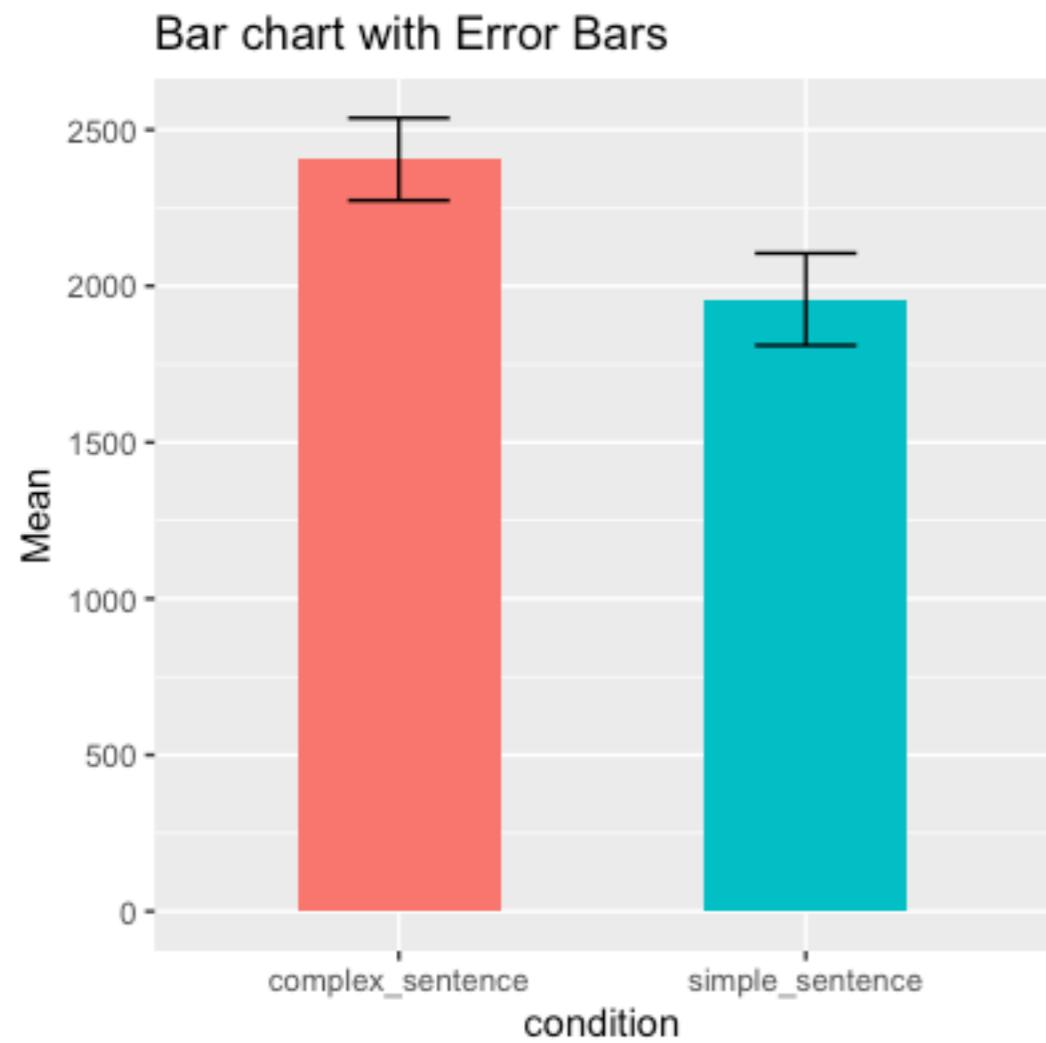


Visualising Your Data

- R has a number of in built (base) graphics functions, but you're more likely to use functions from within the `ggplot2` package. `ggplot2` is part of the `tidyverse` so if you have used `library(tidyverse)` then `ggplot2` will already be loaded.

```
> library(ggplot2)
```

Bar Graphs (bleurgh!)

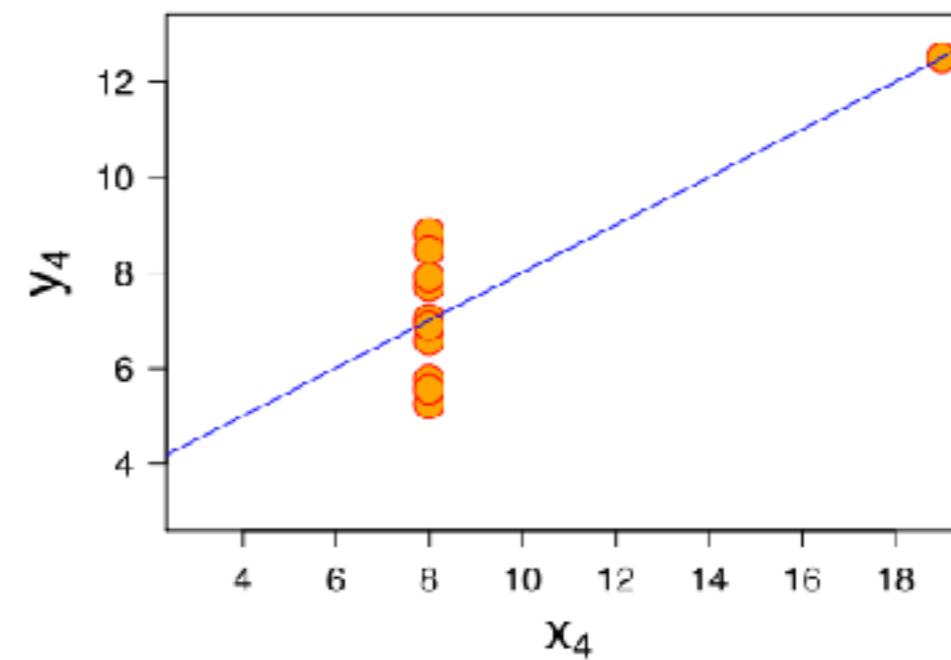
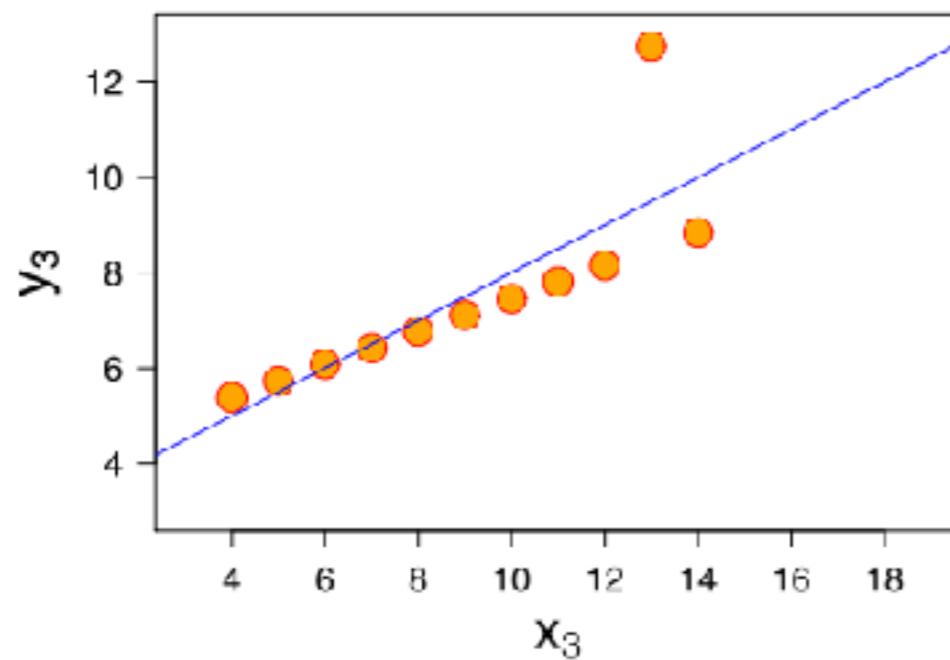
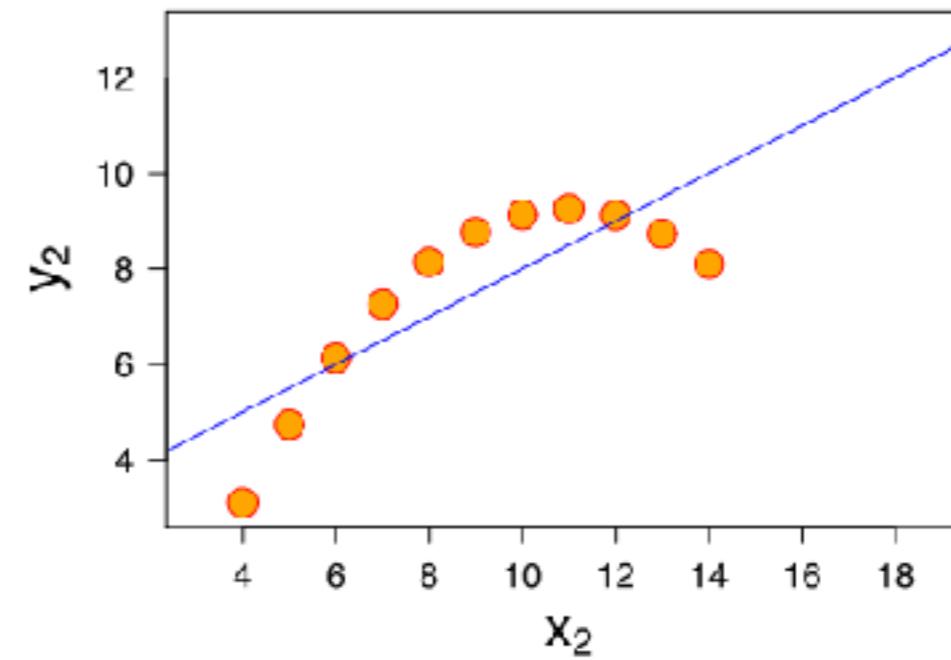
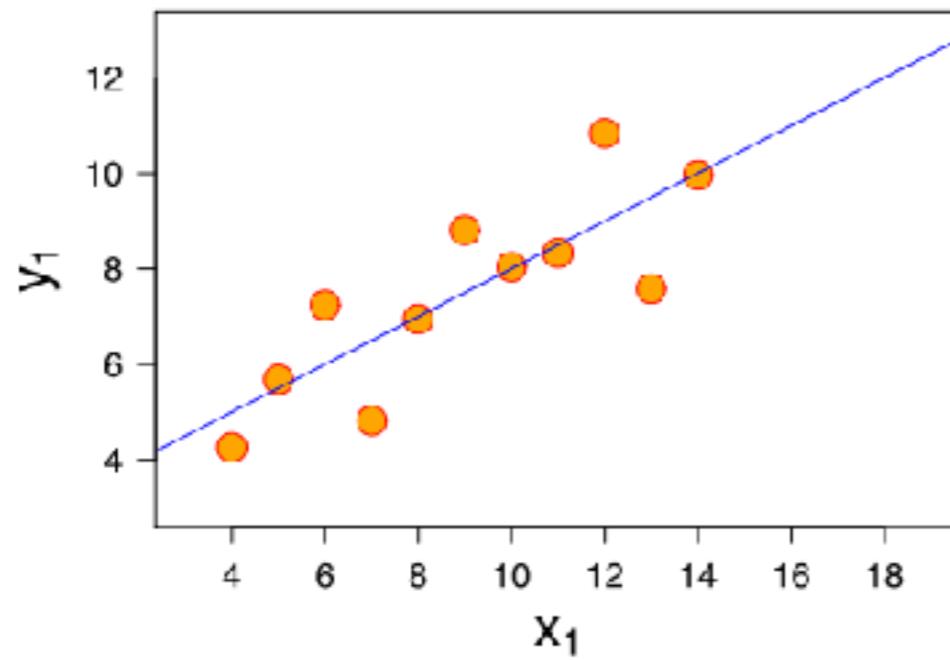


Bar graphs tend to be quite limited in terms of what they communicate. Here they communicate the means for levels of a factor and information about variance. But they don't tell us anything about the *distribution* of the data.

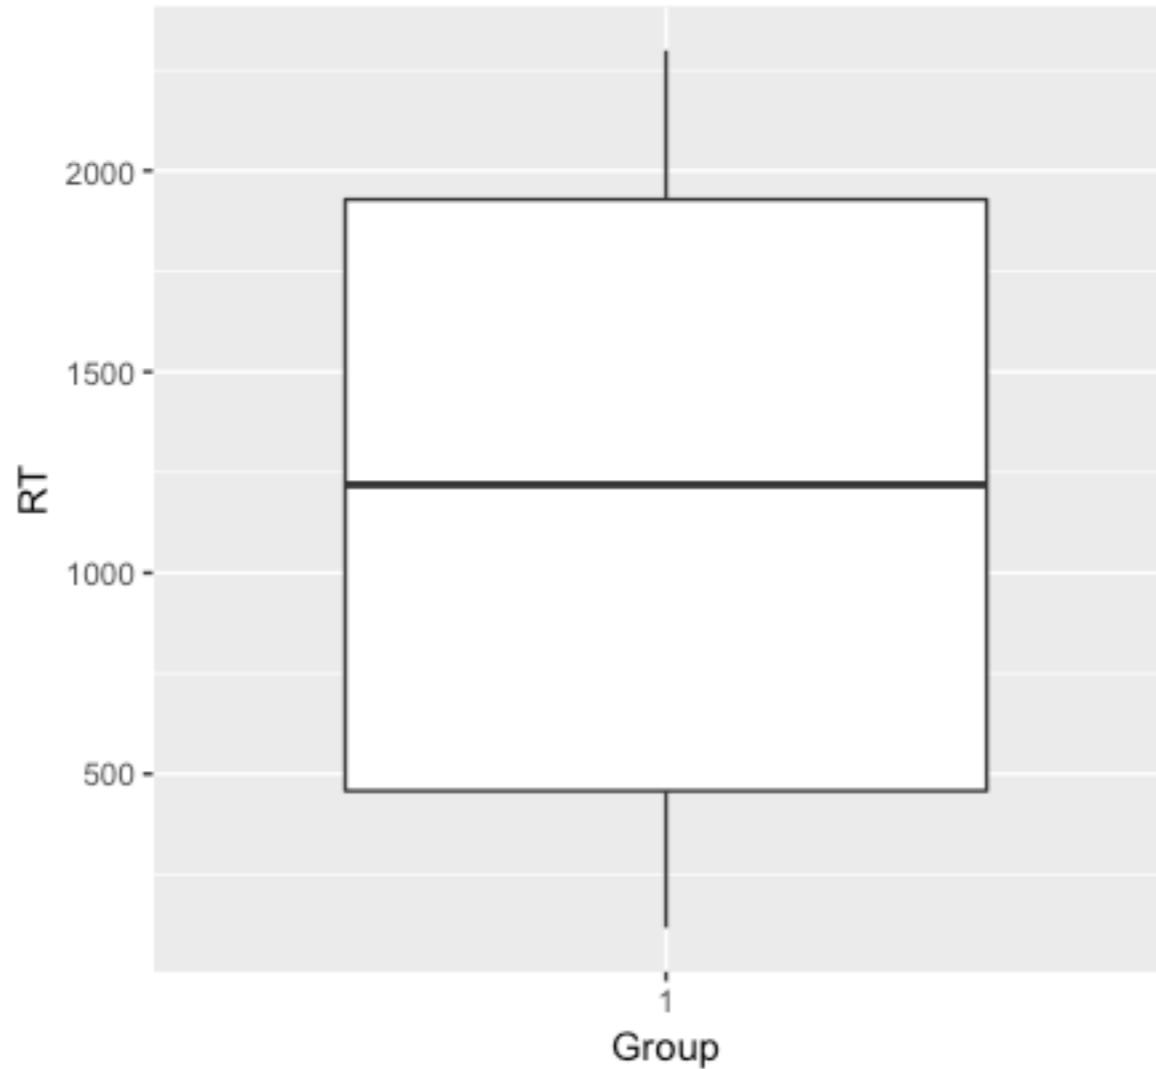
```
data_summ <- data_long %>% group_by(condition) %>% summarise(Mean = mean(rt), sd = sd(rt))

ggplot(data_summ, aes(x = condition, y = Mean, group = condition,
                      fill = condition, ymin = Mean - sd, ymax = Mean + sd)) +
  geom_bar(stat = "identity", width = .5) +
  geom_errorbar(width = .25) +
  ggtitle("Bar chart with Error Bars") +
  guides(fill = FALSE)
```

Anscombe's Quartet

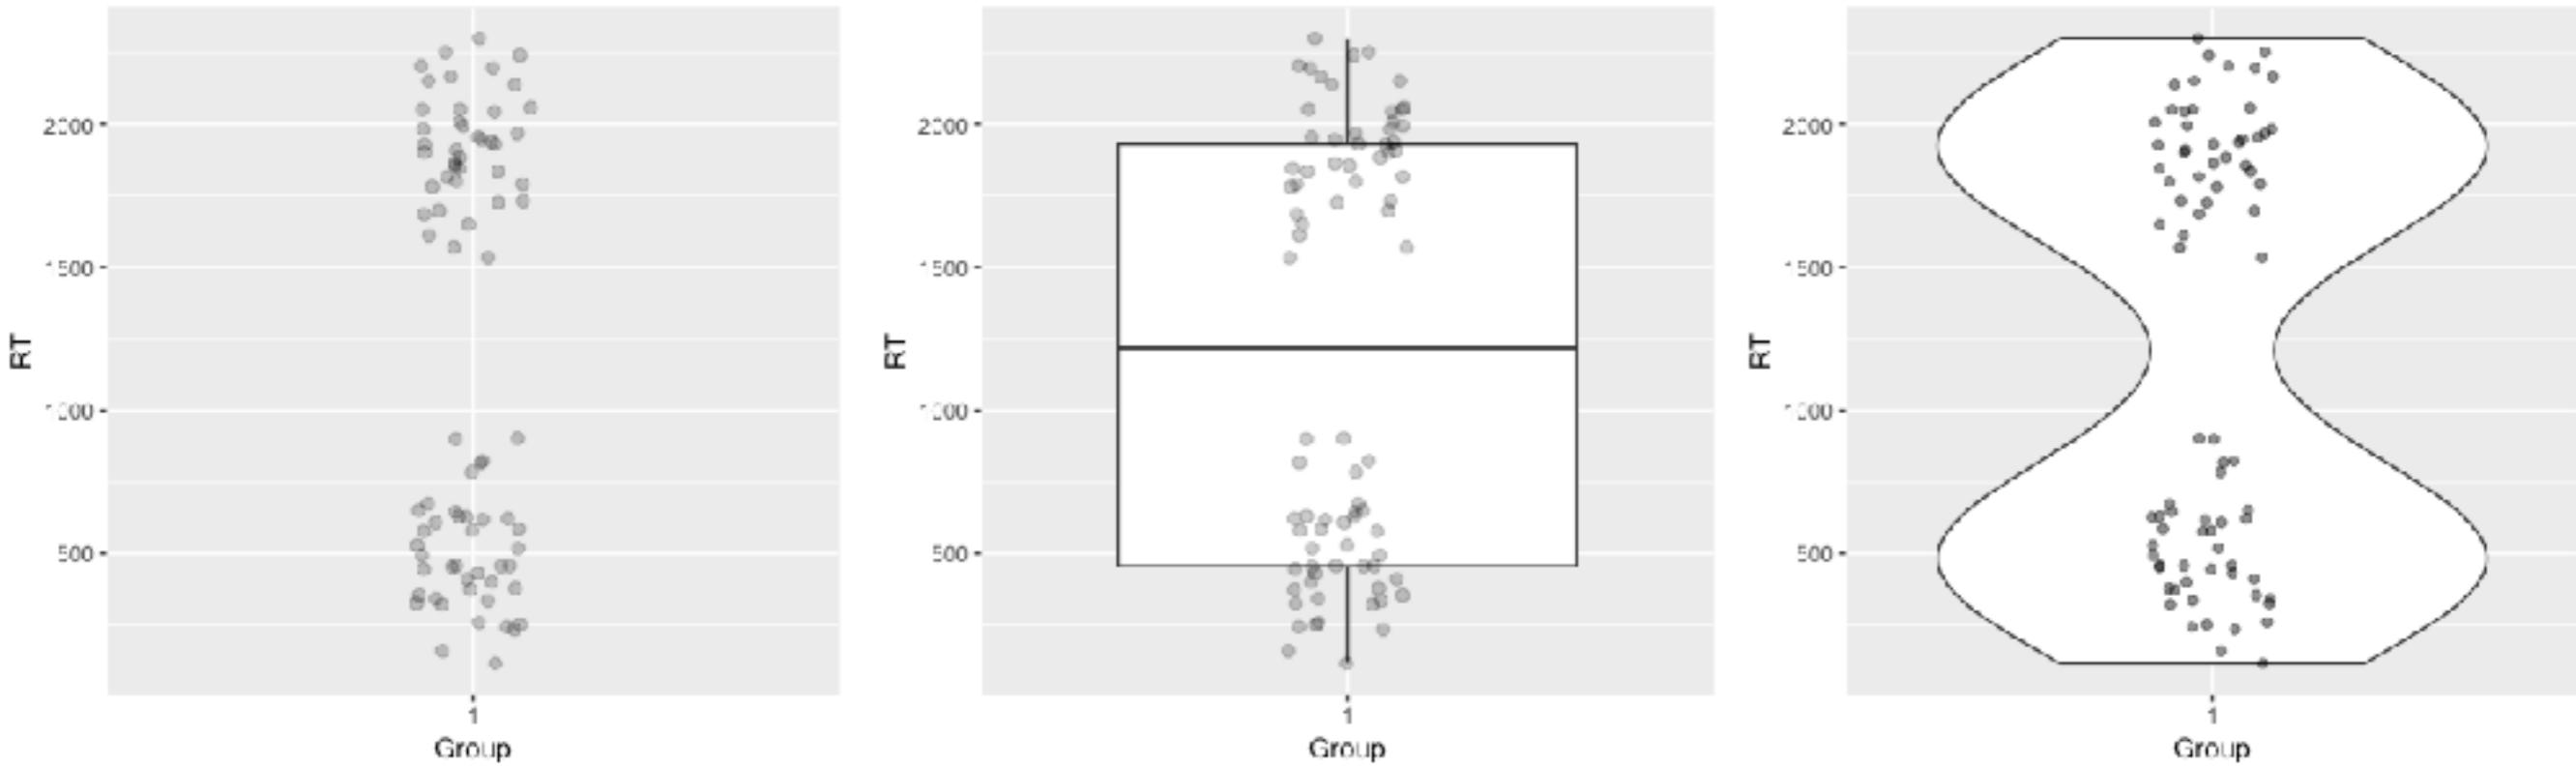


Plots Based on Aggregated Data Can Mislead...



You might make one set of inferences based on this boxplot - maybe a median around 1,250 with the 25th and 75th percentiles being ~480 to ~1,980...

But look more closely at the actual data...

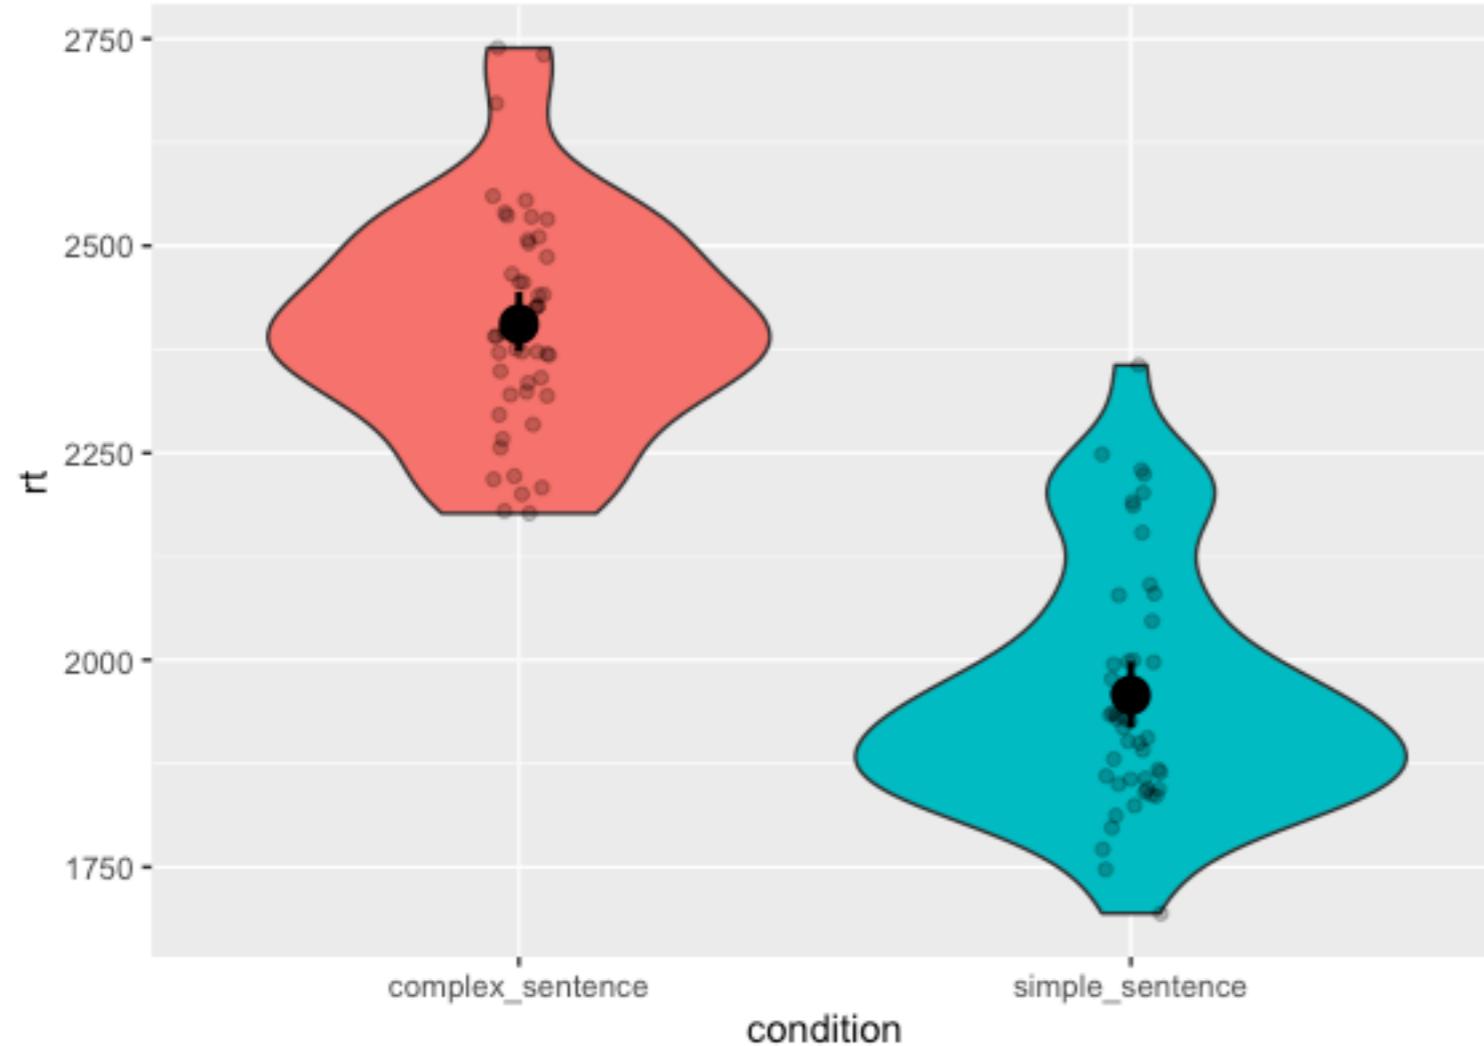


The data are clearly bimodal with no actual data point near the mean.
Distribution shape matters and we need to capture that in our data visualisations.

You try...

```
data2 <- read_csv("https://bit.ly/2YYr6dm")  
  
ggplot(data2, aes(x = group, y = rt)) +  
  geom_boxplot()  
  
ggplot(data2, aes(x = group, y = rt)) +  
  geom_jitter(size = 2, width = .1, alpha = .25)  
  
ggplot(data2, aes(x = group, y = rt)) +  
  geom_boxplot() +  
  geom_jitter(size = 2, width = .1, alpha = .25)  
  
ggplot(data2, aes(x = group, y = rt)) +  
  geom_violin() +  
  geom_jitter(width = .1, alpha = .5)
```

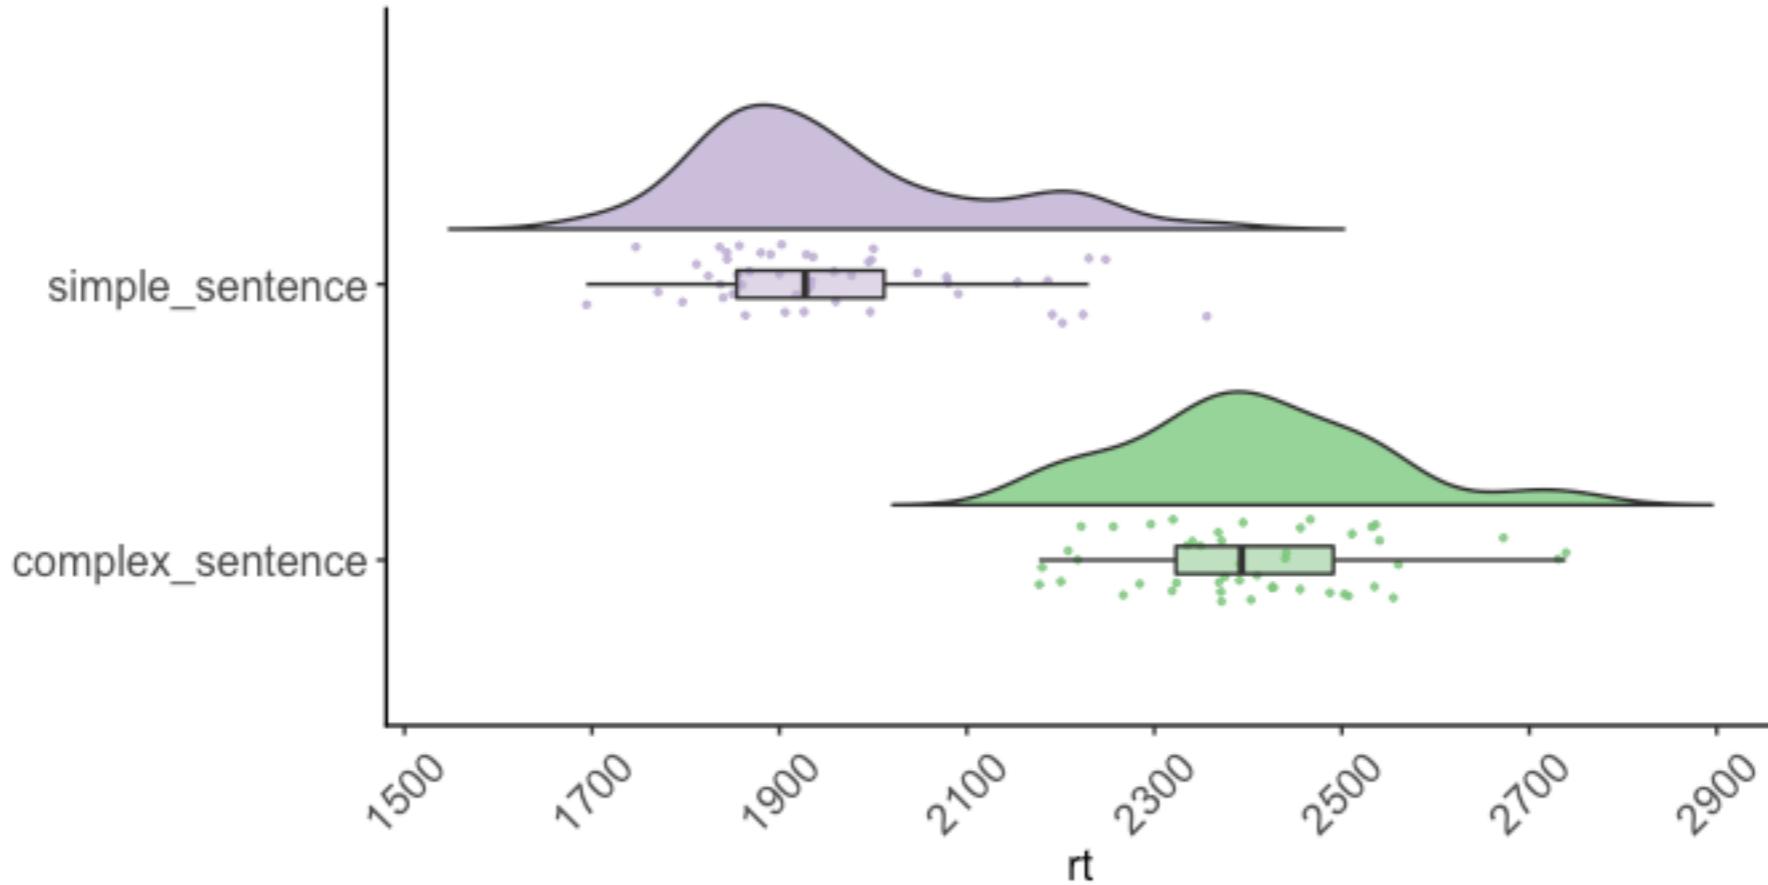
Violin Plots



Violin plots tell us about the distribution of the data. The width at any point corresponds to the *density* of the data at that value.

```
ggplot(data_long, aes(x = condition, y = rt,
  group = condition, fill = condition)) +
  geom_violin() +
  geom_jitter(alpha = .25, position = position_jitter(0.05)) +
  guides(colour = FALSE, fill = FALSE) +
  stat_summary(fun.data = "mean_cl_boot", colour = "black", size = 1)
```

Raincloud Plots



Developed by Micah Allen (UCL), raincloud plots allow you to see the raw data, and the shape of the distribution alongside a box plot (capturing the median, 25th and 75th percentiles as hinges, and $1.5 * \text{IQR}$ from the hinges as the whisker length.)

A Variety of Plots Using the Same Dataset

We're going to use the built-in dataset 'mpg' to build a variety of plots. First, let's find out about the data by using the head function to view the first part of the data.

```
> head(mpg)
# A tibble: 6 x 11
  manufacturer model displ year cyl trans   drv   cty   hwy fl class
  <chr>        <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
1 audi         a4     1.8  1999    4 auto (15) f      18     29 p   compact
2 audi         a4     1.8  1999    4 manual (m5) f      21     29 p   compact
3 audi         a4     2.0  2008    4 manual (m6) f      20     31 p   compact
4 audi         a4     2.0  2008    4 auto (av)   f      21     30 p   compact
5 audi         a4     2.8  1999    6 auto (15) f      16     26 p   compact
6 audi         a4     2.8  1999    6 manual (m5) f      18     26 p   compact
```

We can explore the data further by asking for all the possibilities in each column using the `unique` function. For example, we can check to see how many different types of cars there are. Note that the \$ after the dataset name allows us to refer to a column in the mpg dataset.

```
> unique(mpg$manufacturer)
[1] "audi"        "chevrolet"    "dodge"       "ford"        "honda"       "hyundai"     "jeep"
[8] "land rover" "lincoln"      "mercury"     "nissan"     "pontiac"     "subaru"     "toyota"
[15] "volkswagen"
```

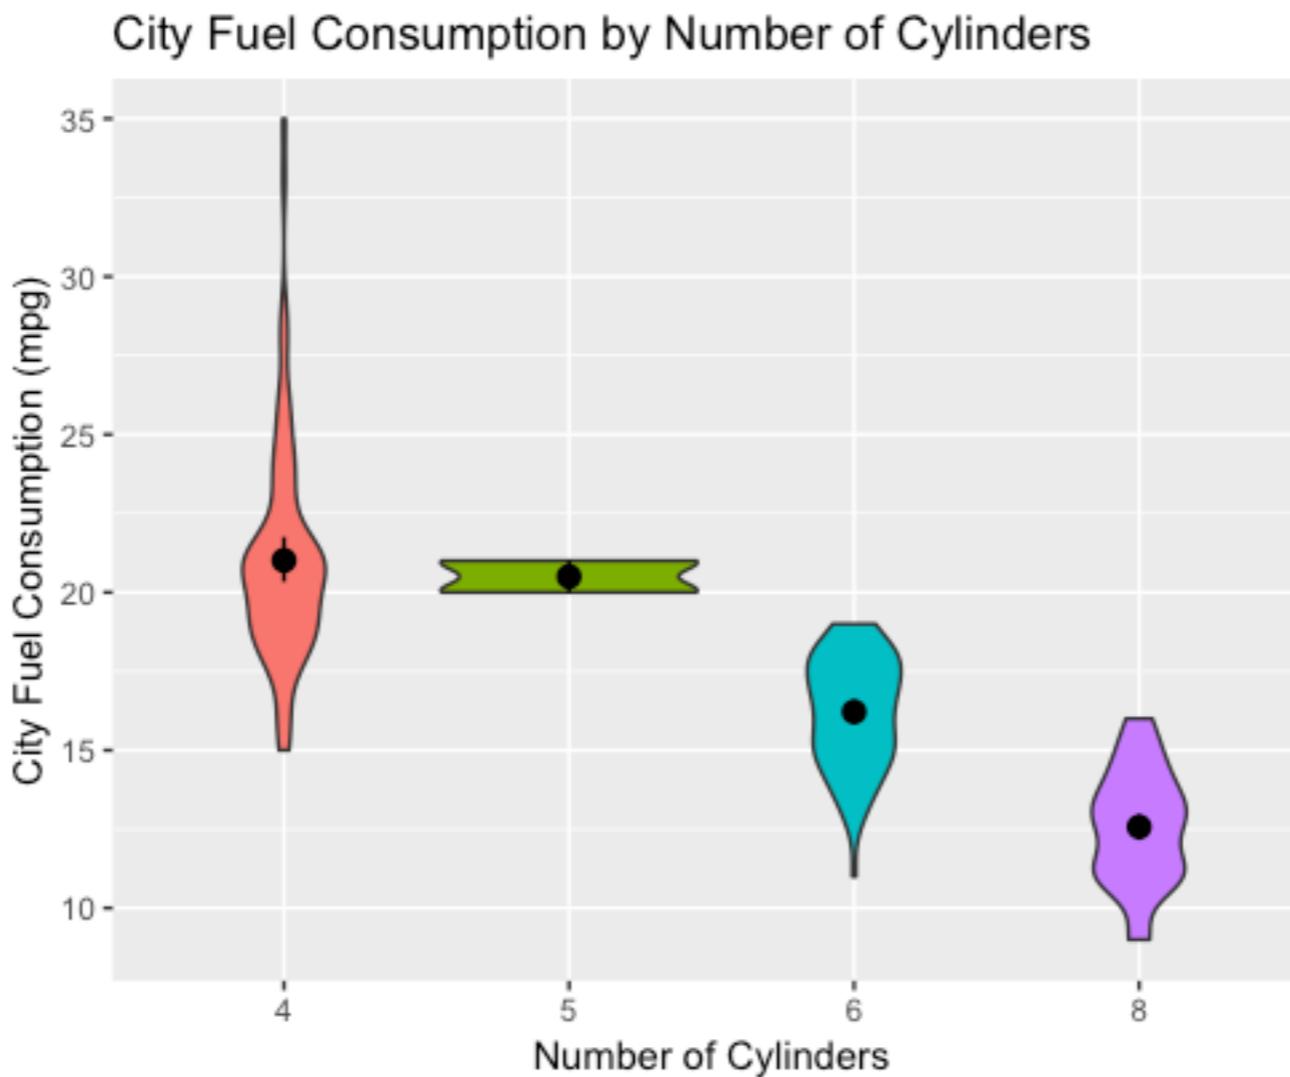
We can use the `length` function to give us the total number of unique possibilities:

```
> length(unique(mpg$manufacturer))
[1] 15
```

Let's look at a whole bunch of different visualisations using the mpg data set...

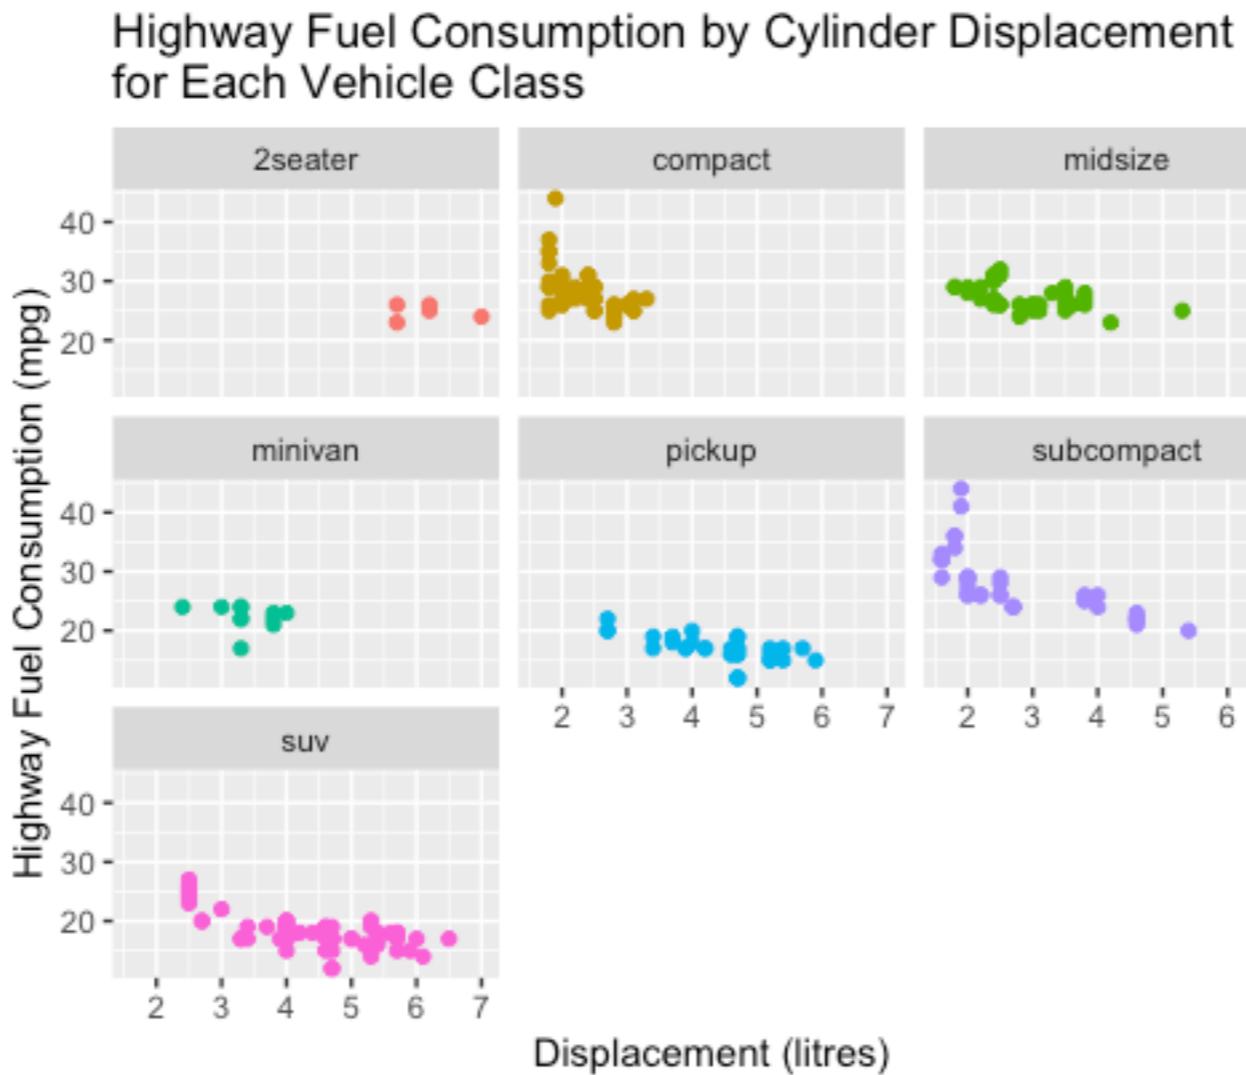
This illustrates the idea that there is not one 'correct' way to visualise the data, but rather that your choice of visualisation will be influenced by the question you're investigating, or the story you're wanting to tell...

Violin Plots



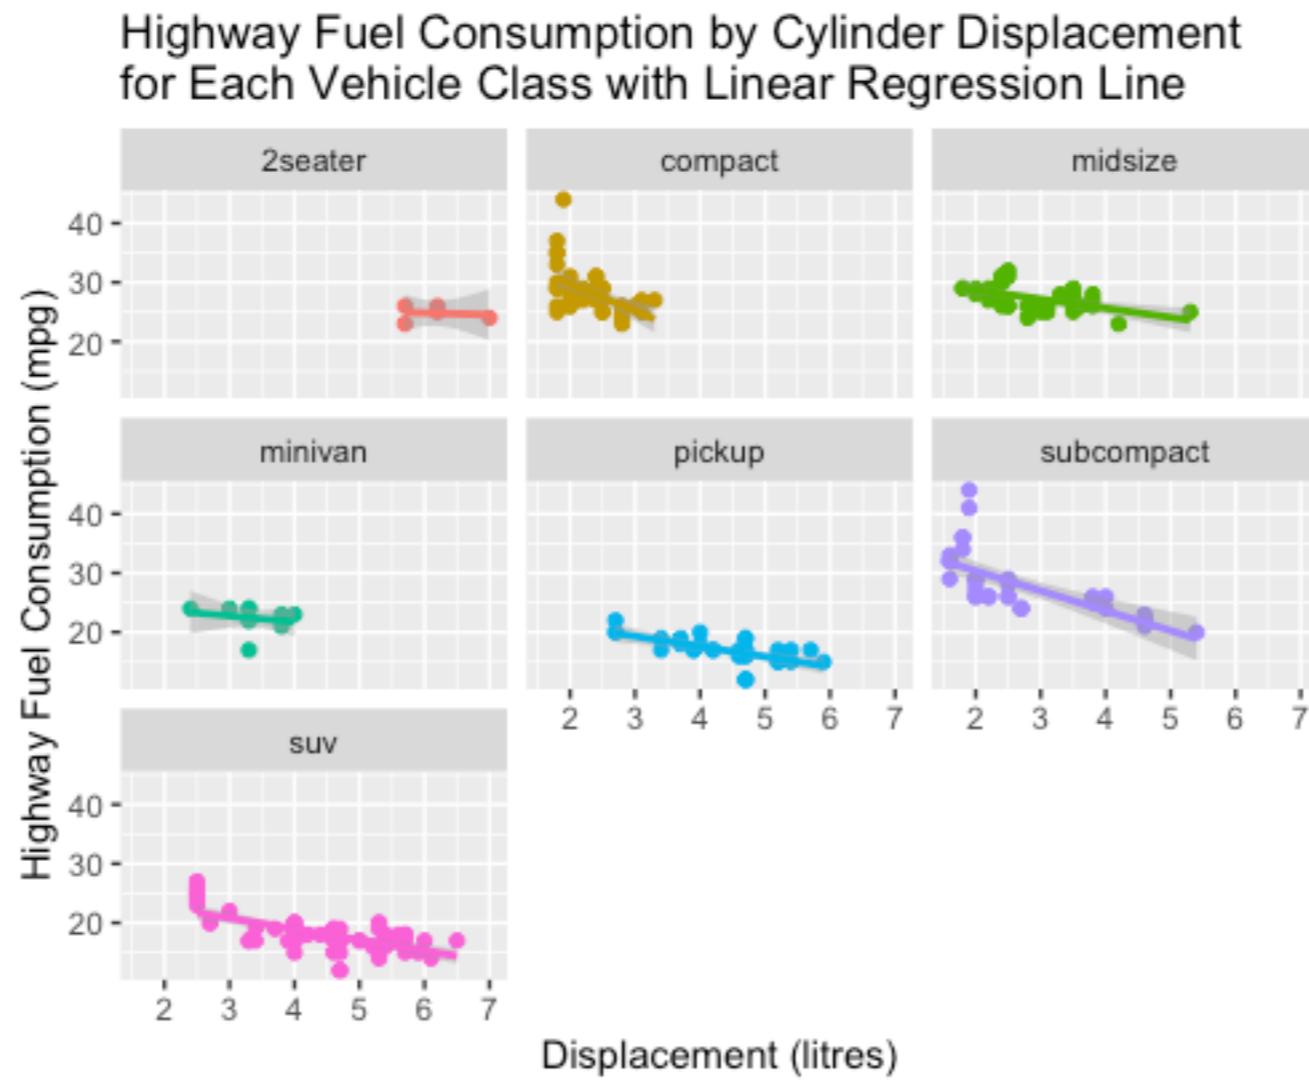
```
ggplot(mpg, aes(x = factor(cyl), y = cty, fill = factor(cyl))) +  
  geom_violin() +  
  guides(colour = FALSE, fill = FALSE) +  
  stat_summary(fun.data = mean_cl_boot, colour = "black", size = .5) +  
  labs(title = "City Fuel Consumption by Number of Cylinders",  
       x = "Number of Cylinders",  
       y = "City Fuel Consumption (mpg)")
```

Faceting



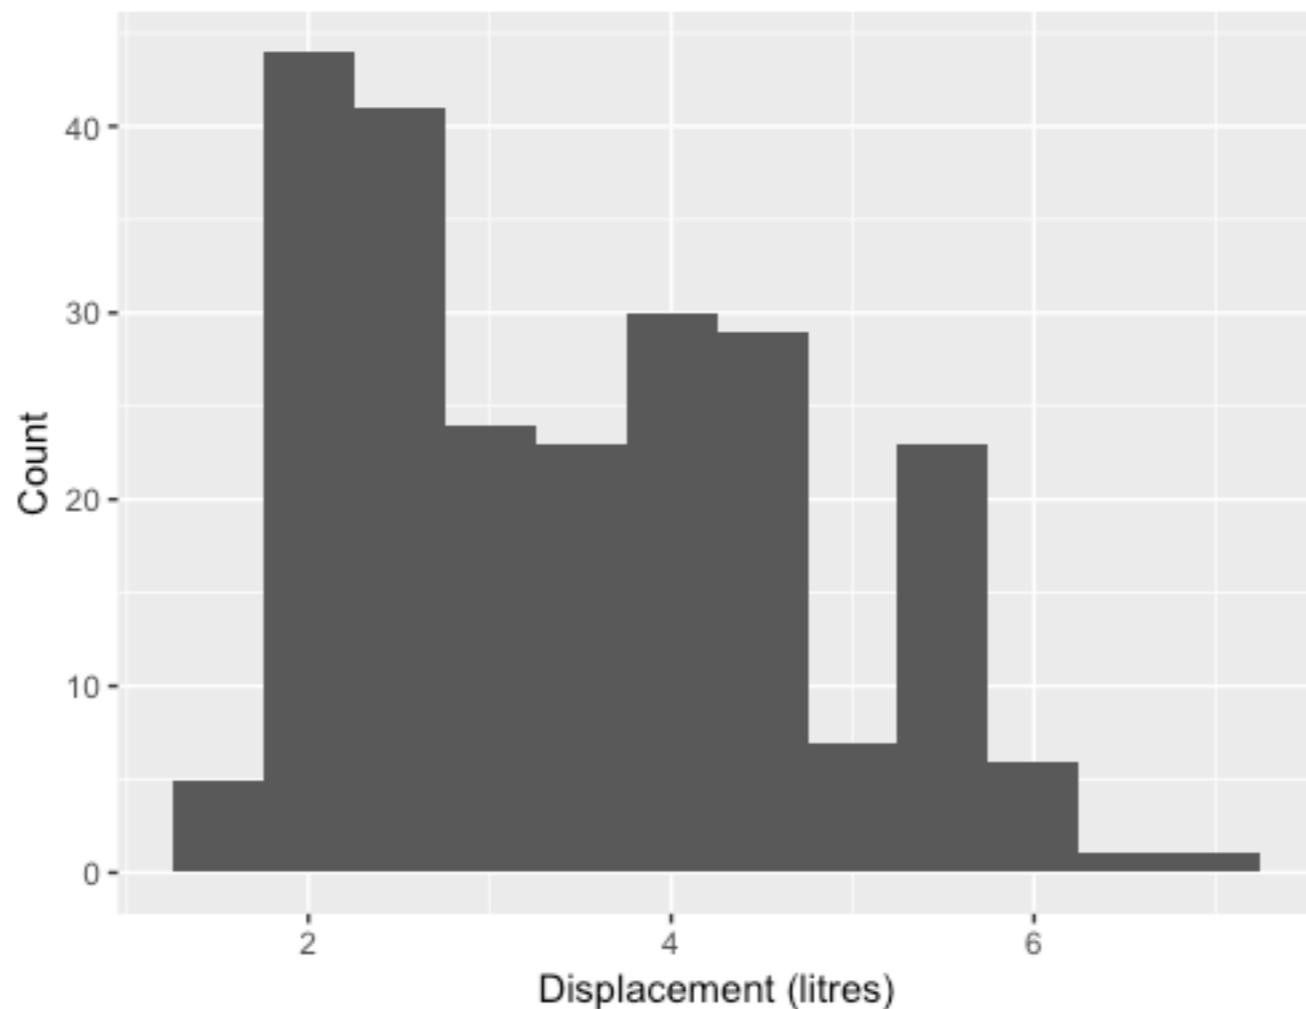
```
ggplot(mpg, aes(x = displ, y = hwy, colour = class)) +  
  geom_point() +  
  facet_wrap(~ class) +  
  guides(colour = FALSE) +  
  labs(title = "Highway Fuel Consumption by Cylinder Displacement \nfor Each Vehicle Class",  
       x = "Displacement (litres)",  
       y = "Highway Fuel Consumption (mpg)")
```

Adding a regression line



```
ggplot(mpg, aes(x = displ, y = hwy, colour = class)) +  
  geom_point() +  
  facet_wrap(~ class) +  
  guides(colour = FALSE) +  
  geom_smooth(method = "lm") +  
  labs(title = "Highway Fuel Consumption by Cylinder Displacement \nfor Each Vehicle Class",  
       x = "Displacement (litres)",  
       y = "Highway Fuel Consumption (mpg)")
```

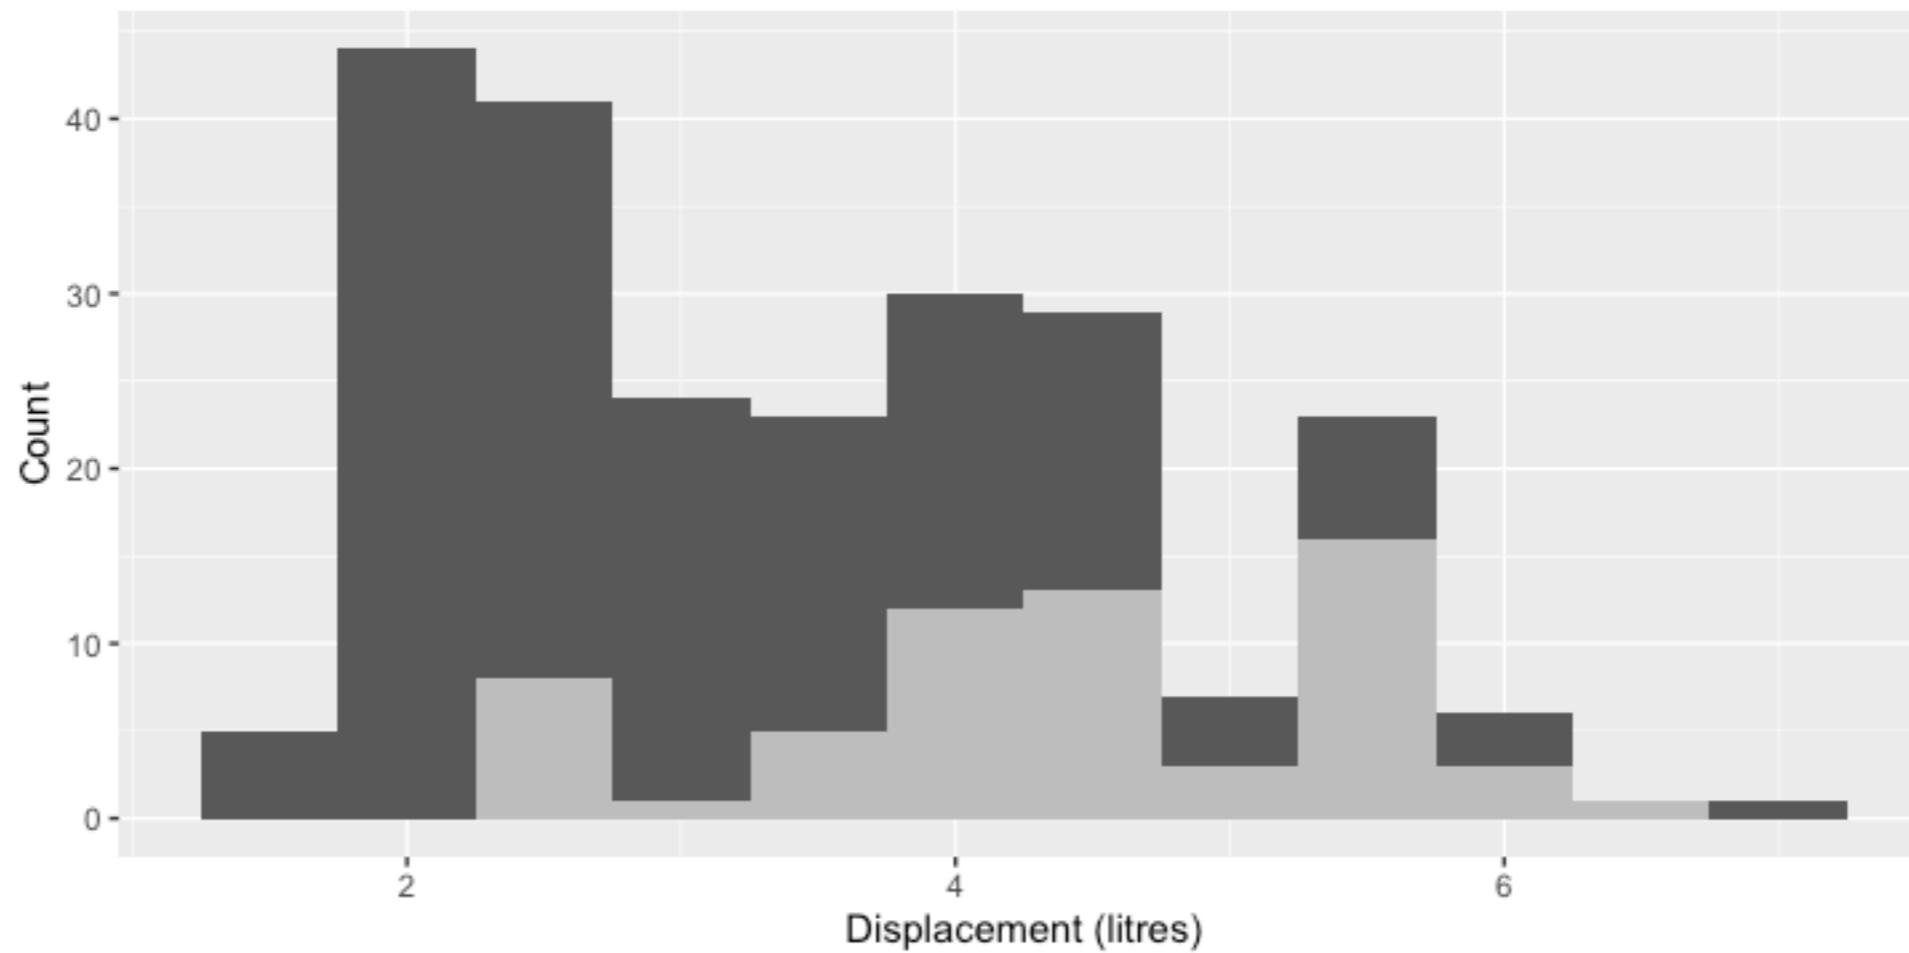
Histogram of Cylinder Displacement



```
ggplot(mpg, aes(x = displ)) +  
  geom_histogram(binwidth = .5) +  
  guides(fill = FALSE) +  
  labs(title = "Histogram of Cylinder Displacement",  
       x = "Displacement (litres)",  
       y = "Count")
```

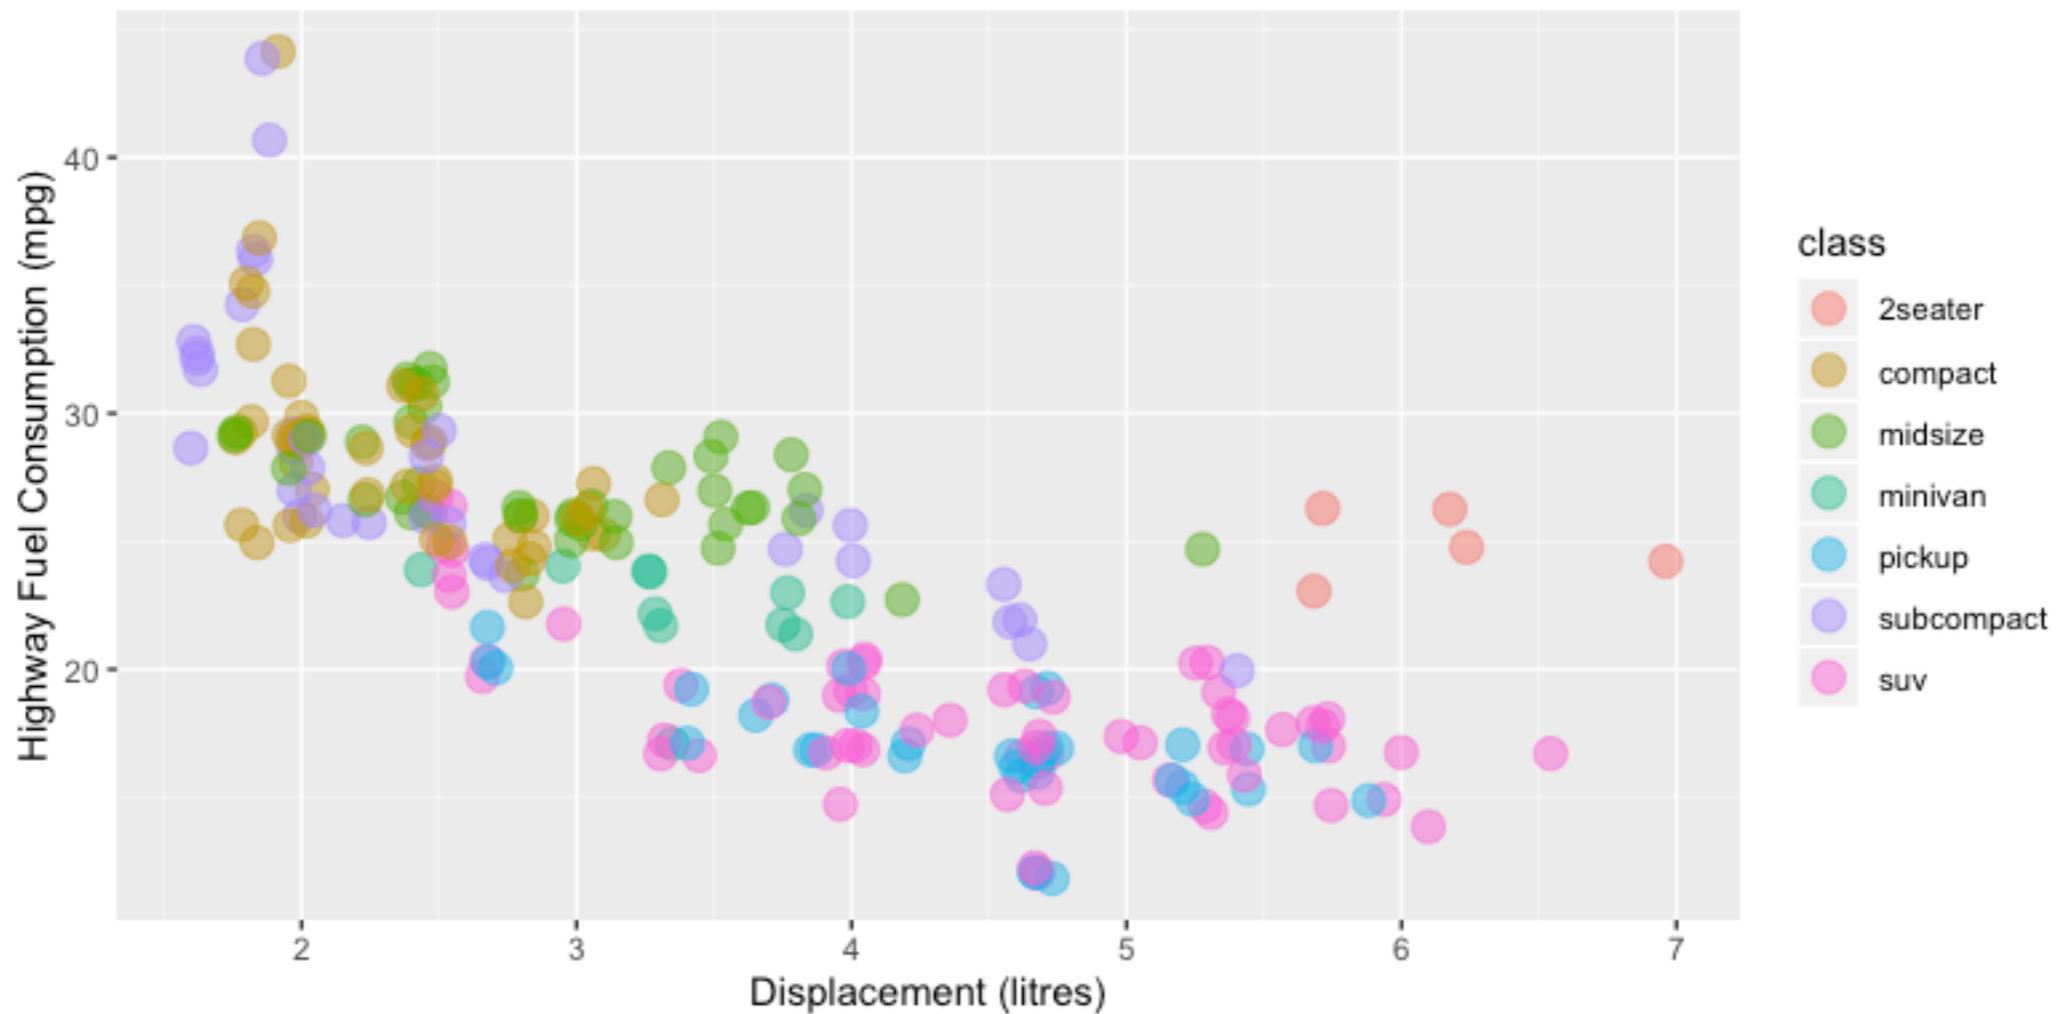
Histogram of Cylinder Displacement

SUVs highlighted



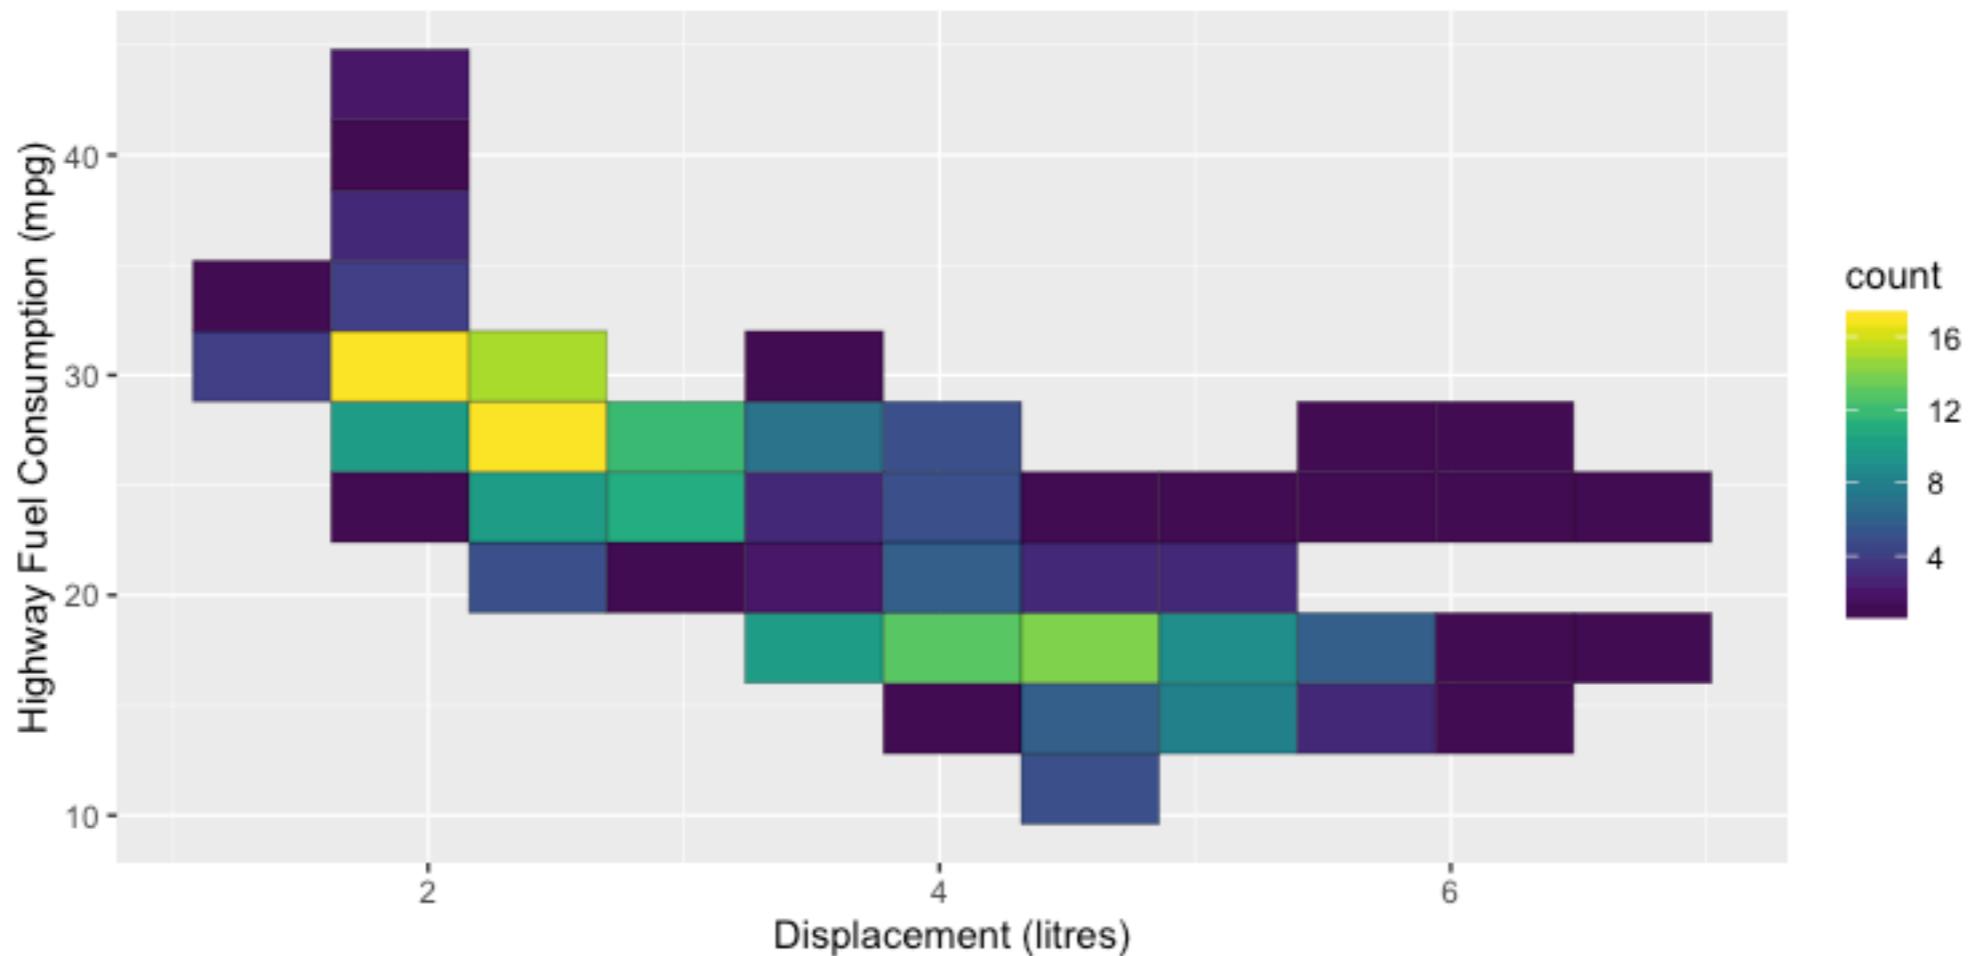
```
ggplot(mpg, aes(x = displ)) +  
  geom_histogram(binwidth = .5) +  
  geom_histogram(data = filter(mpg, class == "suv"), fill = "grey", binwidth = .5) +  
  guides(fill = FALSE) +  
  labs(title = "Histogram of Cylinder Displacement",  
       subtitle = "SUVs highlighted",  
       x = "Displacement (litres)",  
       y = "Count")
```

Scatterplot of Highway Fuel Consumption against
Engine Displacement Grouped by Class



```
ggplot(mpg, aes(x = displ, y = hwy, colour = class)) +  
  geom_jitter(width = 0.05, alpha = .5, size = 4) +  
  labs(title = "Scatterplot of Highway Fuel Consumption against\nEngine  
Displacement Grouped by Class",  
    x = "Displacement (litres)",  
    y = "Highway Fuel Consumption (mpg)")
```

Density heatmap of Highway Fuel Consumption against Engine Displacement



```
ggplot(mpg, aes(x = displ, y = hwy)) +  
  stat_bin2d(bins = 10, colour = "black") +  
  scale_fill_viridis() +  
  labs(title = "Density heatmap of Highway Fuel Consumption against Engine Displacement",  
       x = "Displacement (litres)",  
       y = "Highway Fuel Consumption (mpg)")
```

Plotting Time Series Data

We can install the `gapminder` package which contains lots of interesting data about life expectancy, population size, GDP for lots of countries collected over lots of years.

```
> gapminder
# A tibble: 1,704 x 6
  country      continent    year  lifeExp      pop  gdpPercap
  <fct>        <fct>     <int>   <dbl>     <int>      <dbl>
1 Afghanistan Asia      1952    28.8  8425333    779.
2 Afghanistan Asia      1957    30.3  9240934    821.
3 Afghanistan Asia      1962    32.0  10267083   853.
4 Afghanistan Asia      1967    34.0  11537966   836.
5 Afghanistan Asia      1972    36.1  13079460   740.
6 Afghanistan Asia      1977    38.4  14880372   786.
7 Afghanistan Asia      1982    39.9  12881816   978.
8 Afghanistan Asia      1987    40.8  13867957   852.
9 Afghanistan Asia      1992    41.7  16317921   649.
10 Afghanistan Asia     1997    41.8  22227415   635.
# ... with 1,694 more rows
```

Animated visualisations

For datasets with time series information, we might think it could be easier to tell our data story if we animate by time.

The `ggridge` package allows us to create animated visualisations from within R which we can import or embed in R Markdown documents.

We need to install it via the usual route:

```
> install.packages("ggridge")
```

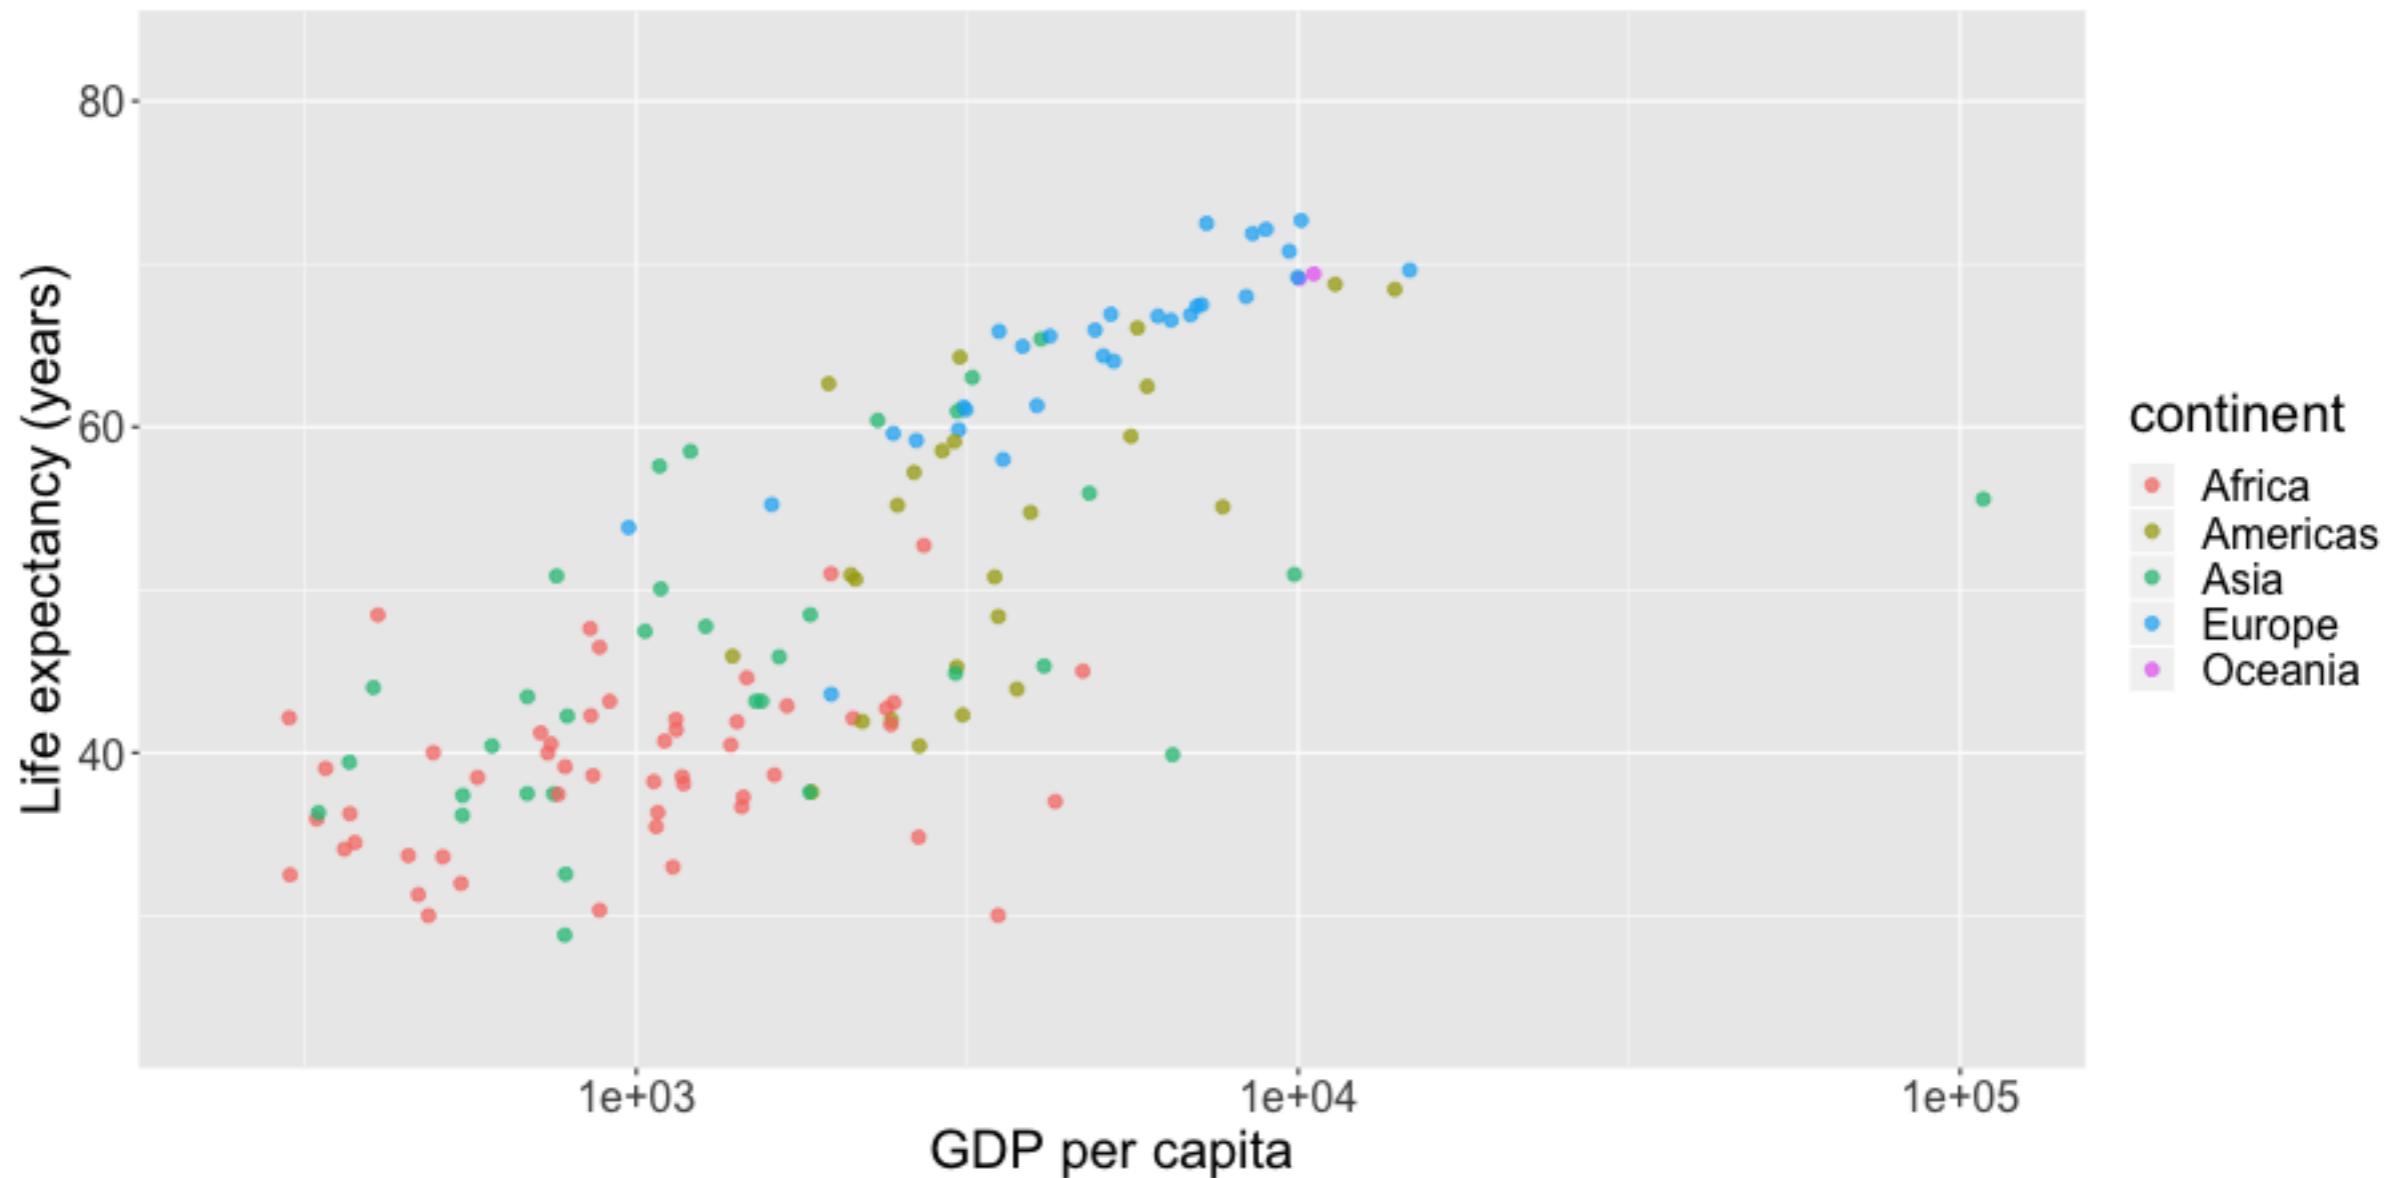
and then load it when we want to use it:

```
> library(ggridge)
```

Animated Time Series Data

Gapminder dataset

Year: 1952



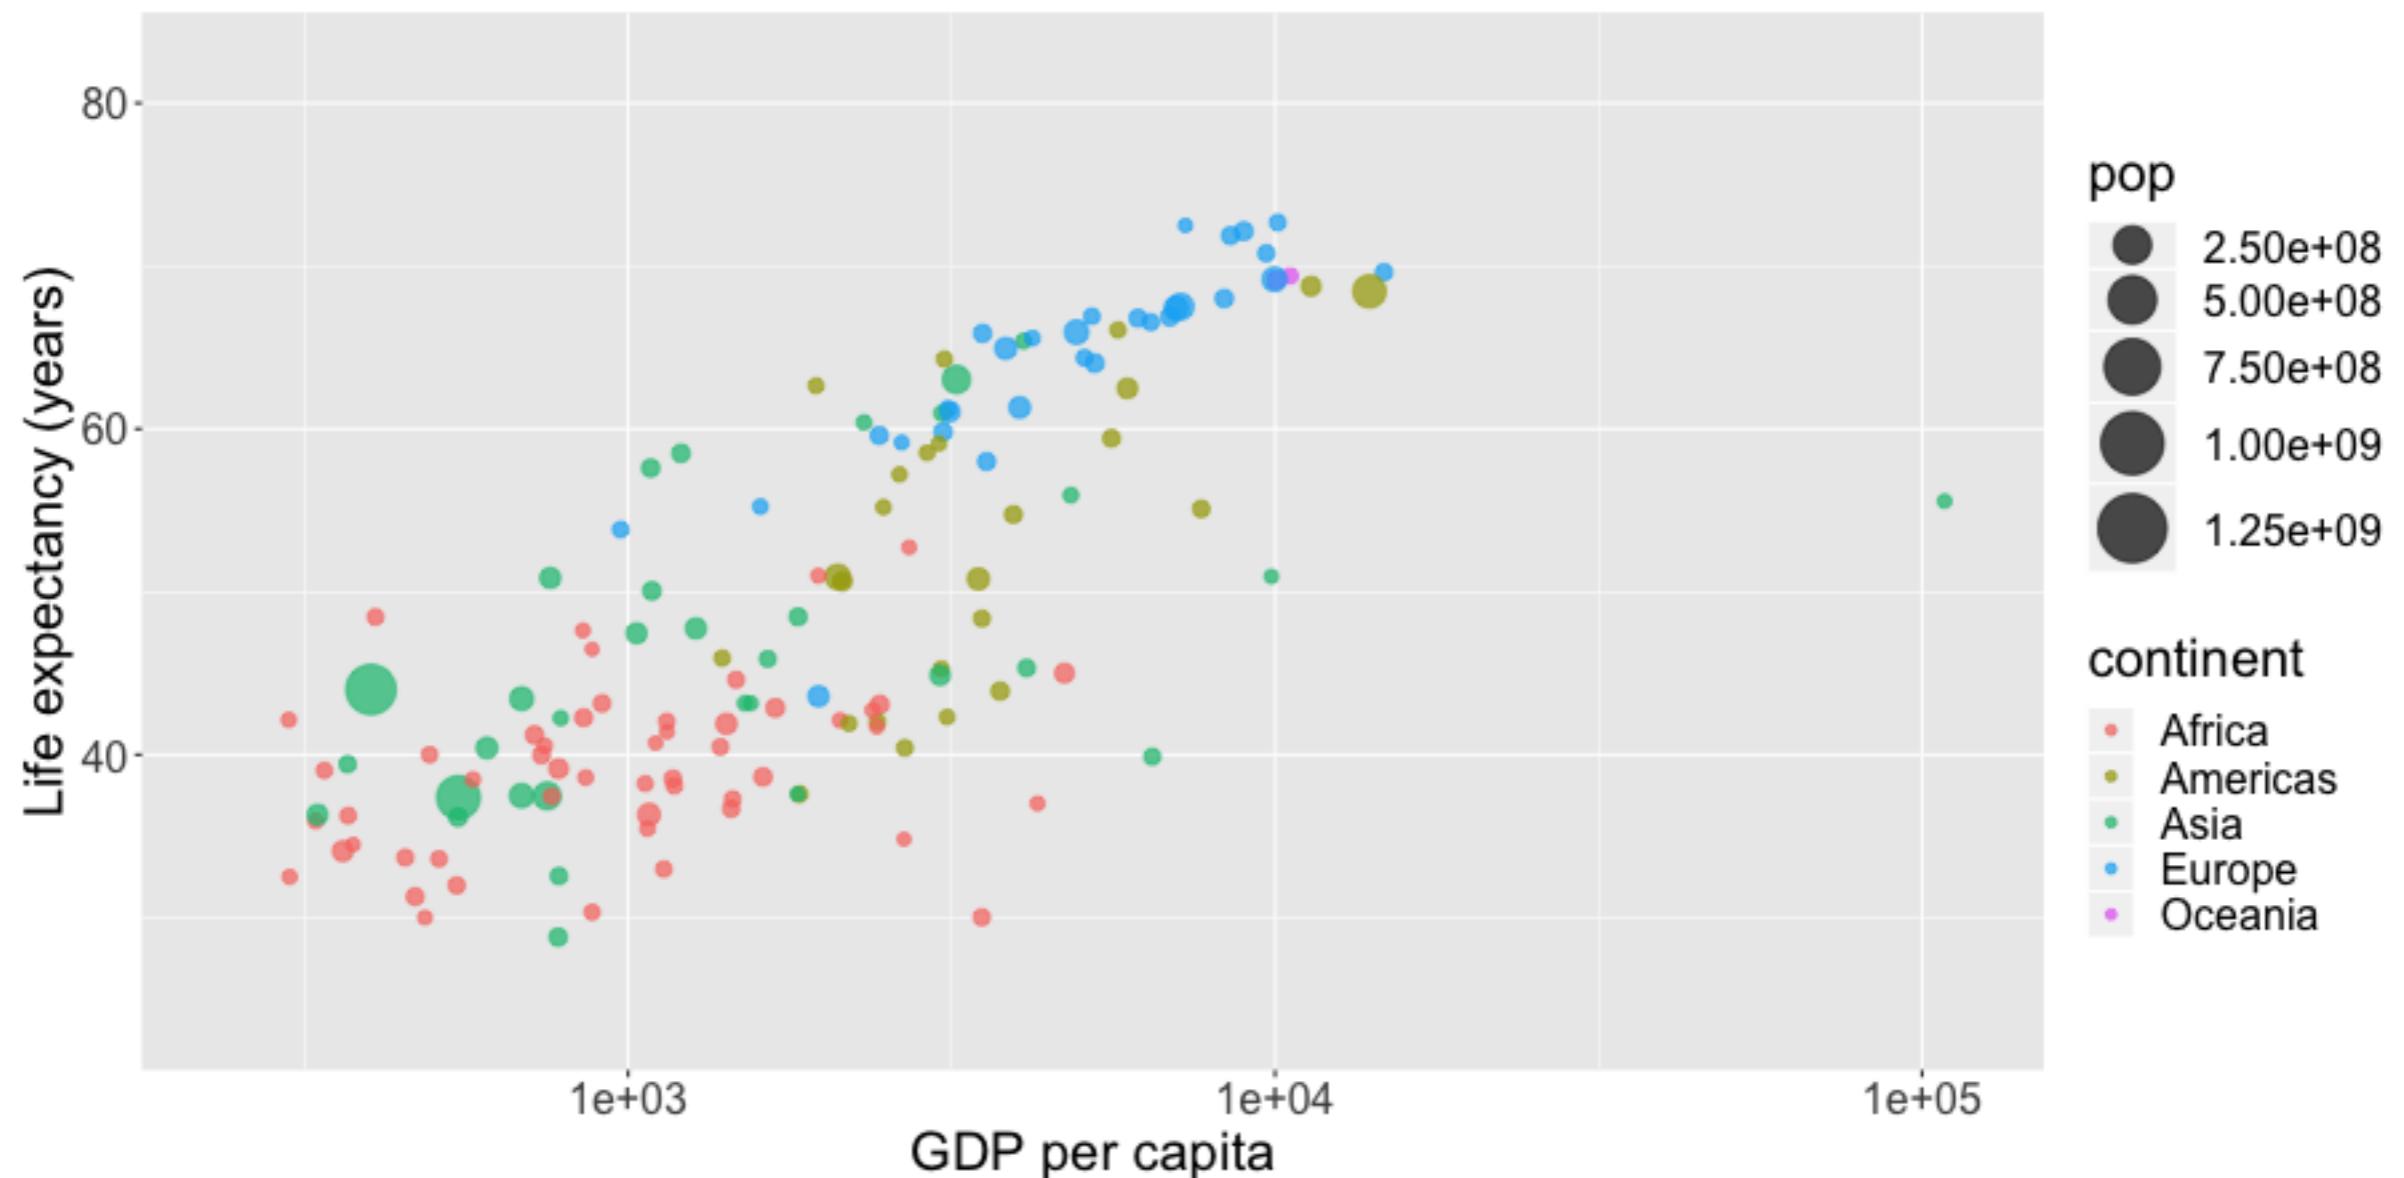
Visualising Data with 5 Variables Simultaneously

Animated Time Series Data

Now with a representation of population size.

Gapminder dataset

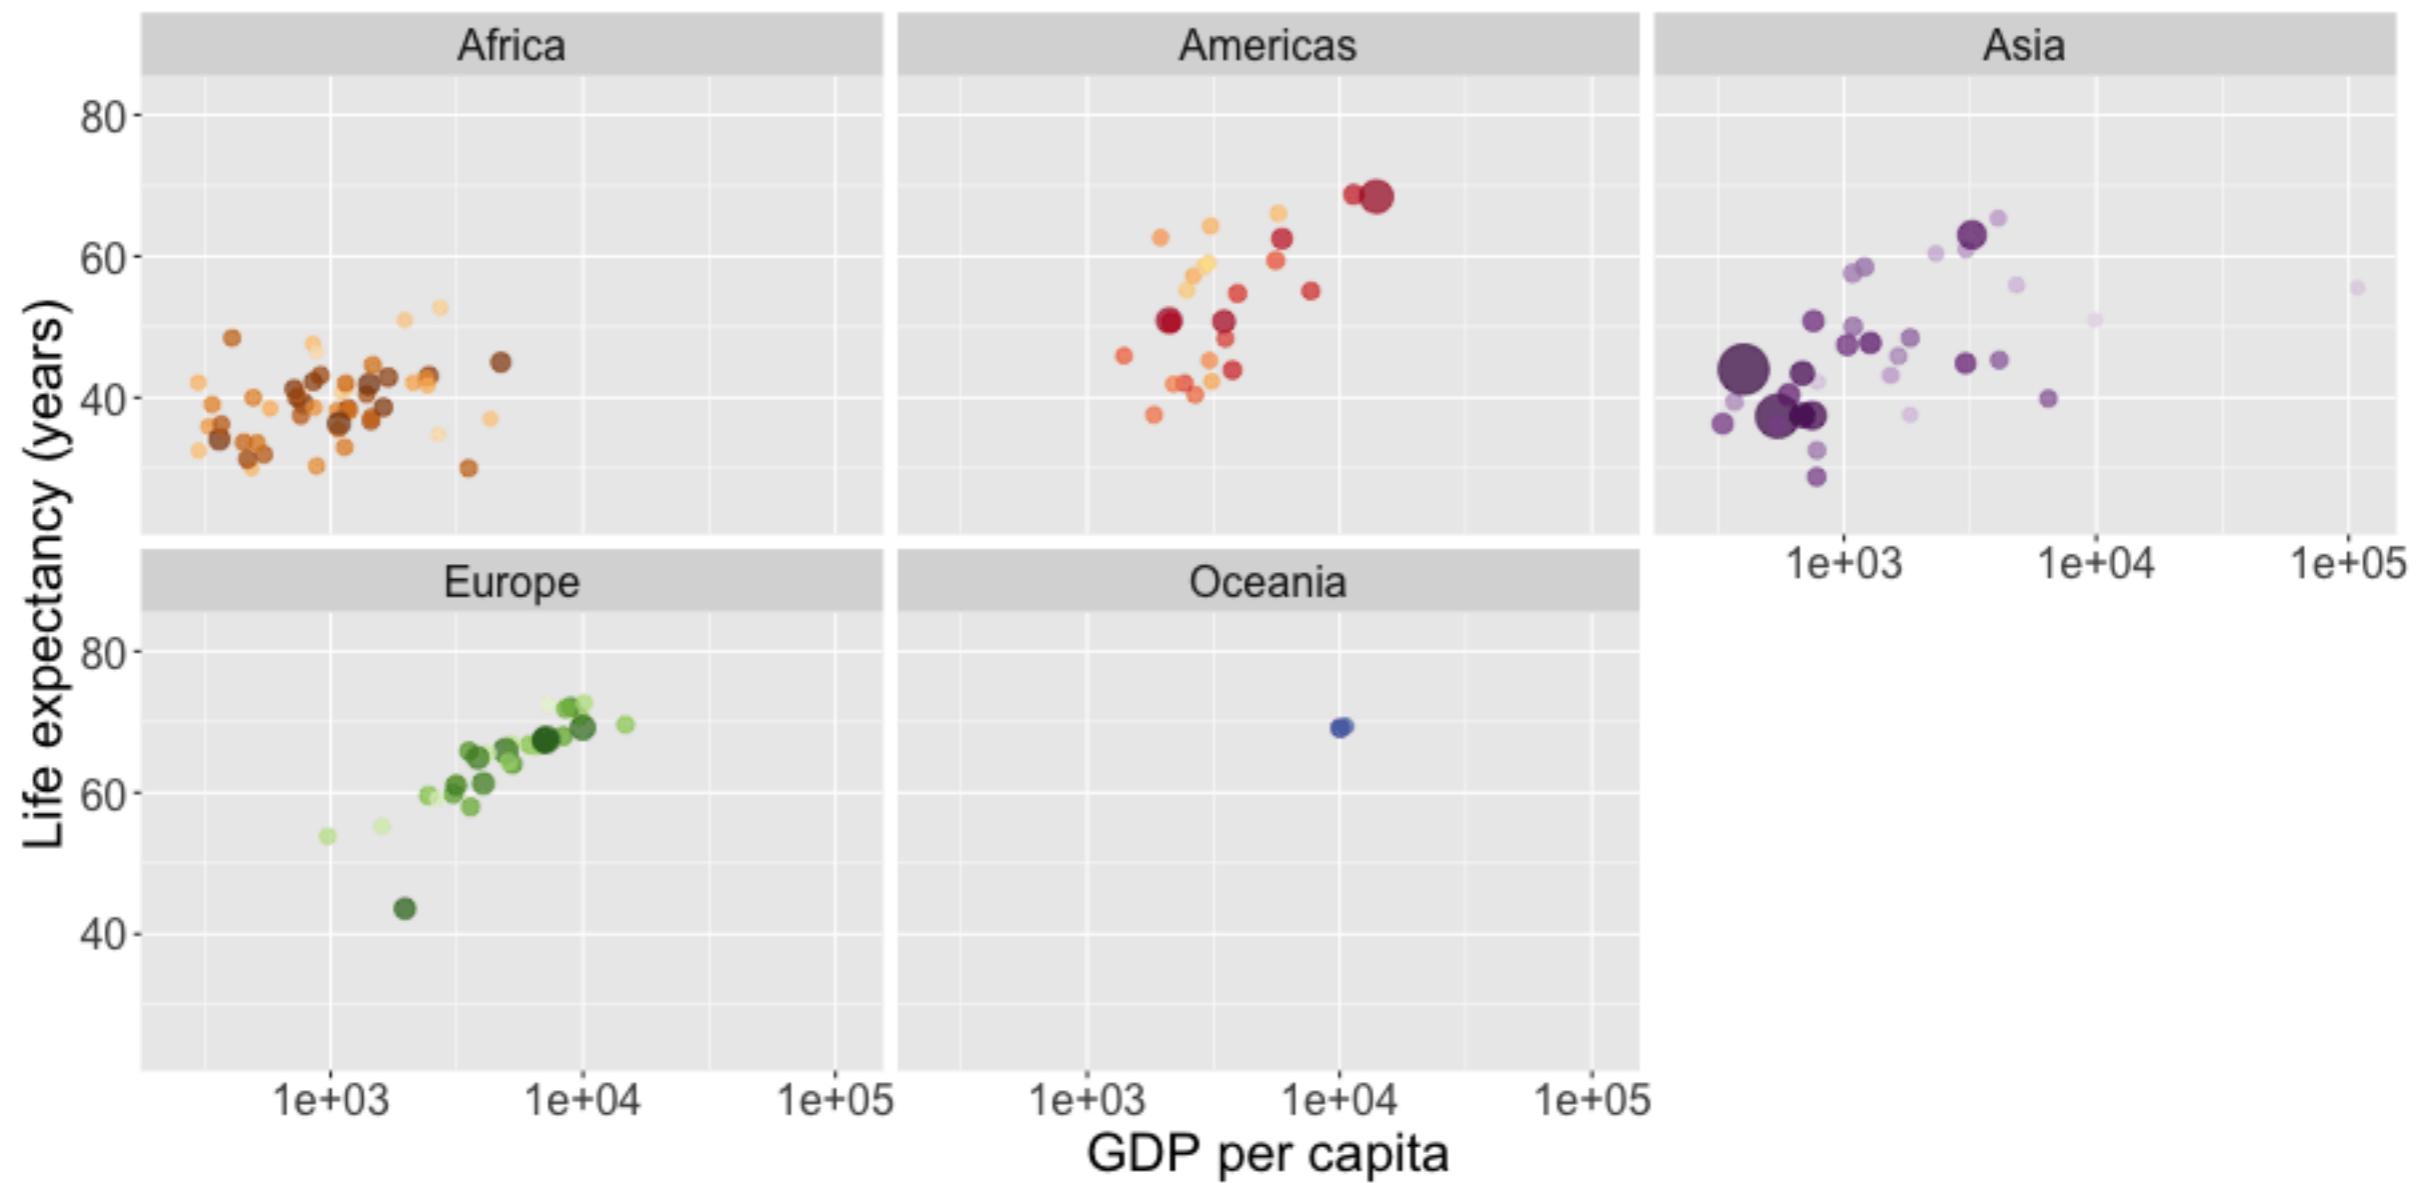
Year: 1952



Separately by Continent

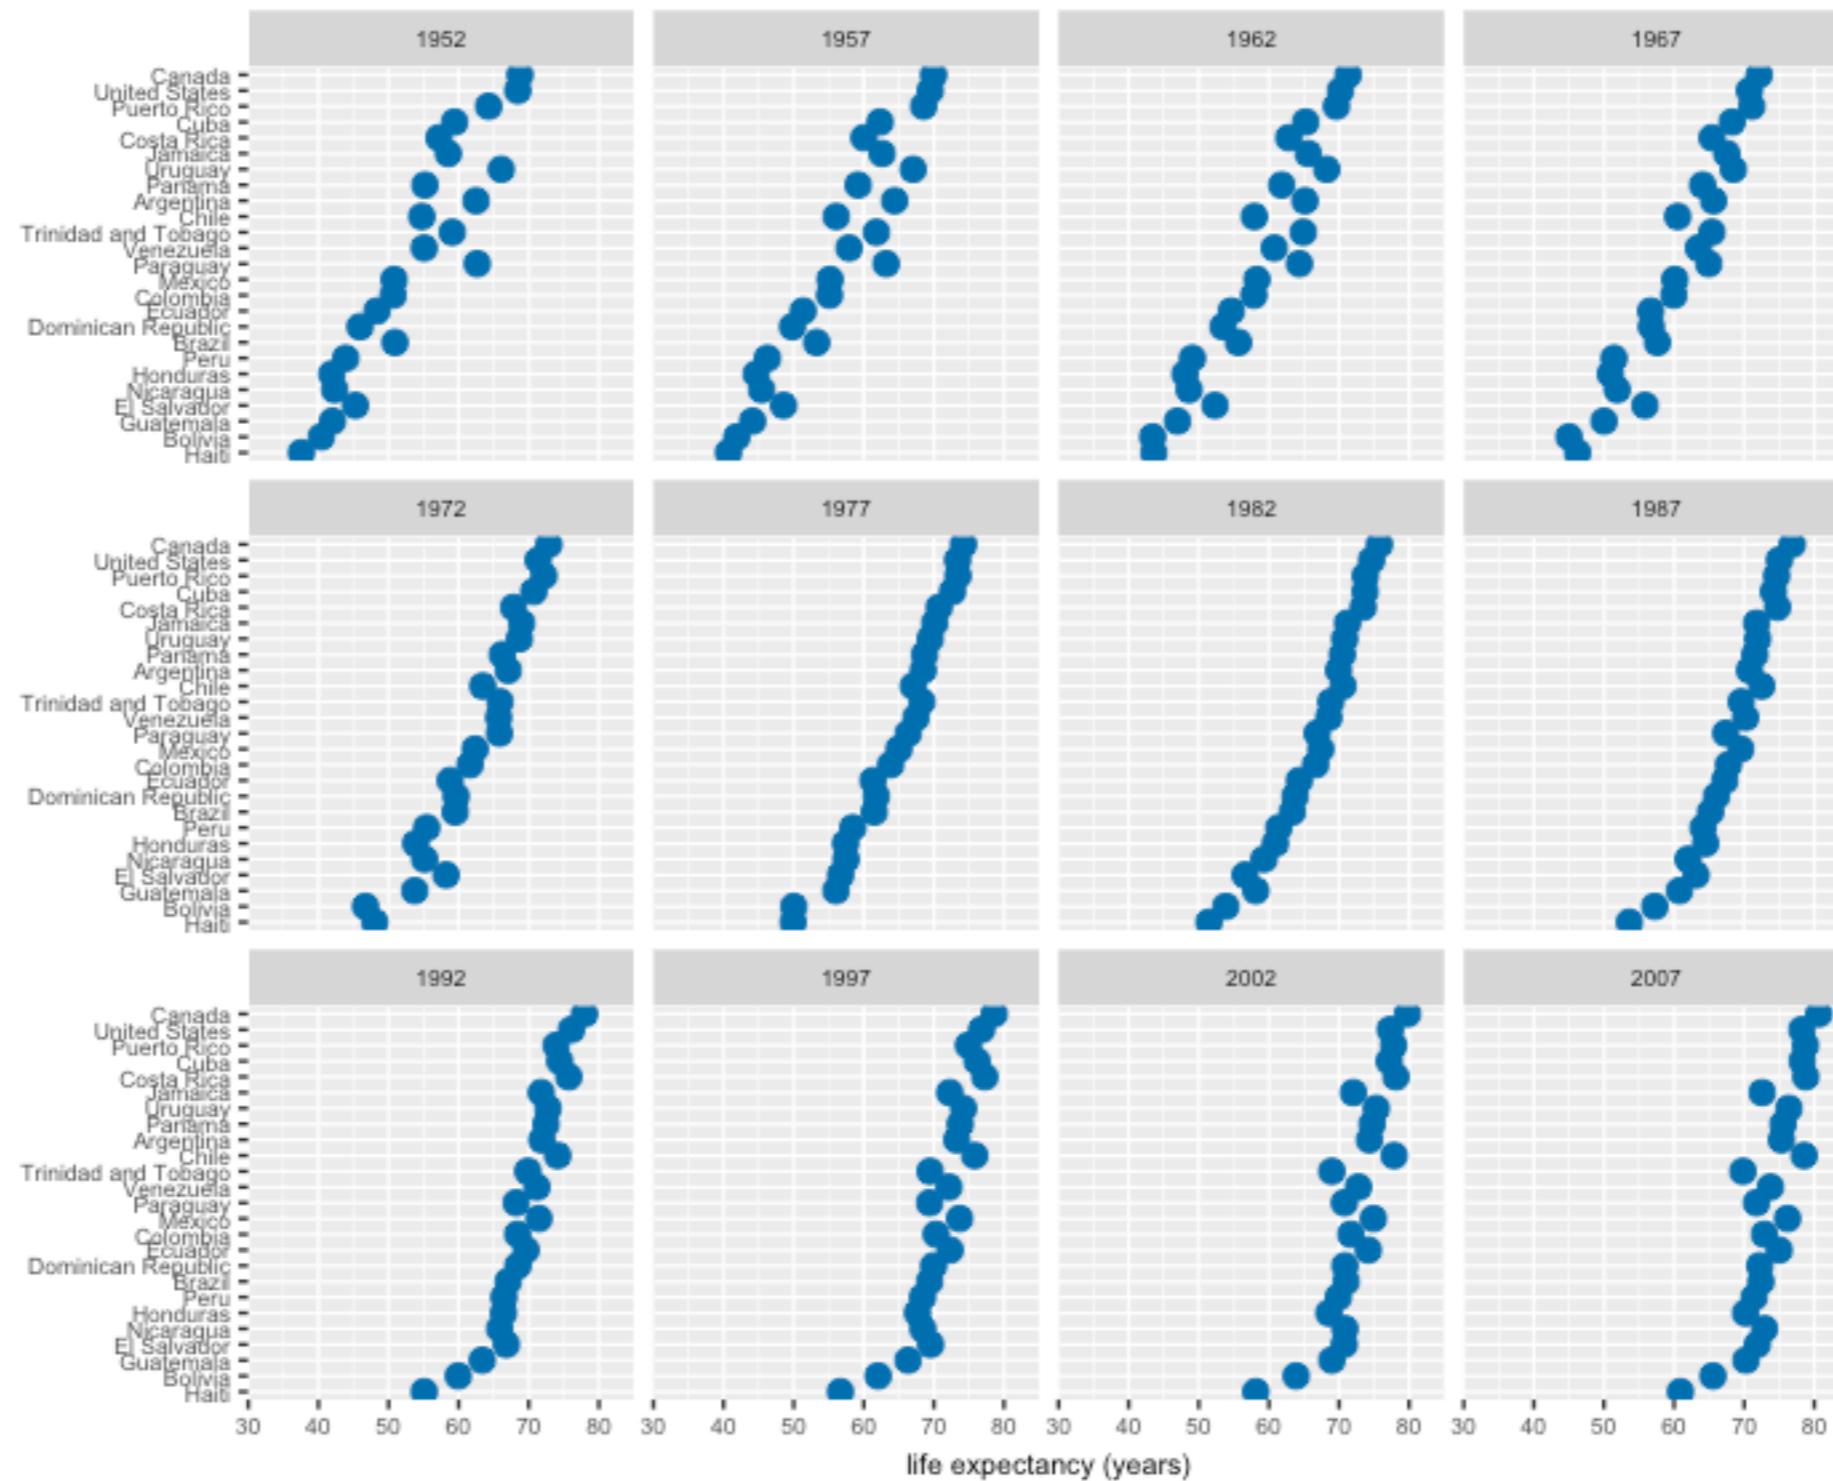
Gapminder dataset

Year: 1952

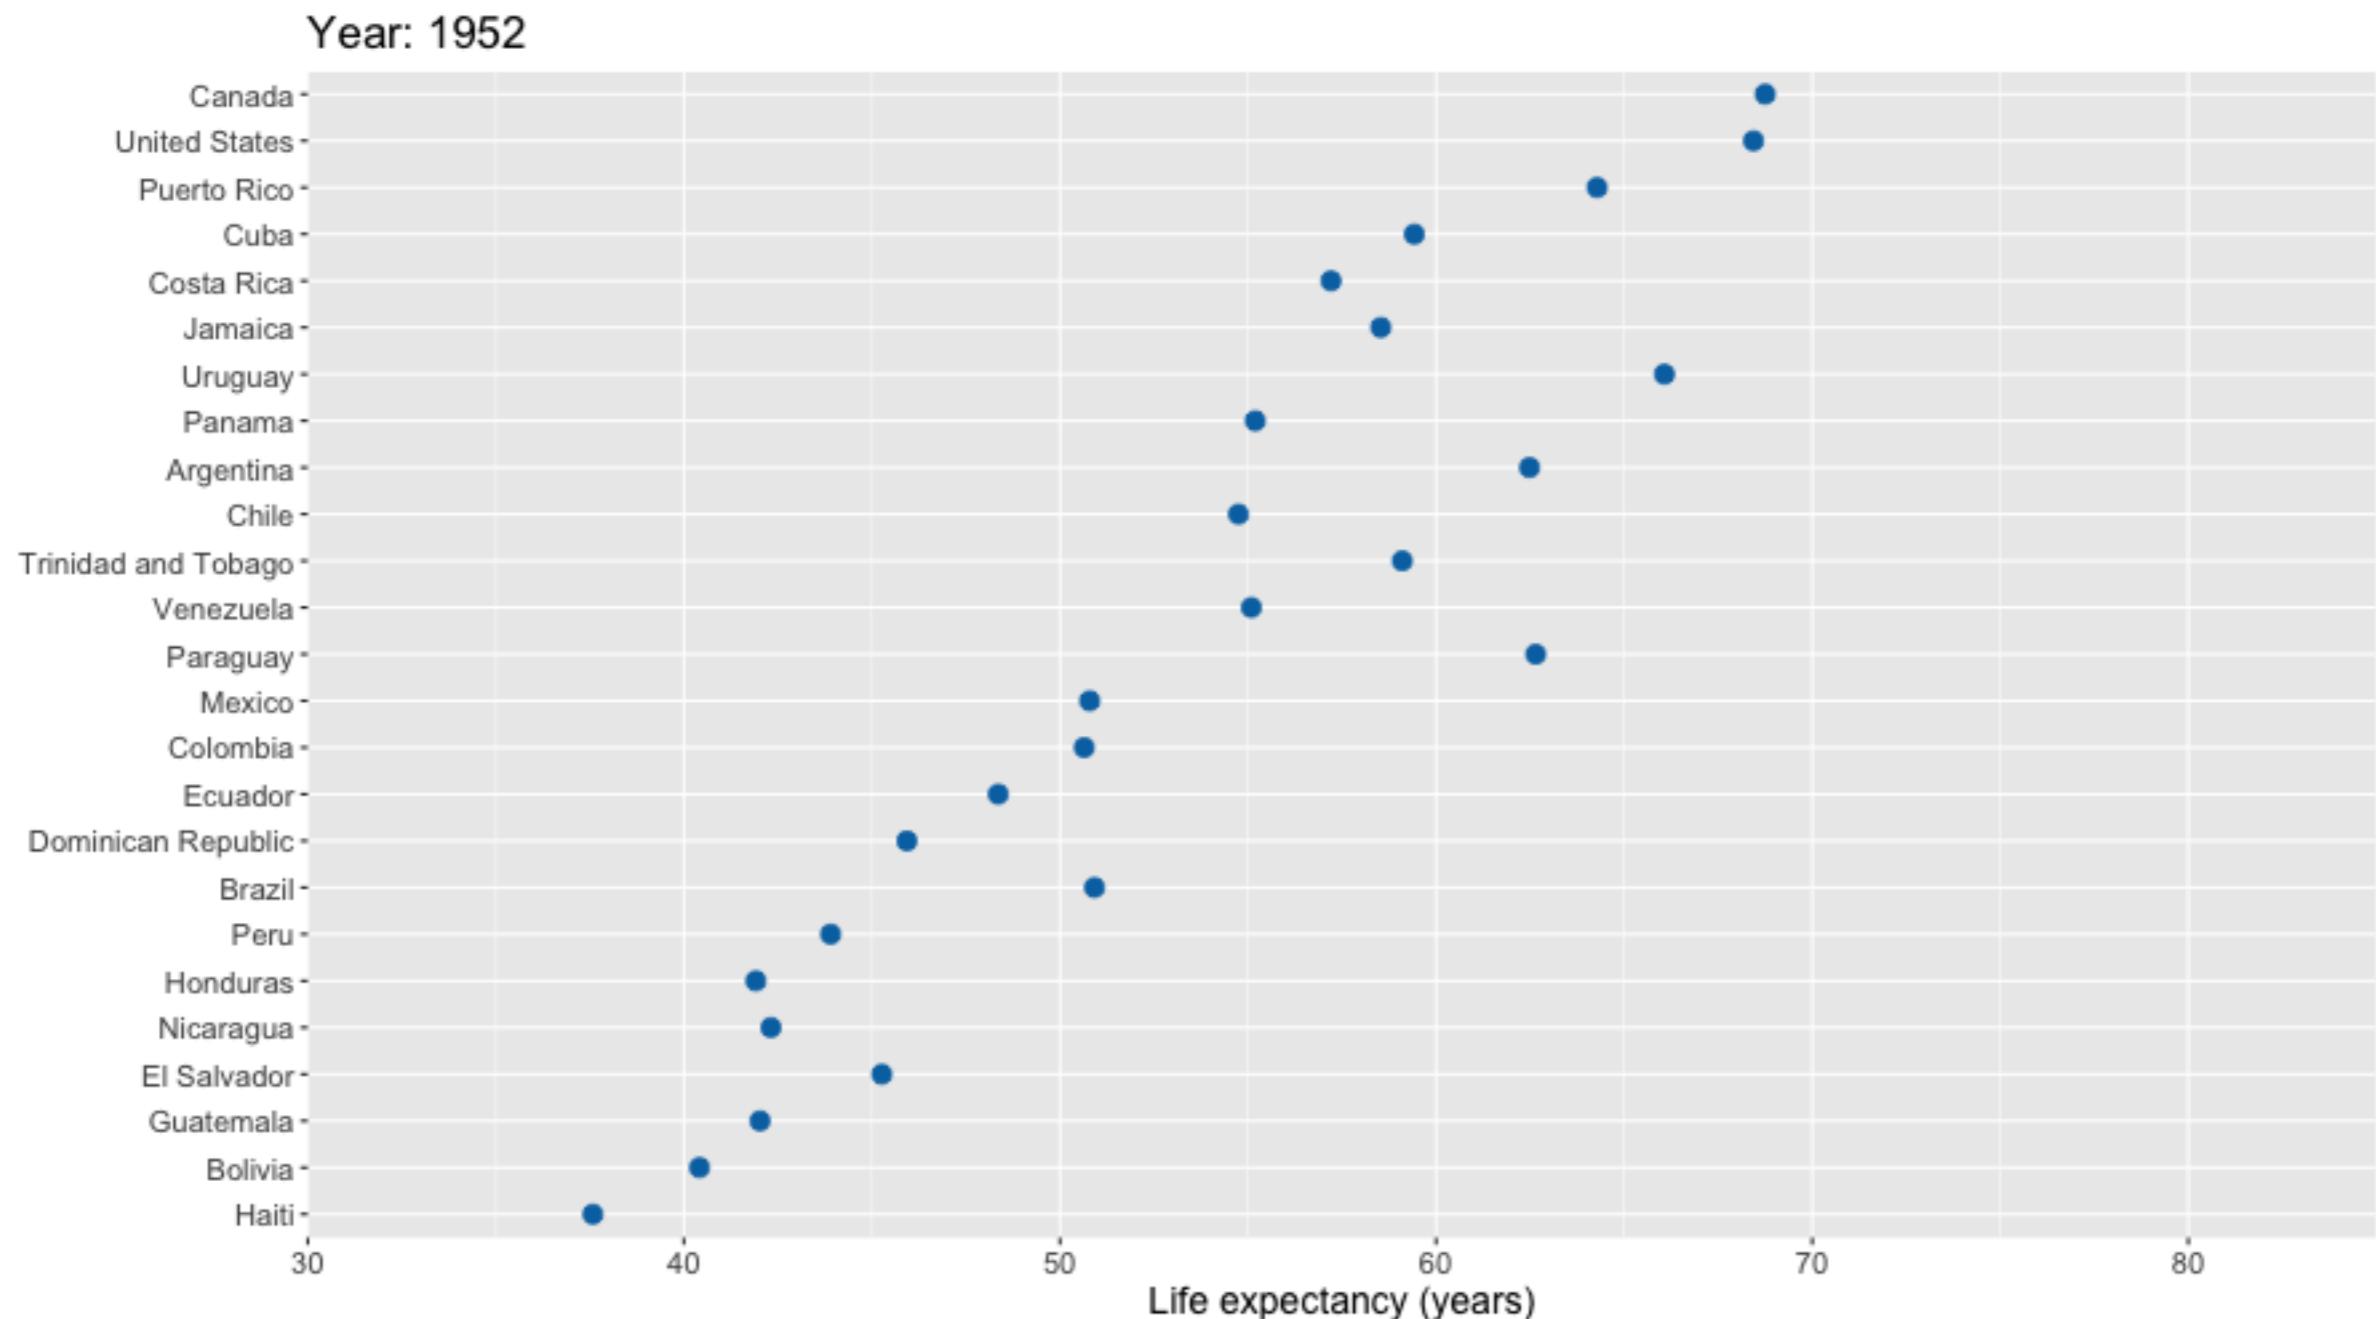


<https://github.com/thomasp85/gganimate>

Life Expectancy - Americas - Static



Life Expectancy - Americas - Animated



Although we have data only once every 5 years, the `gganimate` package interpolates between each census date to provide a smoother animation.

You try...

```
library(gganimate)

df_Americas <- gapminder %>% filter(continent == "Americas")

p <- ggplot(df_Americas, aes(x = lifeExp, y = fct_reorder(country, lifeExp))) +
  geom_point(color = "#0072B2", size = 3) +
  scale_x_continuous(name = "Life expectancy (years)",
                      limits = c(30, 85),
                      expand = c(0, 0)) +
  scale_y_discrete(name = NULL, expand = c(0, 0.5)) +
  labs(title = 'Year: {frame_time}') +
  theme(text = element_text(size = 14)) +
  transition_time(year) +
  ease_aes('linear')

animate(p, height = 500, width = 900)
```

Worksheet 1